

[New Book] Click to get *Python for Machine Learning!*Use the offer code **20offearlybird** to get 20% off.

Navigation



Machine Learning Mastery

Making Developers Awesome at Machine Learning

[Click to Take the FREE Deep Learning Time Series Crash-Course](#)

Search...



Time Series Forecasting with the Long Short-Term Memory Network in Python

by Jason Brownlee on April 7, 2017 in [Deep Learning for Time Series](#)[Tweet](#)[Tweet](#)[Share](#)[Share](#)

Last Updated on August 28, 2020

The Long Short-Term Memory recurrent neural network has the promise of learning long sequences of observations.

It seems a perfect match for [time series forecasting](#), and in fact, it may be.

In this tutorial, you will discover how to develop an LSTM forecast model for a one-step univariate time series forecasting problem.

After completing this tutorial, you will know:

- How to develop a baseline of performance for a forecast problem.
- How to design a robust test harness for one-step time series forecasting.
- How to prepare data, develop, and evaluate an LSTM recurrent neural network for time series forecasting.

Kick-start your project with my new book [Deep Learning for Time Series Forecasting](#), including *step-by-step tutorials* and the *Python source code* files for all examples.

Let's get started.

- **Update May/2017:** Fixed bug in `invert_scale()` function, thanks Max.
- **Updated Apr/2019:** Updated the link to dataset.

[Start Machine Learning](#)



Time Series Forecasting with the Long Short-Term Memory Network
Photo by [Matt MacGillivray](#)

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.
Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

Tutorial Overview

This is a big topic and we are going to cover a lot of ground. Strap in.

This tutorial is broken down into 9 parts; they are:

1. Shampoo Sales Dataset
2. Test Setup
3. Persistence Model Forecast
4. LSTM Data Preparation
5. LSTM Model Development
6. LSTM Forecast
7. Complete LSTM Example
8. Develop a Robust Result
9. Tutorial Extensions

Python Environment

This tutorial assumes you have a Python SciPy environment installed. You can use either Python 2 or 3 with this tutorial.

You must have Keras (2.0 or higher) installed with [either the TensorFlow or Theano backend](#)

[Start Machine Learning](#)

The tutorial also assumes you have scikit-learn, Pandas, NumPy and Matplotlib installed.

If you need help with your environment, see this post:

- [How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda](#)

Need help with Deep Learning for Time Series?

Take my free 7-day email crash course now (with sample code).

Click to sign-up and also get a free

[Download Your](#)

Start Machine Learning X

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Shampoo Sales Dataset

This dataset describes the monthly number of sales of shampoo over a 3-year period.

The units are a sales count and there are 36 observations. The data is from Makridakis, Wheelwright, and Hyndman (1998).

- [Download the dataset.](#)

Download the dataset to your current working directory with the name “shampoo-sales.csv”.

The example below loads and creates a plot of the loaded dataset.

```

1 # load and plot dataset
2 from pandas import read_csv
3 from pandas import datetime
4 from matplotlib import pyplot
5 # load dataset
6 def parser(x):
7     return datetime.strptime('190'+x, '%Y-%m')
8 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True)
9 # summarize first few rows
10 print(series.head())
11 # line plot
12 series.plot()
13 pyplot.show()
```

Running the example loads the dataset as a Pandas Series and prints the first 5 rows.

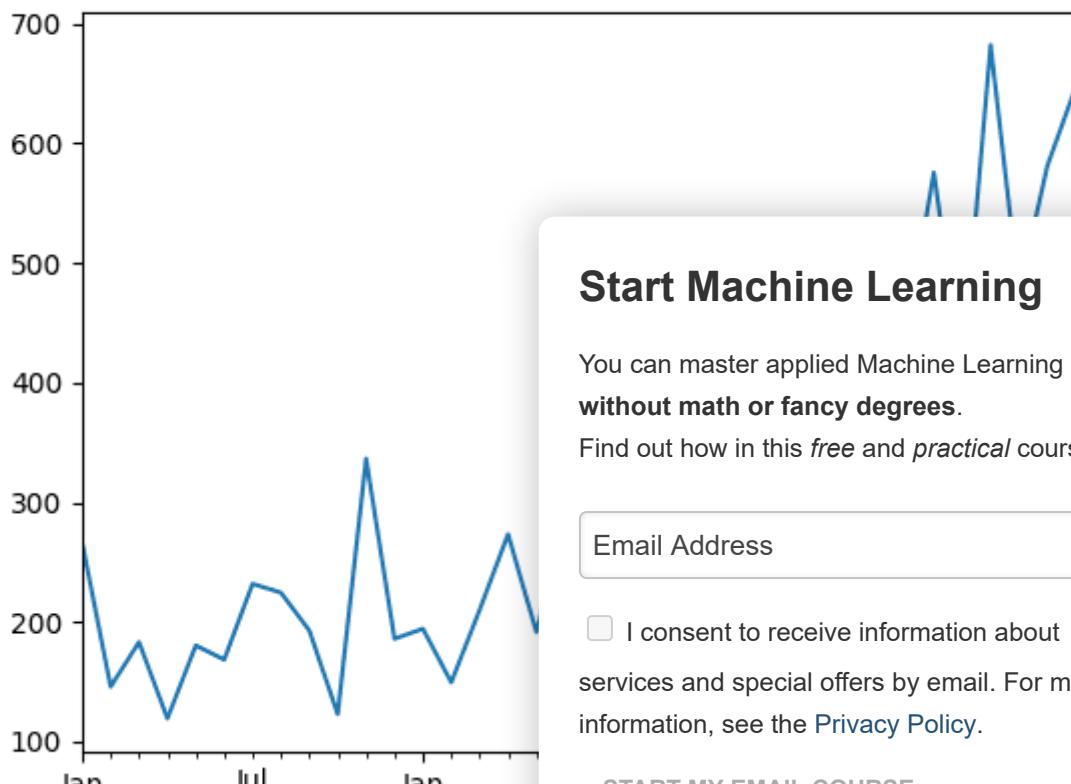
1 Month
2 1901-01-01 266.0
3 1901-02-01 145.9
4 1901-03-01 183.1
5 1901-04-01 119.3

[Start Machine Learning](#)

6 1901-05-01 180.3

7 Name: Sales, dtype: float64

A line plot of the series is then created showing a clear increasing trend.



Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Line Plot of Monthly Shampoo Sales Dataset

Experimental Test Setup

We will split the Shampoo Sales dataset into two parts: a training and a test set.

The first two years of data will be taken for the training dataset and the remaining one year of data will be used for the test set.

For example:

```
1 # split data into train and test
2 X = series.values
3 train, test = X[0:-12], X[-12:]
```

Models will be developed using the training dataset and will make predictions on the test dataset.

A rolling forecast scenario will be used, also called walk-forward model validation.

Each time step of the test dataset will be walked one at a time. A model will be used to make a forecast for the time step, then the actual expected value from the test set will be taken and made available to the model for the forecast on the next time step.

Start Machine Learning

For example:

```
1 # walk-forward validation
2 history = [x for x in train]
3 predictions = list()
4 for i in range(len(test)):
5     # make prediction...
```

This mimics a real-world scenario where new Shampoo Sales observations would be available each month and used in the forecasting of the following month.

Finally, all forecasts on the test dataset will be collected and an error score calculated to summarize the skill of the model. The root mean squared error (RMSE) will be used as it punishes large errors and results in a score that is in the same units as the forecast.

For example:

```
1 from sklearn.metrics import mean_squared_error
2 rmse = sqrt(mean_squared_error(test, predictions))
3 print('RMSE: %.3f' % rmse)
```

Persistence Model Forecast

A good baseline forecast for a time series with a linear trend is the persistence forecast.

The persistence forecast is where the observation is the same as the previous observation at the current time step (t).

We can implement this by taking the last observation from the training data and history accumulated by walk-forward validation and using that to predict the current time step.

For example:

```
1 # make prediction
2 yhat = history[-1]
```

We will accumulate all predictions in an array so that they can be directly compared to the test dataset.

The complete example of the persistence forecast model on the Shampoo Sales dataset is listed below.

```
1 from pandas import read_csv
2 from pandas import datetime
3 from sklearn.metrics import mean_squared_error
4 from math import sqrt
5 from matplotlib import pyplot
6 # load dataset
7 def parser(x):
8     return datetime.strptime('190'+x, '%Y-%m')
9 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True)
10 # split data into train and test
11 X = series.values
12 train, test = X[0:-12], X[-12:]
13 # walk-forward validation
```

[Start Machine Learning](#)

```

14 history = [x for x in train]
15 predictions = list()
16 for i in range(len(test)):
17     # make prediction
18     predictions.append(history[-1])
19     # observation
20     history.append(test[i])
21 # report performance
22 rmse = sqrt(mean_squared_error(test, predictions))
23 print('RMSE: %.3f' % rmse)
24 # line plot of observed vs predicted
25 pyplot.plot(test)
26 pyplot.plot(predictions)
27 pyplot.show()

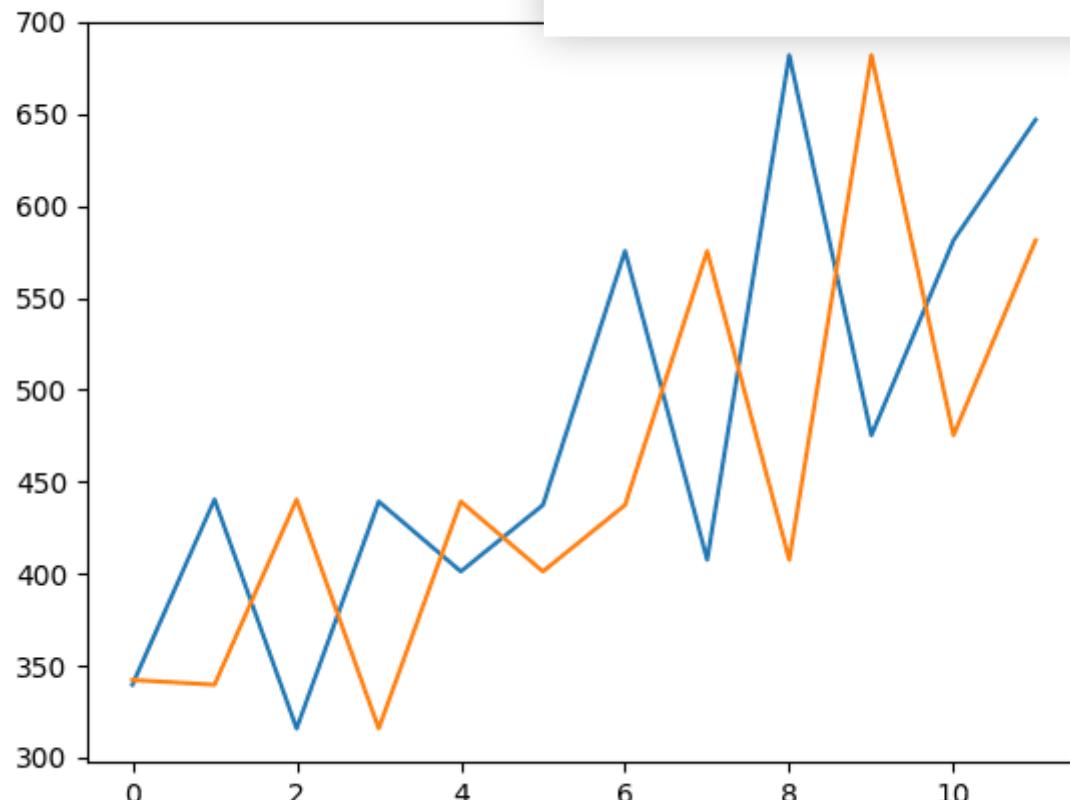
```

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the statistics.

Running the example prints the RMSE of about 136.761 for the test dataset.

1 RMSE: 136.761

A line plot of the test dataset (blue) compared to the persistence model forecast in context.



Persistence Forecast of Observed vs Predicted for Shampoo Sales Dataset

[Start Machine Learning](#)

For more on the persistence model for time series forecasting, see this post:

- [How to Make Baseline Predictions for Time Series Forecasting with Python](#)

Now that we have a baseline of performance on the dataset, we can get started developing an LSTM model for the data.

Need help with LSTMs for Sequence Prediction?

Take my free 7-day email course and discover 6 different LSTM architectures (with code).

Click to sign-up and also get a free

[Start Your FREE](#)

Start Machine Learning X

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

LSTM Data Preparation

Before we can fit an LSTM model to the dataset, we must prepare the data.

This section is broken down into three steps:

1. Transform the time series into a supervised learning problem.
2. Transform the time series data so that it is stationary.
3. Transform the observations to have a specific scale.

Transform Time Series to Supervised Learning

The LSTM model in Keras assumes that your data is divided into input (X) and output (y) components.

For a time series problem, we can achieve this by using the observation from the last time step ($t-1$) as the input and the observation at the current time step (t) as the output.

We can achieve this using the `shift()` function in Pandas that will push all values in a series down by a specified number places. We require a shift of 1 place, which will become the input variables. The time series as it stands will be the output variables.

We can then concatenate these two series together to create a DataFrame ready for supervised learning. The pushed-down series will have a new position at the top with no value. A NaN (not a number) value will be used in this position. We will replace these NaN values with 0 values, which the LSTM model will have to learn as “the start of the series” or “I have no data here,” as a month with zero sales on this dataset has not been observed.

The code below defines a helper function to do this called `timeseries_to_supervised()`. It takes a NumPy array of the raw time series data and a lag or num

[Start Machine Learning](#)

```

1 # frame a sequence as a supervised learning problem
2 def timeseries_to_supervised(data, lag=1):
3     df = DataFrame(data)
4     columns = [df.shift(i) for i in range(1, lag+1)]
5     columns.append(df)
6     df = concat(columns, axis=1)
7     df.fillna(0, inplace=True)
8     return df

```

We can test this function with our loaded Shampoo Sales dataset and convert it into a supervised learning problem.

```

1 from pandas import read_csv
2 from pandas import datetime
3 from pandas import DataFrame
4 from pandas import concat
5
6 # frame a sequence as a supervised learning
7 def timeseries_to_supervised(data, lag=1):
8     df = DataFrame(data)
9     columns = [df.shift(i) for i in range(1, lag+1)]
10    columns.append(df)
11    df = concat(columns, axis=1)
12    df.fillna(0, inplace=True)
13    return df
14
15 # load dataset
16 def parser(x):
17     return datetime.strptime('190'+x, '%Y-%m')
18 series = read_csv('shampoo-sales.csv', header=0, parse_date=True)
19 # transform to supervised learning
20 X = series.values
21 supervised = timeseries_to_supervised(X, 1)
22 print(supervised.head())

```

Start Machine Learning X

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this free and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

Running the example prints the first 5 rows of the new supervised learning problem.

1	0	0
2	0.000000	266.000000
3	266.000000	145.899994
4	145.899994	183.100006
5	183.100006	119.300003
6	119.300003	180.300003

For more information on transforming a time series problem into a supervised learning problem, see the post:

- [Time Series Forecasting as Supervised Learning](#)

Transform Time Series to Stationary

The Shampoo Sales dataset is not stationary.

This means that there is a structure in the data that is dependent on the time. Specifically, there is an increasing trend in the data.

Stationary data is easier to model and will very likely result in more skillful forecasts

[Start Machine Learning](#)

The trend can be removed from the observations, then added back to forecasts later to return the prediction to the original scale and calculate a comparable error score.

A standard way to remove a trend is by differencing the data. That is the observation from the previous time step ($t-1$) is subtracted from the current observation (t). This removes the trend and we are left with a difference series, or the changes to the observations from one time step to the next.

We can achieve this automatically using the `diff()` function in pandas. Alternatively, we can get finer grained control and write our own function to do this, which is preferred for its flexibility in this case.

Below is a function called `difference()` that calculates a differenced series. Note that the first observation in the series is skipped as there is no prior observation with which to calculate a differenced value.

```
1 # create a differenced series
2 def difference(dataset, interval=1):
3     diff = list()
4     for i in range(interval, len(dataset)):
5         value = dataset[i] - dataset[i - interval]
6         diff.append(value)
7     return Series(diff)
```

We also need to invert this process in order to take the difference series back to their original scale.

The function below, called `inverse_difference()`, inverts the differencing operation.

```
1 # invert differenced value
2 def inverse_difference(history, yhat, interval=1):
3     return yhat + history[-interval]
```

We can test out these functions by differencing the whole series, then returning it to the original scale, as follows:

```
1 from pandas import read_csv
2 from pandas import datetime
3 from pandas import Series
4
5 # create a differenced series
6 def difference(dataset, interval=1):
7     diff = list()
8     for i in range(interval, len(dataset)):
9         value = dataset[i] - dataset[i - interval]
10        diff.append(value)
11    return Series(diff)
12
13 # invert differenced value
14 def inverse_difference(history, yhat, interval=1):
15     return yhat + history[-interval]
16
17 # load dataset
18 def parser(x):
19     return datetime.strptime('190' + x, '%Y-%m')
20 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True)
21 print(series.head())
22 # transform to be stationary
23 differenced = difference(series, 1)
24 print(differenced.head())
```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this free and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

```

25 # invert transform
26 inverted = list()
27 for i in range(len(differenced)):
28     value = inverse_difference(series, differenced[i], len(series)-i)
29     inverted.append(value)
30 inverted = Series(inverted)
31 print(inverted.head())

```

Running the example prints the first 5 rows of the loaded data, then the first 5 rows of the differenced series, then finally the first 5 rows with the difference operation inverted.

Note that the first observation in the original dataset was removed from the inverted difference data. Besides that, the last set of data matches the first as expected.

```

1 Month
2 1901-01-01    266.0
3 1901-02-01    145.9
4 1901-03-01    183.1
5 1901-04-01    119.3
6 1901-05-01    180.3
7
8 Name: Sales, dtype: float64
9 0   -120.1
10 1    37.2
11 2   -63.8
12 3    61.0
13 4   -11.8
14 dtype: float64
15
16 0    145.9
17 1    183.1
18 2    119.3
19 3    180.3
20 4    168.5
21 dtype: float64

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

For more information on making the time series stationary and differencing, see the posts:

- [How to Check if Time Series Data is Stationary with Python](#)
- [How to Difference a Time Series Dataset with Python](#)

Transform Time Series to Scale

Like other neural networks, LSTMs expect data to be within the scale of the activation function used by the network.

The default activation function for LSTMs is the hyperbolic tangent (*tanh*), which outputs values between -1 and 1. This is the preferred range for the time series data.

To make the experiment fair, the scaling coefficients (min and max) values must be calculated on the training dataset and applied to scale the test dataset and any forecasts. This is to avoid contaminating the experiment with knowledge from the test dataset, which might give the model a small edge.

We can transform the dataset to the range [-1, 1] using the [MinMaxScaler class](#). Like other scikit-learn transform classes, it requires data provided in a matrix format with rows and columns. Therefore, we must reshape our NumPy arrays before transforming~

[Start Machine Learning](#)

For example:

```

1 # transform scale
2 X = series.values
3 X = X.reshape(len(X), 1)
4 scaler = MinMaxScaler(feature_range=(-1, 1))
5 scaler = scaler.fit(X)
6 scaled_X = scaler.transform(X)

```

Again, we must invert the scale on forecasts to return the values back to the original scale so that the results can be interpreted and a comparable error score can be calculated.

```

1 # invert transform
2 inverted_X = scaler.inverse_transform(scaled_X)

```

Putting all of this together, the example below transforms the loaded dataset.

```

1 from pandas import read_csv
2 from pandas import datetime
3 from pandas import Series
4 from sklearn.preprocessing import MinMaxScaler
5 # load dataset
6 def parser(x):
7     return datetime.strptime('190'+x, '%Y-%m')
8 series = read_csv('shampoo-sales.csv', header=0, parse_date=True,
9 print(series.head())
10 # transform scale
11 X = series.values
12 X = X.reshape(len(X), 1)
13 scaler = MinMaxScaler(feature_range=(-1, 1))
14 scaler = scaler.fit(X)
15 scaled_X = scaler.transform(X)
16 scaled_series = Series(scaled_X[:, 0])
17 print(scaled_series.head())
18 # invert transform
19 inverted_X = scaler.inverse_transform(scaled_X)
20 inverted_series = Series(inverted_X[:, 0])
21 print(inverted_series.head())

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this free and *practical* course.

Email Address

I consent to receive information about
services and special offers by email. For more
information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

Running the example first prints the first 5 rows of the loaded data, then the first 5 rows of the scaled data, then the first 5 rows with the scale transform inverted, matching the original data.

```

1 Month
2 1901-01-01    266.0
3 1901-02-01    145.9
4 1901-03-01    183.1
5 1901-04-01    119.3
6 1901-05-01    180.3
7
8 Name: Sales, dtype: float64
9 0   -0.478585
10 1   -0.905456
11 2   -0.773236
12 3   -1.000000
13 4   -0.783188
14 dtype: float64
15
16 0    266.0
17 1    145.9
18 2    183.1

```

[Start Machine Learning](#)

```
19 3    119.3
20 4    180.3
21 dtype: float64
```

Now that we know how to prepare data for the LSTM network, we can start developing our model.

LSTM Model Development

The Long Short-Term Memory network (LSTM) is a type of Recurrent Neural Network (RNN).

A benefit of this type of network is that it can learn and remember over long sequences and does not rely on a pre-specified window lagged observation as input.

In Keras, this is referred to as stateful, and involves defining an LSTM layer.

By default, an LSTM layer in Keras maintains state across a fixed-sized number of rows from the training data by updating the weights of the network. State in the LSTM layer is therefore cleared, by calling the `reset_states()` method.

The LSTM layer expects input to be in a matrix with:

- **Samples:** These are independent observations.
- **Time steps:** These are separate time steps of the sequence.
- **Features:** These are separate measures observed at each time step.

Start Machine Learning

You can master applied Machine Learning as a **stateless** machine learning developer without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

We have some flexibility in how the Shampoo Sales dataset is framed for the network. We will keep it simple and frame the problem as each time step in the original sequence is one separate sample, with one timestep and one feature.

Given that the training dataset is defined as X inputs and y outputs, it must be reshaped into the Samples/TimeSteps/Features format, for example:

```
1 X, y = train[:, 0:-1], train[:, -1]
2 X = X.reshape(X.shape[0], 1, X.shape[1])
```

The shape of the input data must be specified in the LSTM layer using the “`batch_input_shape`” argument as a tuple that specifies the expected number of observations to read each batch, the number of time steps, and the number of features.

The batch size is often much smaller than the total number of samples. It, along with the number of epochs, defines how quickly the network learns the data (how often the weights are updated).

The final import parameter in defining the LSTM layer is the number of neurons, also called the number of memory units or blocks. This is a reasonably simple problem and a number between 1 and 5 should be sufficient.

[Start Machine Learning](#)

The line below creates a single LSTM hidden layer that also specifies the expectations of the input layer via the “*batch_input_shape*” argument.

```
1 layer = LSTM(neurons, batch_input_shape=(batch_size, X.shape[1], X.shape[2]), stateful=True)
```

The network requires a single neuron in the output layer with a linear activation to predict the number of shampoo sales at the next time step.

Once the network is specified, it must be compiled into an efficient symbolic representation using a backend mathematical library, such as TensorFlow or Theano.

In compiling the network, we must specify a loss function “*mean_squared_error*” as the loss function as it closely matches the error of a linear model and the efficient ADAM optimization algorithm.

Using the Sequential Keras API to define the network is simple. We can define the layers of the network.

```
1 model = Sequential()
2 model.add(LSTM(neurons, batch_input_shape=(batch_size, X.shape[1], X.shape[2]), stateful=True))
3 model.add(Dense(1))
4 model.compile(loss='mean_squared_error', optimizer='adam')
```

Once compiled, it can be fit to the training data. Because the LSTM is a stateful model, the internal state is reset. Therefore, we must manually manage the internal state of the model over time across the desired number of epochs.

By default, the samples within an epoch are shuffled prior to being exposed to the network. Again, this is undesirable for the LSTM because we want the network to build up state as it learns across the sequence of observations. We can disable the shuffling of samples by setting “*shuffle*” to “*False*”.

Also by default, the network reports a lot of debug information about the learning progress and skill of the model at the end of each epoch. We can disable this by setting the “*verbose*” argument to the level of “0”.

We can then reset the internal state at the end of the training epoch, ready for the next training iteration.

Below is a loop that manually fits the network to the training data.

```
1 for i in range(nb_epoch):
2     model.fit(X, y, epochs=1, batch_size=batch_size, verbose=0, shuffle=False)
3     model.reset_states()
```

Putting this all together, we can define a function called *fit_lstm()* that trains and returns an LSTM model. As arguments, it takes the training dataset in a supervised learning format, a batch size, a number of epochs, and a number of neurons.

```
1 def fit_lstm(train, batch_size, nb_epoch,
2             X, y = train[:, 0:-1], train[:, -1])
```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this free and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

```

3     X = X.reshape(X.shape[0], 1, X.shape[1])
4     model = Sequential()
5     model.add(LSTM(neurons, batch_input_shape=(batch_size, X.shape[1], X.shape[2]), stateful=True))
6     model.add(Dense(1))
7     model.compile(loss='mean_squared_error', optimizer='adam')
8     for i in range(nb_epoch):
9         model.fit(X, y, epochs=1, batch_size=batch_size, verbose=0, shuffle=False)
10    model.reset_states()
11
return model

```

The batch_size must be set to 1. This is because it must be a factor of the size of the training and test datasets.

The `predict()` function on the model is also constrained by the batch size; there it must be set to 1 because we are interested in making one-step forecasts on the test data.

We will not tune the network parameters in this tutorial. Instead, we will find what works best by trial and error:

- Batch Size: 1
- Epochs: 3000
- Neurons: 4

As an extension to this tutorial, you might like to experiment with different parameter values to improve performance.

- **Update:** Consider trying 1500 epochs and 1 neuron.

Next, we will look at how we can use a fit LSTM model to make forecasts.

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this free and practical course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

LSTM Forecast

Once the LSTM model is fit to the training data, it can be used to make forecasts.

Again, we have some flexibility. We can decide to fit the model once on all of the training data, then predict each new time step one at a time from the test data (we'll call this the fixed approach), or we can re-fit the model or update the model each time step of the test data as new observations from the test data are made available (we'll call this the dynamic approach).

In this tutorial, we will go with the fixed approach for its simplicity, although, we would expect the dynamic approach to result in better model skill.

To make a forecast, we can call the `predict()` function on the model. This requires a 3D NumPy array input as an argument. In this case, it will be an array of one value, the observation at the previous time step.

The `predict()` function returns an array of predictions, one for each input row provided. Because we are providing a single input, the output will be a 2D NumPy array with one value.

We can capture this behavior in a function named `forecast()` listed below. Given a fit model, a batch-size used when fitting the model (e.g. 1), and a row from the test data, we can take the last input data from the test row, reshape it, and return the prediction.

```

1 def forecast(model, batch_size, row):
2     X = row[0:-1]
3     X = X.reshape(1, 1, len(X))
4     yhat = model.predict(X, batch_size=batch_size)
5     return yhat[0,0]

```

During training, the internal state is reset after each epoch. While forecasting, we will not want to reset the internal state between forecasts. In fact, **we would like the model to build up state as we forecast each time step in the test dataset.**

This raises the question as to what would be a good initial state for the network prior to forecasting the test dataset.

In this tutorial, **we will seed the state by making a prediction on the first time step.** This way, the internal state should be set up ready to receive the second time step.

We now have all of the pieces to fit an LSTM Network to the Shampoo Sales dataset and evaluate its performance.

In the next section, we will put all of these pieces together to complete the example.

Complete LSTM Example

In this section, we will fit an LSTM to the Shampoo Sales dataset.

This will involve drawing together all of the elements from the previous sections. Let's review:

1. Load the dataset from CSV file.
2. Transform the dataset to make it suitable for the LSTM model, including:
 1. Transforming the data to a supervised learning problem.
 2. Transforming the data to be stationary.
 3. Transforming the data so that it has the scale -1 to 1.
3. Fitting a stateful LSTM network model to the training data.
4. Evaluating the static LSTM model on the test data.
5. Report the performance of the forecasts.

Some things to note about the example:

- The scaling and inverse scaling behaviors have been moved to the functions `scale()` and `invert_scale()` for brevity.
- The test data is scaled using the fit of the scaler on the training data, as is required to ensure the min/max values of the test data do not influence the model.
- The order of data transforms was adjusted for convenience to first make the data stationary, then a supervised learning problem, then scaled.
- Differencing was performed on the entire dataset prior to splitting into train and test sets for convenience. We could just as easily collect observations during the walk-forward validation and difference them as we go. I decided against it for readability.

Start Machine Learning

You can master applied Machine Learning without **math or fancy degrees**.

Find out how in this **free** and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

so

The complete example is listed below.

```

1 from pandas import DataFrame
2 from pandas import Series
3 from pandas import concat
4 from pandas import read_csv
5 from pandas import datetime
6 from sklearn.metrics import mean_squared_error
7 from sklearn.preprocessing import MinMaxScaler
8 from keras.models import Sequential
9 from keras.layers import Dense
10 from keras.layers import LSTM
11 from math import sqrt
12 from matplotlib import pyplot
13 import numpy
14
15 # date-time parsing function for loading
16 def parser(x):
17     return datetime.strptime('190'+x, '%Y-%m-%d %H:%M:%S')
18
19 # frame a sequence as a supervised learning problem
20 def timeseries_to_supervised(data, lag=1):
21     df = DataFrame(data)
22     columns = [df.shift(i) for i in range(lag, 0, -1)]
23     columns.append(df)
24     df = concat(columns, axis=1)
25     df.fillna(0, inplace=True)
26     return df
27
28 # create a differenced series
29 def difference(dataset, interval=1):
30     diff = list()
31     for i in range(interval, len(dataset)):
32         value = dataset[i] - dataset[i - interval]
33         diff.append(value)
34     return Series(diff)
35
36 # invert differenced value
37 def inverse_difference(history, yhat, interval=1):
38     return yhat + history[-interval]
39
40 # scale train and test data to [-1, 1]
41 def scale(train, test):
42     # fit scaler
43     scaler = MinMaxScaler(feature_range=(-1, 1))
44     scaler = scaler.fit(train)
45     # transform train
46     train = train.reshape(train.shape[0], train.shape[1])
47     train_scaled = scaler.transform(train)
48     # transform test
49     test = test.reshape(test.shape[0], test.shape[1])
50     test_scaled = scaler.transform(test)
51     return scaler, train_scaled, test_scaled
52
53 # inverse scaling for a forecasted value
54 def invert_scale(scaler, X, value):
55     new_row = [x for x in X] + [value]
56     array = numpy.array(new_row)
57     array = array.reshape(1, len(array))
58     inverted = scaler.inverse_transform(array)
59     return inverted[0, -1]
60
61 # fit an LSTM network to training data
62 def fit_lstm(train, batch_size, nb_epoch, neurons):
63     X, y = train[:, :-1], train[:, -1]
64     X = X.reshape(X.shape[0], 1, X.shape[1])
65     model = Sequential()

```

Start Machine Learning



You can master applied Machine Learning
without math or fancy degrees.

Find out how in this free and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

[Start Machine Learning](#)

```

66     model.add(LSTM(neurons, batch_input_shape=(batch_size, X.shape[1], X.shape[2]), stateful=True))
67     model.add(Dense(1))
68     model.compile(loss='mean_squared_error', optimizer='adam')
69     for i in range(nb_epoch):
70         model.fit(X, y, epochs=1, batch_size=batch_size, verbose=0, shuffle=False)
71         model.reset_states()
72     return model
73
74 # make a one-step forecast
75 def forecast_lstm(model, batch_size, X):
76     X = X.reshape(1, 1, len(X))
77     yhat = model.predict(X, batch_size=batch_size)
78     return yhat[0, 0]
79
80 # load dataset
81 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True)
82
83 # transform data to be stationary
84 raw_values = series.values
85 diff_values = difference(raw_values, 1)
86
87 # transform data to be supervised learning
88 supervised = timeseries_to_supervised(diff_values)
89 supervised_values = supervised.values
90
91 # split data into train and test-sets
92 train, test = supervised_values[0:-12], supervised_values[-12:]
93
94 # transform the scale of the data
95 scaler, train_scaled, test_scaled = scale(train, test)
96
97 # fit the model
98 lstm_model = fit_lstm(train_scaled, 1, 30)
99 # forecast the entire training dataset to
100 train_reshaped = train_scaled[:, 0].reshape(-1, 1, 1)
101 lstm_model.predict(train_reshaped, batch_size=1)
102
103 # walk-forward validation on the test data
104 predictions = list()
105 for i in range(len(test_scaled)):
106     # make one-step forecast
107     X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
108     yhat = forecast_lstm(lstm_model, 1, X)
109     # invert scaling
110     yhat = invert_scale(scaler, X, yhat)
111     # invert differencing
112     yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)
113     # store forecast
114     predictions.append(yhat)
115     expected = raw_values[len(train) + i + 1]
116     print('Month=%d, Predicted=%f, Expected=%f' % (i+1, yhat, expected))
117
118 # report performance
119 rmse = sqrt(mean_squared_error(raw_values[-12:], predictions))
120 print('Test RMSE: %.3f' % rmse)
121 # line plot of observed vs predicted
122 pyplot.plot(raw_values[-12:])
123 pyplot.plot(predictions)
124 pyplot.show()

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this free and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

Running the example prints the expected and predicted values for each of the 12 months in the test dataset.

Note: Your results may vary given the stochastic nature of the algorithm or evaluation procedure, or differences in numerical precision. Consider running the example a few times and compare the average outcome.

[Start Machine Learning](#)

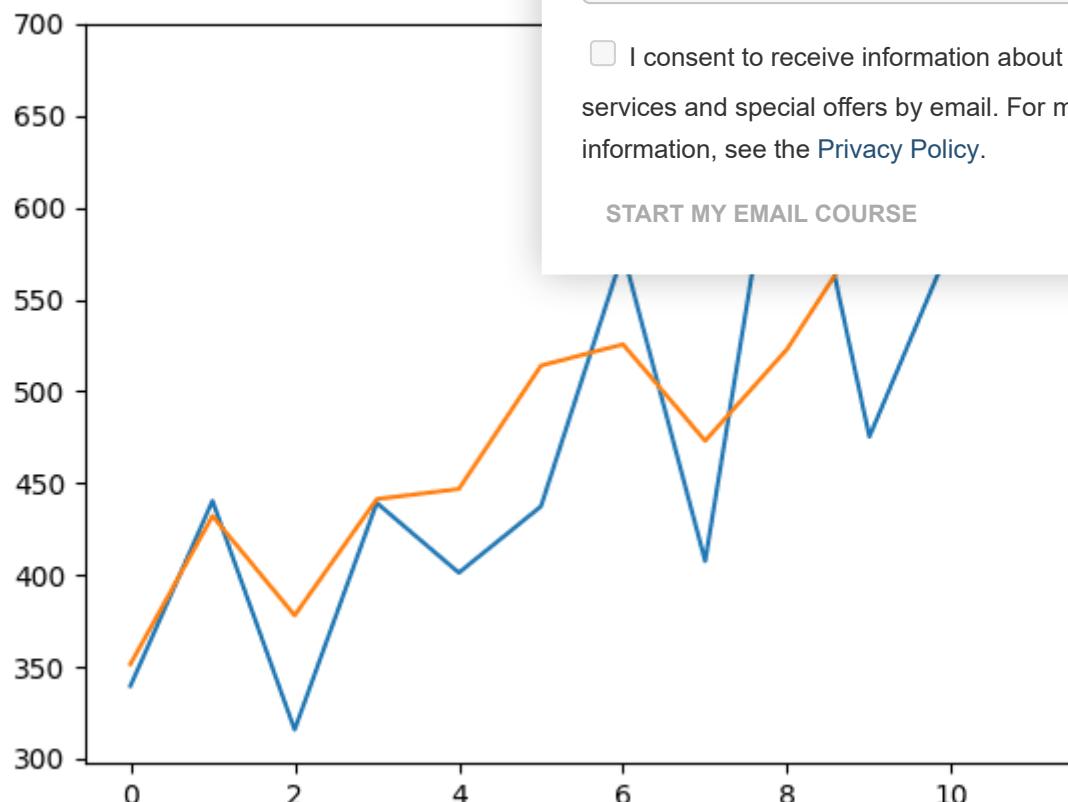
The example also prints the RMSE of all forecasts. The model shows an RMSE of 71.721 monthly shampoo sales, which is better than the persistence model that achieved an RMSE of 136.761 shampoo sales.

```

1 Month=1, Predicted=351.582196, Expected=339.700000
2 Month=2, Predicted=432.169667, Expected=440.400000
3 Month=3, Predicted=378.064505, Expected=315.900000
4 Month=4, Predicted=441.370077, Expected=439.300000
5 Month=5, Predicted=446.872627, Expected=401.300000
6 Month=6, Predicted=514.021244, Expected=437.400000
7 Month=7, Predicted=525.608903, Expected=575.500000
8 Month=8, Predicted=473.072365, Expected=407.600000
9 Month=9, Predicted=523.126979, Expected=682.000000
10 Month=10, Predicted=592.274106, Expected=475.300000
11 Month=11, Predicted=589.299863, Expected=5
12 Month=12, Predicted=584.149152, Expected=6
13 Test RMSE: 71.721

```

A line plot of the test data (blue) vs the predicted values (orange) will help you visualize the model skill.



Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this free and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

As an alternate, you can do a quick experiment to build your trust in the test harness and all of the transforms and inverse transforms.

Comment out the line that fits the LSTM model in walk-forward validation:

```
1 yhat = forecast_lstm(lstm_model, 1, X)
```

Start Machine Learning

And replace it with the following:

```
1 yhat = y
```

This should produce a model with perfect skill (e.g. a model that predicts the expected outcome as the model output).

The results should look as follows, showing that if the LSTM model could predict the series perfectly, the inverse transforms and error calculation would show it correctly.

```
1 Month=1, Predicted=339.700000, Expected=339.700000
2 Month=2, Predicted=440.400000, Expected=44
3 Month=3, Predicted=315.900000, Expected=31
4 Month=4, Predicted=439.300000, Expected=43
5 Month=5, Predicted=401.300000, Expected=40
6 Month=6, Predicted=437.400000, Expected=43
7 Month=7, Predicted=575.500000, Expected=57
8 Month=8, Predicted=407.600000, Expected=40
9 Month=9, Predicted=682.000000, Expected=68
10 Month=10, Predicted=475.300000, Expected=4
11 Month=11, Predicted=581.300000, Expected=5
12 Month=12, Predicted=646.900000, Expected=6
13 Test RMSE: 0.000
```

Develop a Robust Result

A difficulty with neural networks is that they give different results each time they are run.

One approach might be to fix the random number seed so the results are always reproducible. Another approach would be to control for the random initial conditions using a different experimental setup.

We can repeat the experiment from the previous section multiple times, then take the average RMSE as an indication of how well the configuration would be expected to perform on unseen data on average.

This is often called multiple repeats or multiple restarts.

We can wrap the model fitting and walk-forward validation in a loop of fixed number of repeats. Each iteration the RMSE of the run can be recorded. We can then summarize the distribution of RMSE scores.

```
1 # repeat experiment
2 repeats = 30
3 error_scores = list()
4 for r in range(repeats):
5     # fit the model
6     lstm_model = fit_lstm(train_scaled, 1, 3000, 4)
7     # forecast the entire training dataset to build up state for forecasting
8     train_reshaped = train_scaled[:, 0].reshape(len(train_scaled), 1, 1)
9     lstm_model.predict(train_reshaped, batch_size=1)
10    # walk-forward validation on the test data
11    predictions = list()
12    for i in range(len(test_scaled)):
```

[Start Machine Learning](#)

```

13     # make one-step forecast
14     X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
15     yhat = forecast_lstm(lstm_model, 1, X)
16     # invert scaling
17     yhat = invert_scale(scaler, X, yhat)
18     # invert differencing
19     yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)
20     # store forecast
21     predictions.append(yhat)
22 # report performance
23 rmse = sqrt(mean_squared_error(raw_values[-12:], predictions))
24 print('%d) Test RMSE: %.3f' % (r+1, rmse))
25 error_scores.append(rmse)

```

The data preparation would be the same as before.

We will use 30 repeats as that is sufficient to provide a good estimate.

The complete example is listed below.

```

1 from pandas import DataFrame
2 from pandas import Series
3 from pandas import concat
4 from pandas import read_csv
5 from pandas import datetime
6 from sklearn.metrics import mean_squared_error
7 from sklearn.preprocessing import MinMaxScaler
8 from keras.models import Sequential
9 from keras.layers import Dense
10 from keras.layers import LSTM
11 from math import sqrt
12 from matplotlib import pyplot
13 import numpy
14
15 # date-time parsing function for loading the dataset
16 def parser(x):
17     return datetime.strptime('190'+x, '%Y-%m')
18
19 # frame a sequence as a supervised learning problem
20 def timeseries_to_supervised(data, lag=1):
21     df = DataFrame(data)
22     columns = [df.shift(i) for i in range(1, lag+1)]
23     columns.append(df)
24     df = concat(columns, axis=1)
25     df.fillna(0, inplace=True)
26     return df
27
28 # create a differenced series
29 def difference(dataset, interval=1):
30     diff = list()
31     for i in range(interval, len(dataset)):
32         value = dataset[i] - dataset[i - interval]
33         diff.append(value)
34     return Series(diff)
35
36 # invert differenced value
37 def inverse_difference(history, yhat, interval=1):
38     return yhat + history[-interval]
39
40 # scale train and test data to [-1, 1]
41 def scale(train, test):
42     # fit scaler
43     scaler = MinMaxScaler(feature_range=(-1, 1))
44     scaler = scaler.fit(train)
45     # transform train
46     train = train.reshape(train.shape[0],

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this free and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

```

47     train_scaled = scaler.transform(train)
48     # transform test
49     test = test.reshape(test.shape[0], test.shape[1])
50     test_scaled = scaler.transform(test)
51     return scaler, train_scaled, test_scaled
52
53 # inverse scaling for a forecasted value
54 def invert_scale(scaler, X, value):
55     new_row = [x for x in X] + [value]
56     array = numpy.array(new_row)
57     array = array.reshape(1, len(array))
58     inverted = scaler.inverse_transform(array)
59     return inverted[0, -1]
60
61 # fit an LSTM network to training data
62 def fit_lstm(train, batch_size, nb_epoch, neurons):
63     X, y = train[:, 0:-1], train[:, -1]
64     X = X.reshape(X.shape[0], 1, X.shape[1])
65     model = Sequential()
66     model.add(LSTM(neurons, batch_input_shape=(batch_size, 1)))
67     model.add(Dense(1))
68     model.compile(loss='mean_squared_error')
69     for i in range(nb_epoch):
70         model.fit(X, y, epochs=1, batch_size=batch_size)
71         model.reset_states()
72     return model
73
74 # make a one-step forecast
75 def forecast_lstm(model, batch_size, X):
76     X = X.reshape(1, 1, len(X))
77     yhat = model.predict(X, batch_size=batch_size)
78     return yhat[0, 0]
79
80 # load dataset
81 series = read_csv('shampoo-sales.csv', header=0)
82
83 # transform data to be stationary
84 raw_values = series.values
85 diff_values = difference(raw_values, 1)
86
87 # transform data to be supervised learning
88 supervised = timeseries_to_supervised(diff_values, 1)
89 supervised_values = supervised.values
90
91 # split data into train and test-sets
92 train, test = supervised_values[0:-12], supervised_values[-12:]
93
94 # transform the scale of the data
95 scaler, train_scaled, test_scaled = scale(train, test)
96
97 # repeat experiment
98 repeats = 30
99 error_scores = list()
100 for r in range(repeats):
101     # fit the model
102     lstm_model = fit_lstm(train_scaled, 1, 3000, 4)
103     # forecast the entire training dataset to build up state for forecasting
104     train_reshaped = train_scaled[:, 0].reshape(len(train_scaled), 1, 1)
105     lstm_model.predict(train_reshaped, batch_size=1)
106     # walk-forward validation on the test data
107     predictions = list()
108     for i in range(len(test_scaled)):
109         # make one-step forecast
110         X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
111         yhat = forecast_lstm(lstm_model, 1, X)
112         # invert scaling
113         yhat = invert_scale(scaler, X, yhat)
114         # invert differencing
115         yhat = inverse_difference(raw_val=yhat[-1], last_ob=yhat[-2], difference=1)
116     error_scores.append(mse(test_scaled, predictions))
117
118 print('Walk Forward MSE: %f' % (sum(error_scores) / float(repeats)))

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this free and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

```

116      # store forecast
117      predictions.append(yhat)
118  # report performance
119  rmse = sqrt(mean_squared_error(raw_values[-12:], predictions))
120  print('%d) Test RMSE: %.3f' % (r+1, rmse))
121  error_scores.append(rmse)
122
123 # summarize results
124 results = DataFrame()
125 results['rmse'] = error_scores
126 print(results.describe())
127 results.boxplot()
128 pyplot.show()

```

Running the example prints the RMSE score each repeat. The end of the run provides summary statistics of the collected RMSE scores.

Note: Your results may vary given the stochastic nature of the problem. Also differences in numerical precision. Consider running the example a few times and compare the results.

We can see that the mean and standard deviation are approximately 136.761 and 108.829 respectively.

This is a very useful result as it suggests the results of the experiment suggest that the model is probably about 136.761, if not slightly worse.

This indicates that, at the very least, further model

```

1 1) Test RMSE: 136.191
2 2) Test RMSE: 169.693
3 3) Test RMSE: 176.553
4 4) Test RMSE: 198.954
5 5) Test RMSE: 148.960
6 6) Test RMSE: 103.744
7 7) Test RMSE: 164.344
8 8) Test RMSE: 108.829
9 9) Test RMSE: 232.282
10 10) Test RMSE: 110.824
11 11) Test RMSE: 163.741
12 12) Test RMSE: 111.535
13 13) Test RMSE: 118.324
14 14) Test RMSE: 107.486
15 15) Test RMSE: 97.719
16 16) Test RMSE: 87.817
17 17) Test RMSE: 92.920
18 18) Test RMSE: 112.528
19 19) Test RMSE: 131.687
20 20) Test RMSE: 92.343
21 21) Test RMSE: 173.249
22 22) Test RMSE: 182.336
23 23) Test RMSE: 101.477
24 24) Test RMSE: 108.171
25 25) Test RMSE: 135.880
26 26) Test RMSE: 254.507
27 27) Test RMSE: 87.198
28 28) Test RMSE: 122.588
29 29) Test RMSE: 228.449
30 30) Test RMSE: 94.427
31          rmse

```

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this free and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

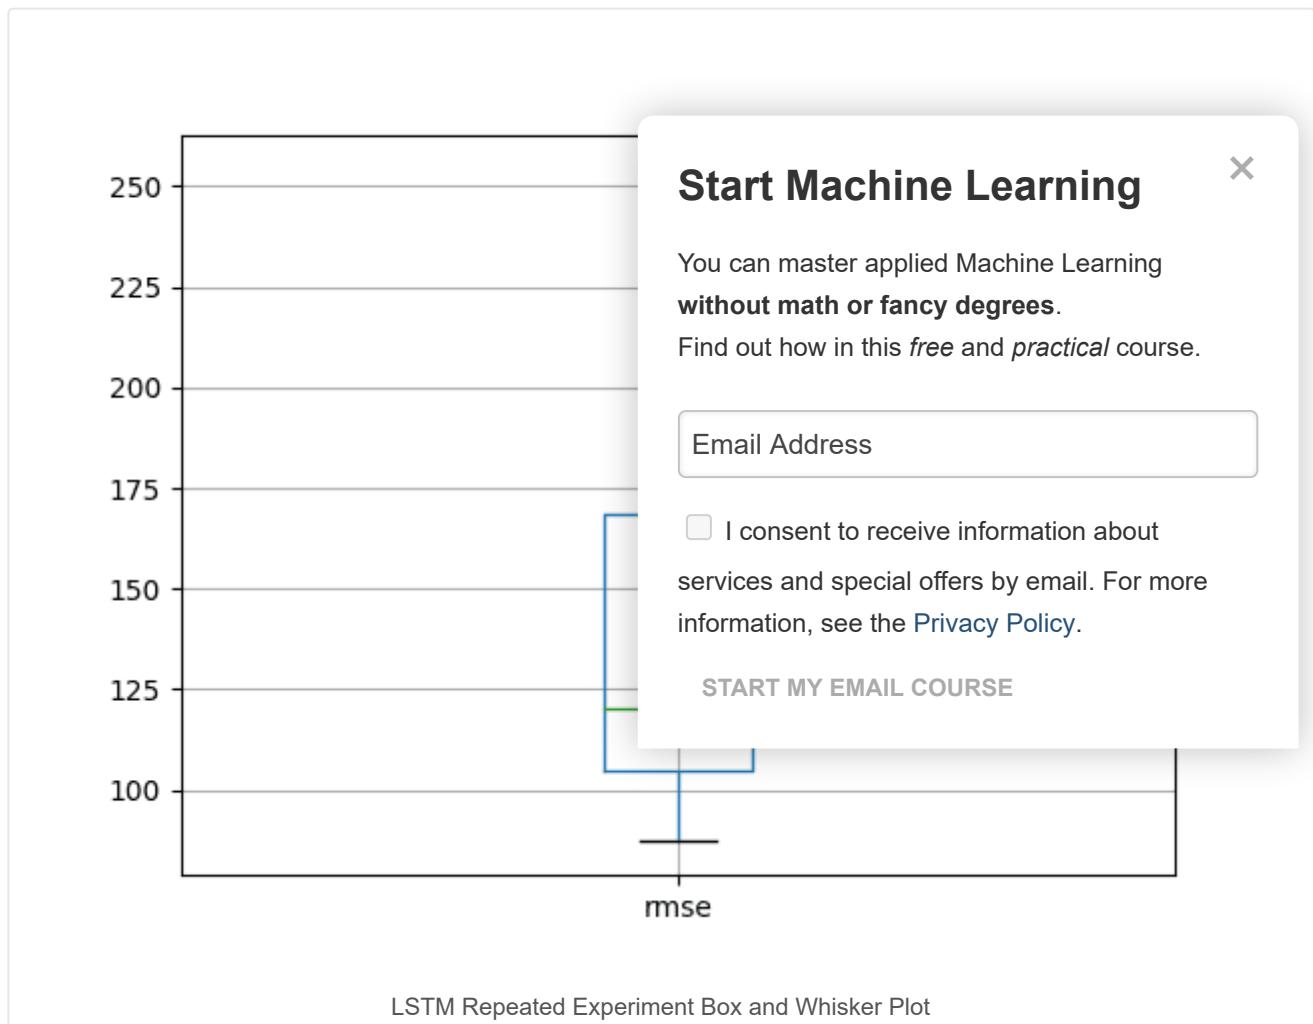
[START MY EMAIL COURSE](#)

```

32 count    30.000000
33 mean     138.491905
34 std      46.313783
35 min      87.198493
36 25%     104.679391
37 50%     120.456233
38 75%     168.356040
39 max      254.507272

```

A box and whisker plot is created from the distribution shown below. This captures the middle of the data as well as the extents and outlier results.



This is an experimental setup that could be used to compare one configuration of the LSTM model or set up to another.

Tutorial Extensions

There are many extensions to this tutorial that we may consider.

Perhaps you could explore some of these yourself and post your discoveries in the comments below.

- **Multi-Step Forecast.** The experimental setup could be changed to predict the next n -time steps rather than the next single time step. This would also permit a larger batch size and faster training. Note that we are basically performing a type of 12 one-step forecast in this tutorial given the model is not updated, although new observations are available and are used as input variables.

[Start Machine Learning](#)

- **Tune LSTM model.** The model was not tuned; instead, the configuration was found with some quick trial and error. I believe much better results could be achieved by tuning at least the number of neurons and number of training epochs. I also think early stopping via a callback might be useful during training.
- **Seed State Experiments.** It is not clear whether seeding the system prior to forecasting by predicting all of the training data is beneficial. It seems like a good idea in theory, but this needs to be demonstrated. Also, perhaps other methods of seeding the model prior to forecasting would be beneficial.
- **Update Model.** The model could be updated in each time step of the walk-forward validation. Experiments are needed to determine if it would be better to refit the model from scratch or update the weights with a few more training epochs including the new sample.
- **Input Time Steps.** The LSTM input supports multiple time steps. Experiments are needed to see if including lag observations as additional inputs provides any benefit.
- **Input Lag Features.** Lag observations may be included as additional inputs to the LSTM model to see if including lag features provide any benefit.
- **Input Error Series.** An error series may be computed and used as an additional input feature, not unlike a lag feature, to see if this provides any benefit.
- **Learn Non-Stationary.** The LSTM network may be able to learn non-stationary trends and make reasonable predictions. Experiments are needed to determine if learning non-stationary trends and seasonality, left in data can be learned by the LSTM.
- **Contrast Stateless.** Stateful LSTMs were used as the default in this tutorial, but experiments with stateless LSTM configurations.
- **Statistical Significance.** The multiple repeats of the experiment and evaluation process include statistical significance tests to demonstrate that the differences in the RMSE results with different configurations are statistically significant.

Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this free and practical course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Summary

In this tutorial, you discovered how to develop an LSTM model for time series forecasting.

Specifically, you learned:

- How to prepare time series data for developing an LSTM model.
- How to develop an LSTM model for time series forecasting.
- How to evaluate an LSTM model using a robust test harness.

Can you get a better result?

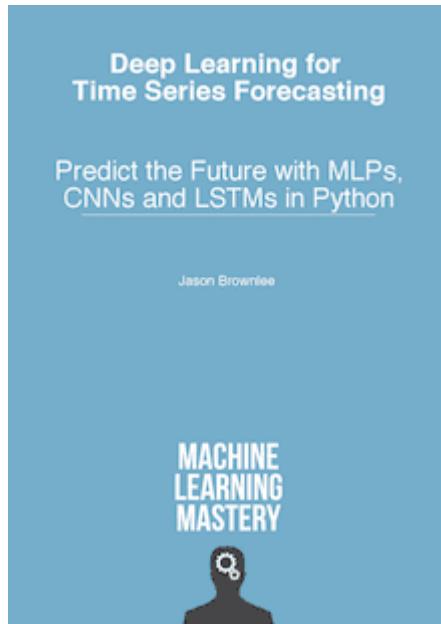
Share your findings in the comments below.

Develop Deep Learning models for Time Series Today!

Develop Your Own Forecasting models in Minutes

...with just a few lines of python code

[Start Machine Learning](#)



Discover how in my new Ebook:
Deep Learning for Time Series Forecasting

It provides **self-study tutorials** on topics like:
CNNs, LSTMs, Multivariate Forecasting, Multi-Step Forecasting and much more...

Finally Bring Deep Learning to your Time Series Forecasting Projects

Skip the Academics. Just Results.

SEE WHAT'S INSIDE

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

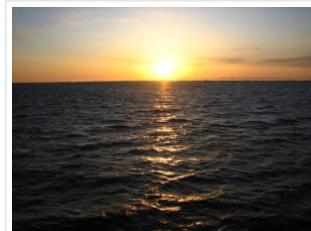
Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

More On This Topic



How to Get Started with Deep Learning for Time...



How to Develop LSTM Models for Time Series Forecasting

Start Machine Learning



Multi-Step LSTM Time Series Forecasting Models for...



Comparing Classical and...

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

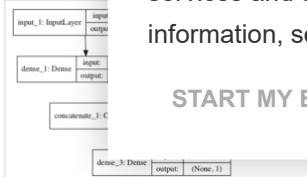
Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

How to Develop Convolu...



How to Develop Multilayer Perceptron Models for Time...



About Jason Brownlee

Jason Brownlee, PhD is a machine learning specialist who teaches developers how to get results with modern machine learning methods via hands-on tutorials.

[View all posts by Jason Brownlee →](#)

◀ Seasonal Persistence Forecasting With Python

How to Seed State for LSTMs for Time Series Forecasting in Python >

699 Responses to *Time Series Forecasting with the Long Short-Term Memory Network in Python*



Chang April 7, 2017 at 4:42 pm #

REPLY ↗

[Start Machine Learning](#)

I've been working on multi-step-ahead forecast after reading your ebook and following one step ahead tutorials. But still struggling getting nothing. It seems that seq2seq model is used, but I want to configure simple lstm for multi step ahead prediction. Can you help me in getting basic idea to do this?



Jason Brownlee April 9, 2017 at 2:53 pm #

REPLY ↗

Yes, I have some seq2seq examples scheduled to come out on the blog soon.



Shawn December 29, 2021 at 10:20 am

This is all fine and well but how do we use this model to simply look ahead x number of steps? Every blog I see just shows complete dataset? Every blog I see just shows complete



James Carmichael December 29, 2021 at 10:20 am #

Hi Shawn,

Perhaps the following may help clear up your question:

<https://www.youtube.com/watch?v=tep>



Supriya April 9, 2017 at 9:49 am #

REPLY ↗

Hi Jason,

Thank you for this blog. I am working on multi-step-ahead forecast using recursive prediction technique and I have some difficulty. Blog on this particular topic would be really helpful.
Also, is it possible to somehow implement recursive technique in ARIMA?.



Jason Brownlee April 9, 2017 at 3:01 pm #

REPLY ↗

Absolutely, you can take the predictions as history and re-fit the ARIMA.

Thanks for the suggestion, it would make a good blog post.



Jack Dan July 28, 2017 at 5:13 am #

REPLY ↗

I am wondering if I have predictions (not good prediction) and refitting without validating with the test set, wouldn't it give false prediction again?

[Start Machine Learning](#)

**Jason Brownlee** July 28, 2017 at 8:34 am #

REPLY ↗

It may, design an experiment to test the approach.

**Shifu Jain** December 27, 2017 at 11:50 pm #

REPLY ↗

Hi Jason,

Thanks for this blog. Can you please share an example of multi-step-ahead forecast using recursive prediction technique using LSTM.

Thanks,
Shifu

**Jason Brownlee** December 28, 2017 at 12:01 pm #

Thanks for the suggestion.

I have an example of a direct model for
<https://machinelearningmastery.com/memory-networks-python/>

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

**Peter Marelas** April 9, 2017 at 2:45 pm #

Hi Jason,

LSTM remember sequences but is there a way to encode calendar effects in the network so that it remembers or learns events that occur at different intervals within the sequence and each cycles? A concrete example would be a time series that exhibits specific events that repeat themselves on specific times in the year, example first Monday of every month and/or last day of every month. I am thinking whether we can label this data in advance to help the LSTM predict these events better?

**Guillaume** July 10, 2017 at 8:12 pm #

REPLY ↗

Hello Peter,

You might want to check out the X-11 method to separate trend, seasonal, and random change to your sequence. Then apply an algorithm to each part.

You can look at the following article :

Study of the Long-term Performance Prediction Methods Using the Spacecraft Telemetry Data from Hongzeng Fang

(Sorry but I can't find a free dl page anymore ..).

Start Machine Learning



Jason Brownlee July 11, 2017 at 10:29 am #

REPLY ↗

Thanks for the tip.



Yasmine Sayed May 18, 2018 at 11:58 pm #

REPLY ↗

Hey Jason, in the case of taking the last observation in the dataset in order to predict the next future timestep (for which we don't have current values for), does the model use only that 1 observation to predict the next??? Or because it's an LSTM , does it use both the memory from all the preceding observations during training + the last observation in order to predict the next timestep?

If it's the latter – it would make more sense to use all the previous observations, nothing but 1 observation to predict on



Jason Brownlee May 19,

The way the model is defined, it takes the last observation from the previous time step.

You can define the mapping any way you like.



Kunpeng Zhang April 10, 2017 at 5:04 am #

Hi Jason,

Nice post. Is there any tutorial available on multivariate time series forecasting problem? I got two sets of data: traffic flow data and weather data. I am thinking to predict the traffic flow using these two data sets.

I'd like to learn if I get weather condition involved what will happen for my model.

Could you kindly give me some advice?

Thank you.



Jason Brownlee April 10, 2017 at 7:40 am #

REPLY ↗

I do not have a multivariate regression example, but I hope to post one soon.



Kunpeng Zhang April 12, 2017 at 10:42 am #

REPLY ↗

Great. Thank you in advance.

Start Machine Learning



galdo trouchy April 12, 2017 at 7:46 pm #

REPLY ↗

That would be awesome, indeed



Wenhan Wang September 29, 2017 at 10:08 pm #

REPLY ↗

Hi Jason. It is really a nice example of using LSTM. I'm working on it.

I also have the same question, are you going to give an example on multivariate time series forecasting? I'm struggling on it.



Jason Brownlee September 30, 2017 at 10:08 pm #

I have an example here:

<https://machinelearningmastery.com/multivariate-time-series-forecasting-lstm-python/>



Gabriel Beauplet April 13, 2017 at 12:21 am #

You did a great job. This is very detailed !



Jason Brownlee April 13, 2017 at 10:02 am #

REPLY ↗

I'm glad you found it useful Gabriel.



Hand April 19, 2017 at 1:39 pm #

REPLY ↗

Great tutorial,

How do we get a prediction on a currently not existing point in the future?



Jason Brownlee April 20, 2017 at 9:22 am #

REPLY ↗

`yhat = model.predict(X)`



Hans April 20, 2017 at 1:44 am #

REPLY ↗

I mean without a y.

[Start Machine Learning](#)

Error measurements are cool, but I also want to make a prediction for a next step.

That's my main intention to read this tutorial.

Since there are other tutorials out there, lacking the same issue, it would be great to have a complete example, with a real live output/result/value.

I'm aware of the fact that every problem should have its own solution/model, but with a real result (or a comprehensible way to it), the code would be more practical/reusable- especially for python-ml beginners trying to predict some values in the future.



Jason Brownlee April 20, 2017 at 9:30 am #

REPLY ↗

Fit your model on the entire training dataset:

```
1 yhat = model.predict(X)
```



Donna May 22, 2017 at 7:07 am #

Hi Jason, thanks for your tutorial.
train your model on the entire training set and do the prediction for future



Jason Brownlee May 22, 2017

That is just a little too much hand holding Donna. What are you having trouble with exactly?



tuotuo October 16, 2018 at 8:57 pm #

REPLY ↗

Hi, Jsson:

yhat = model.predict(X):X means current values? and yhad means predict value?



Jason Brownlee October 17, 2018 at 6:49 am #

REPLY ↗

Yes, X is the input required to make one or more predictions and yhat are the predictions.



Ran Xu July 2, 2019 at 4:59 pm #

Hi, Jason, just got one question:
you own did prediction one test data

Start Machine Learning

the difference. What about the prediction for future unknown data, how can we do the difference inversion ?



Jason Brownlee July 3, 2019 at 8:21 am #

Yes, inverting the difference for a prediction only requires the last known observation.



Hans April 20, 2017 at 2:14 pm #

Thank you Jason,

I'm new to Python (mainly coding other languages) your outstanding detailed descriptions.

In the last weeks I tried out two other tutorials and appliance with a result).

A) Could you please suggest a location/row number

B) Is there a magic trick available to avoid the date sets?

In the moment I would have to transform my data to

I'm afraid to break the code logic.

But I also try to research every crucial part of the course example for example never
<https://docs.scipy.org/doc/numpy/reference/arrays.indexing.html>

My current Python skills are on the level of this tutorial

<https://www.youtube.com/watch?v=N4mEzFDjqtA> .

Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)



Jason Brownlee April 21, 2017 at 8:30 am #

REPLY ↗

Hi Hans,

What do you mean by "first own test-appliance with a result". I don't follow.

Do you mean make a prediction? If so you can make a prediction by fitting the model no all of your data and calling `model.predict(X)`.

Pandas is really flexible when it comes to loading date data. A good approach for you might be to specify your own date parsing function to use when loading your data. See this post for an example that does this:

<http://machinelearningmastery.com/stateful-stateless-lstm-time-series-forecasting-python/>



Hans April 20, 2017 at 3:39 pm #

[Start Machine Learning](#)

Worth to mention (latest Windows-Info 4.2017):

There is an issue with Keras/Anaconda on Windows.

To run the last example above on Win, we have to manually reinstall Keras.

Further informations can be found here:

<https://github.com/lISourcell/How-to-Predict-Stock-Prices-Easily-Demo/issues/3#issuecomment-288981625>

...otherwise it throws compiling errors.



Jason Brownlee April 21, 2017 at 8:30 am #

REPLY ↗

Does this tutorial work for you?

<http://machinelearningmastery.com/setup-python-environment-for-machine-learning-with-anaconda/>



Hans April 21, 2017 at 11:12 am #

I already saw this tutorial written by

On Windows (not a Mac) it's a slightly different setup. I used Anaconda.

It begins with the fact that there is no Tensorflow installed. I used the Keras-hint.

We needed several days to discuss it on GitHub and find the right setup in the context of another RNN-script (see the link).

I use virtual conda-environments and your scripts are running on windows if the keras-hint is implemented.

Before I had the same issue with some of your scripts then with the discussed one on Github (compile error).

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



Hans April 20, 2017 at 5:02 pm #

REPLY ↗

Update, what I mean:

I guess one could trigger "forecast_lstm(model, batch_size, X)" or "yhat = model.predict(X)" from the end of the last two example scripts.

But how to do that in regard to the trained model?

"Month=13, Predicted=????"

Do I have to define a new "fictional" X? And if so how?

Jason Brownlee April 21, 2017 at 8:32 am #

Start Machine Learning



You must load the new input data for which a prediction is required as X and use your already fit model to make the prediction by calling the predict() function.

I'm eager to help, but perhaps I don't understand the difficulty exactly?



Christian C. Russo June 17, 2018 at 2:10 pm #

REPLY ↗

Hi again Jason! brilliant tutorial once again,

I believe many people is asking about the model.predict() because it's not really working as expected.

First problem is that doing:

```
yhat = model.predict(X)
```

with the code example previously given, results in:

NameError: name 'model' is not defined

As I understand, this is because the model variable is not defined, so using:

```
yhat = lstm_model.predict(X)
```

works, but returns:

ValueError: Error when checking : expected 2D array, got 1D array instead
with shape (1, 1)

So, personally, what I have done, is using the forecast() function:

```
yhat = forecast_lstm(lstm_model, 1, X)
```

```
print(yhat)
```

0.28453988

Which actually returns a value.

Now the next problem is that the X, is nothing else than the last X of the example, as I never redefine it.

I found that the amount of functions and filters applied to the training data is quite big, hence I need to replicate them to make the shape match.

This is my original training data.

```
series = read_sql_query(seleccion, conn, parse_dates=['creacion'], index_col=['creacion'])

print(series)
sys.exit()

menores
creacion
2018-06-17 03:56:11 0.0
2018-06-17 03:54:03 2.0
2018-06-17 03:52:11 4.0
2018-06-17 03:50:05 6.0
```

Start Machine Learning



You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

```

2018-06-17 03:48:17 4.0
2018-06-17 03:46:04 4.0
2018-06-17 03:44:01 4.0
2018-06-17 03:43:05 1.0
2018-06-17 03:40:12 2.0
2018-06-17 03:38:12 0.0
2018-06-17 03:36:21 4.0
2018-06-17 03:34:32 4.0
2018-06-17 03:32:05 3.0
2018-06-17 03:30:01 2.0
2018-06-17 03:28:23 1.0
2018-06-17 03:26:17 3.0
2018-06-17 03:24:04 0.0
2018-06-17 03:22:34 4.0
2018-06-17 03:20:04 2.0
2018-06-17 03:18:18 2.0
2018-06-17 03:16:00 3.0
2018-06-17 03:14:06 6.0
2018-06-17 03:12:06 4.0
2018-06-17 03:10:04 2.0
2018-06-17 03:08:02 0.0
2018-06-17 03:06:02 4.0
2018-06-17 03:04:02 4.0
2018-06-17 03:02:10 3.0
2018-06-17 03:00:22 4.0
2018-06-17 02:59:13 3.0
...

```

[7161 rows x 1 columns]

Then, this process is applied to “series”:

```

# transform data to be stationary
raw_values = series.values
diff_values = difference(raw_values, 1)

# transform data to be supervised learning
supervised = timeseries_to_supervised(diff_values, 1)
supervised_values = supervised.values

```

if you print “supervised_values” at that point, the original data has been transformed to:

```

[[0 array([4.])]
[array([4.]) array([-3.])]
[array([-3.]) array([1.])]
...
[array([2.]) array([4.])]
[array([4.]) array([-3.])]
[array([-3.]) array([-1.])]]

```

which is clearly less and more condensed information...

Therefore, if I try to apply

```
yhat = forecast_lstm(lstm_model, 1, X)
```

Start Machine Learning



You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

after the new data has been loaded:

```
predecir = read_sql_query(seleccion, conn, parse_dates=['creacion'], index_col=['creacion'])

#limpiamos la cache
conn.commit()

print(predecir)
print(X)

yhat = forecast_lstm(lstm_model, 1, X)
#ynew = ynew[0]

print(yhat)
```

I get the following error:

AttributeError: 'DataFrame' object has no attribute 'reshape'

So kinda lost in how to actually apply the series to make the new prediction!

I'll paste the source my code in case someone needs it:
<https://codepen.io/anon/pen/xzPVxE>

You'll see that I'm loading the data directly from the database with a different approach from a previous post.

I'm not sure either about how to make the prediction.

Thanks a lot for this blog once again... I would appreciate explanations!!!

My best wishes,
 Chris

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.
 Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



Jason Brownlee June 18, 2018 at 6:38 am #

REPLY ↗

Thanks for sharing.



Guglielmo Iozzia May 31, 2019 at 2:23 am #

REPLY ↗

I have solved this the same way as suggested in this Stack Overflow post:
<https://stackoverflow.com/questions/42240376/dataframe-object-has-no-attribute-reshape>
 I am no more directly invoking the reshape method in the pandas training and test DataFrames, but indirectly through their values:
`train = train.values.reshape(train.shape[0], train.shape[1])`
`test = test.values.reshape(test.shape[0], test.shape[1])`

Start Machine Learning



Jason Brownlee May 31, 2019 at 7:51 am #

I'm happy to hear that.



Hans April 21, 2017 at 10:31 am #

REPLY ↗

Hello Jason,

I knew I shouldn't mention the date conversion part :-). Meanwhile I managed it with "return
datetime.strptime(x, "%Y-%m-%d")"

I have all of your example script parts in separate python files to fit my specific requirements.

```
python basic_data_loading.py
python persistence_forecast_model.py
python transform_time_series_to_supervised.py
python transform_time_series_to_stationary.py
python transform_scales.py
python complete_example.py
```

Own data loading is solved. That was a relatively easy part.



Hans April 21, 2017 at 10:56 am #

Since the transforming and performance measurement parts are running (I guess they will do even with integer data)

I now have to build a part lets call it:

"python predict_one_value.py"

Of course I have to load my own data that's clear.

The question is where to trigger the function

```
yhat = model.predict(X)
```

in the context of one of your example-scripts and finally say:
print(yhat). That's all.

I guess a short example snippet could solve the problem.

Could you provide an example- It would help a lot?

Currently I also don't understand the role of X completely.

In the context of "datetime.strptime" it seems to be a date only, If I print it out.

So if I would have training data of:

- 01.12.1977
- 02.12.1977
- 03.12.1977

Start Machine Learning



You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

I would guess I could say something like "yhat = model.predict("1977-12-04")".
The question is where and when in which code context.

Thank you.



Hans April 21, 2017 at 12:06 pm #

REPLY ↗

Update:

Currently I use the the code from "complete example" ("without robust result").
If I comment out from line 106 to line 127 and then at the end of the script say:

```
# report one value in the future
test = datetime.strptime('2017-04-10', '%Y-%m-%d')
yhat = model.predict(test)
print(yhat)
```

I get the error message 'model is not defined'

```
# report one value in the future
test = datetime.strptime('2017-04-10', '%Y-%m-%d')
yhat = lstm_model.predict(test)
print(yhat)
```

...throws the error "Data should be a Numpy array"

I guess maybe I could also append a new column
but I'm not sure If this would be right.

The best way to get this running would be an example snipped in context.

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about
services and special offers by email. For more
information, see the [Privacy Policy](#).

START MY EMAIL COURSE



Jason Brownlee April 22, 2017 at 9:22 am #

REPLY ↗

This post will give you a clearer idea of how to use neural nets in Keras including
making predictions:

<http://machinelearningmastery.com/5-step-life-cycle-neural-network-models-keras/>

Yes, input data to making a prediction must be a 2D numpy array. It will not be date data, it
will be past obs in the case of time series forecasting.



Hans April 22, 2017 at 11:41 am #

Thank you,

I will read it ;-).

Start Machine Learning

**Hans** April 22, 2017 at 12:46 pm #

I have read the tutorial and tried out the pima indians diabetes example.
I guess I got it and understand the 5 steps (mostly).

Unfortunately this does not answer my question. Or do I miss something?
In my problem I have only one input like in the tutorial on this site.

When you say:

`"yhat = model.predict(X)"`

would give a forecast for a next step.

What about a step which is not in the training data nor in the test data?

I have a SVM model which proves
another environment).

Lets say I have 100 items training

It will printout 10 predictions and ac

The last printed prediction is for a 1

How would this be archived in your

Do I have to shift something?

**Jason Brownlee** April 23,

To make predictions beyo
observations from your dataset as

This post might clear up your thinking on X and y:

<http://machinelearningmastery.com/time-series-forecasting-supervised-learning/>

Start Machine Learning X

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about
services and special offers by email. For more
information, see the [Privacy Policy](#).

START MY EMAIL COURSE

**Hans** April 23, 2017 at 1:27 pm #

>To make predictions beyond your dataset, you must feed in the last few
observations from your dataset as input (X) to >predict what happens next (y).

Is there an example available where this is done?

**Jason Brownlee** April 24, 2017 at 5:32 am #

Yes, the “LSTM Forecast” section of this very post.

**Hans** April 24, 2017 at 1:55 pm #

Assuming that “`X = row[0:-1]`”
how do we sample/collect the last

Start Machine Learning



Jason Brownlee April 25, 2017 at 7:43 am #

It depends on your model, if your model expects the last one observation as input, then reframe that value as a 2d array and provide it as X to model.predict(X).

If your model requires the last two lag obs as inputs, retrieve them, define them as a one row, two column matrix and provide them to model.predict().

And so on. I hope that helps.



Arun July 13, 2017 at 10:37 am #

HI Jason,

Even i started with machine learning forecasting we can only get one time prediction i have provided last input model.predict(X) has to be again s

PFB the code:

```
X = test_scaled[3:-1:] (my last observation)
yhat = forecast_lstm(lstm_model, X)
yhat = invert_scale(scaler, X, yhat)
yhat = inverse_difference(raw_values, yhat[-1])
print(yhat)
```

Can you please guide me if i am going in a right way ?

Start Machine Learning

X

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



Jason Brownlee July 13, 2017 at 4:53 pm #

Yes, one step forecasting involves predicting the next time step.

You can make multi-step forecasts, learn more in this post:

<http://machinelearningmastery.com/multi-step-time-series-forecasting/>

Yes, to make use of the prediction you will need to invert any data transforms performed such as scaling and differencing.



Arun Menon July 13, 2017 at 11:32 pm #

Thank you Jason ..



Arun Menon July 13, 2017

Start Machine Learning

Hi Jason,

Being said that, i have another clarification , so when i forecast the next time step using this model , using the below code:

```
X = test_scaled[3,-1:] (my last observation)
yhat = forecast_lstm(lstm_model, 1, X)
yhat = invert_scale(scaler, X, yhat)
yhat = inverse_difference(raw_values, yhat, 1)
print(yhat)
```

in the above code, let yhat be the prediction of future time step, can i use the result of yhat and use the same model to predict one more step ahead in the future ? is this what we call as the recursive multistep forecast ?



Jason Brownlee July 14,

Yes. This is recursive.



Arun Menon July 14, 2017

Hi Jason,

Can i use the below code and use for eg :

yhat value can be used as an input and so on ?

```
X = test_scaled[3,-1:] (my last observation)
yhat = forecast_lstm(lstm_model, 1, X)
yhat = invert_scale(scaler, X, yhat)
yhat = inverse_difference(raw_values, yhat, 1)
print(yhat)
```

Start Machine Learning



You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



Arun Menon July 14, 2017 at 11:59 am #

Hi Jason,

In predictive analytics using this ML technique, how many future steps should we able to predict , is there any ideal forecasting range in future for eg if i have a data for the last 10 days or so , and i want to forecast the future , the less the future time steps are set, the better the result as the error will be minimum right. Can i use the same code for predicting time series data in production for network traffic for future 3 days ? requirement given for me was to predict the network bandwidth for the next entire week given the data for past 1 year.

Your comments and suggestions always welcome 😊

Start Machine Learning

Regards,
Arun



Jason Brownlee July 15, 2017 at 9:37 am #

It depends on the problem and the model.

Generally, the further in the future you want to predict, the worse the performance of the model.



Hans April 21, 2017 at 6:37 pm #

Hello,

can we normalize the RMSE-value(s)?
And if so how?



Jason Brownlee April 22, 2017 at 9:24 am #

Normalize the skill score?

Yes, but you will need to know the largest poss

Start Machine Learning

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



Hans April 22, 2017 at 12:55 pm #

REPLY ↗

I'm feeding your example with values ranged from 1 to 70.

There is no increasing trend in my raw data.

When it comes to predictions the script predicts values above 70.

Regarding to the other tutorial (5 Step Life-Cycle) I think it has to do with the compile part (model.compile).

But I'm not sure. Could you provide a comprehensible hint in regard of the example script on this site?



Fabian April 23, 2017 at 10:33 am #

REPLY ↗

Hi Jason,

assuming you had multiple features (you can add one feature to the shampoo dataset) and wanted to use multiple timesteps, what would the dataset look like that I put into the model? Is it a 3 dimensional array, where the features are lists of values and each observation is a list of these features and the label(s) (which is also a list of values)?

Start Machine Learning



Jason Brownlee April 24, 2017 at 5:31 am #

REPLY ↗

Good question, it would be a 3D array with the dimensions [samples, timesteps, features].



Rahul April 26, 2017 at 5:33 am #

REPLY ↗

Right now the model gives only one step forecast. What if I wanted to create a model which gives forecast for next 60 months.



Jason Brownlee April 26, 2017 at 6:24 am #

I will have a post on this see, until then

<http://machinelearningmastery.com/multi-step-forecasting-tutorial/>



Kunpeng Zhang April 27, 2017 at 1:07 am #

Hi Jason, I have a question.

```
raw_values = series.values
diff_values = difference(raw_values, 1)
print(len(raw_values)) 36
print(len(diff_values)) 35
```

So, after difference we lose the first value?

Start Machine Learning

X

You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



Jason Brownlee April 27, 2017 at 8:43 am #

REPLY ↗

Correct.



Guido van Steen April 27, 2017 at 2:27 am #

REPLY ↗

Dear Jason,

Thank you very much for this extremely useful and interesting post.

I may be missing something, but I think there is one omission: By differencing you loose the trend including its start and end level. Later on you try restore the trend again, but in your code it seems you fail to restore the end level. IMO the end level of the current observations should be added to all the predictions.

Thanks again!

Start Machine Learning



Jason Brownlee April 27, 2017 at 8:44 am #

REPLY ↗

What do you mean by the end level? Sorry I don't follow, perhaps you could restate your comment?



Raul April 27, 2017 at 3:32 am #

REPLY ↗

Great tutorial!

Quick question: On the scaling section of the tutorial you say that

"To make the experiment fair, the scaling coefficient is calculated from the training dataset and applied to scale the test dataset. This prevents the experiment with knowledge from the test dataset."

However, if the max of your sample is on the test dataset, scaling the test dataset will yield a number outside the [-1,1] range. How

Thanks!



Jason Brownlee April 27, 2017 at 8:46 am #

Correct.

One good approach is to estimate the expected range of values in the test set and use these to scale.

If even then you see values out of range, you can clamp them to the bounds 0/1.

Start Machine Learning



You can master applied Machine Learning
without math or fancy degrees.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



Guido van Steen April 27, 2017 at 6:20 am #

REPLY ↗

Dear Jason,

To be more precise: you should add the difference between the start level and the end level of the train set. This is because the current code effectively replicates the train set. By this I mean that it starts at the same level as the train set. However, it should start at the end level of the train set.

Kind regards,

Guido



Guido van Steen April 27, 2017 at 3:11 pm #

REPLY ↗

I will try to restate my comment:

Currently the predictions (of your test set) start at the same level as the observations (in your train set). Therefore, there is a shift between the last

Start Machine Learning

predicted value (of your test set). The size of this shift is equal to: start level of the observations minus end level of the observations (in your train set). You should correct for this shift by adding it to the predicted values.



Jason Brownlee April 28, 2017 at 7:35 am #

REPLY ↗

Isn't this made moot by making the data stationary?



Deepa April 28, 2017 at 4:59 am #

Hello,

1. can you please explain me the below 2 lines in code
- ```
model.add(LSTM(neurons, batch_input_shape=(batch_size, timesteps), return_sequences=True))
model.add(Dense(1))
```
2. I want to know the layer architecture. What are they?
  3. If I want to add one more hidden layer, how should I do?
  4. What could be the reason for test rmse is less than training rmse?



**Jason Brownlee** April 28, 2017 at 7:56 am #

The line defines the input layer and the first LSTM hidden layer with neurons number of memory units.

You can stack LSTMs, the first hidden layer must return sequences (return\_sequence=True) and you can add a second LSTM layer. I have an example here:

<http://machinelearningmastery.com/time-series-prediction-lstm-recurrent-neural-networks-python-keras/>

Better performance on the test data than the training data may be a sign of an unstable/underfit model.



**Fabio** May 2, 2017 at 9:43 pm #

REPLY ↗

Hallo Jason, thank you for this post! I bought the first version of your book and I have seen you have in the meantime deeper analysed this topic. Very good! 😊 Something have I yet not clear.

“Samples: These are independent observations from the domain, typically rows of data.”

“Time steps: These are separate time steps of a given variable for a given observation.”

I could understand the case where the time step parameter is 1 as in your book and in this example but I can't figure out why and how it could be greater than 1...

My hypothesis, sure wrong 😊

[Start Machine Learning](#)

Perhaps when a timestep is made of n observations one could give the value n to it...but then I would expect when one in the model writes (pag. 192):

"model.add(LSTM(4, input\_shape=(1, look\_back)))"

the LSTM would use (look\_back \* timesteps) rows for every step to predict the next row...

I cannot also understand why you say 'of a given variable'...a row is normally built by the values of many variables, isn't it?

Could you give me an example with timesteps > 1? Thank you!



**Jason Brownlee** May 3, 2017 at 7:36 am #

Hi Fabio,

The structure of the input to the LSTM is [samples, time steps, features].

If you have multiple observations at one time step, then:

Does that help?



**Fabio** May 7, 2017 at 10:28 pm #

Hello Jason,

unfortunately If don't have a concrete example. Your posts and your book are clear to me but they do not show an example how could be adapted the scenario described in this post in order to manage a timesteps>1?

Thank you very much!

PS. In the meantime I bought also your book on time series 😊

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)



**Jason Brownlee** May 8, 2017 at 7:44 am #

REPLY ↗

See this post on multi-step forecasting:

<http://machinelearningmastery.com/multi-step-time-series-forecasting/>



**Nitin** May 5, 2017 at 8:11 am #

REPLY ↗

Thanks Jason for the wonderful tutorial!

I am using your tutorial to apply LSTM network on some syslog/network log data.

I have syslog data(a specific event) for each day for last 1 year and so I am using LSTM network for time series analysis.

[Start Machine Learning](#)

As I understand from your tutorial.

1. A batch of data is a fixed-sized number of rows from the training dataset that defines how many patterns to process before updating the weights of the network. Based on the `batch_size` the Model takes random samples from the data for the analysis. For time series this is not desirable, hence the `batch_size` should always be 1.

2. By default, the samples within an epoch are shuffled prior to being exposed to the network. This is undesirable for the LSTM because we want the network to build up state as it learns across the sequence of observations. We can disable the shuffling of samples by setting "shuffle" to "False".

Scenario1 –

Using above two rules/guidelines – I ran several trials with different number of neurons, epoch size and different layers and got better results from the baseline model(persistence model).

Scenario2-

Without using above guidelines/rules – I ran several and different layers and got even better results than

Query – Setting shuffle to True and `Batch_size` value

It seems logical reading your tutorial that the data if want to change the sequence of data, but for my da

At the end what I think, what matters is how I get b

I think I should try and put away "theory" over conc

Kindly enlighten.



**Jason Brownlee** May 5, 2017 at 11:25 am

Random samples are not taken, LSTMs require sequence data to be ordered – they learn order dependence.

The "shuffle" argument must be set to "False" on sequence prediction problems if sequence data is spread across samples.



**Masum** May 7, 2017 at 12:44 am #

REPLY ↗

Dear Jason,

I have two hours time series data which consists of 120 observations, using LSTM RNN how can I predict next 30 observation while putting all my data to training section.

We normally split the original data set in two data set ( test dataset and validation dataset) for checking our model performance. I would like to see that my model is only taking help from training dataset to produce an output that does match with validation data set. What I understand from several of your blogs that we are forecasting single value and using that single forecasting along with the help of validation dataset we are forecasting rest of values. I believe I getting lost there? How it is going to work when have only past and current data ( suppose no validation data) and we want to predict the next half an hour.

For example, suppose I have data of a product price observations and using these past 90 observations

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

to 2:00pm(otherwise next 30 observations) .

Would you please help me to solve the confusion that I have? By the way I am going through your books time series forecasting with Python and deep learning with Python.



**Jason Brownlee** May 7, 2017 at 5:42 am #

REPLY ↲

You can make a multi-step model with an LSTM by using it directly (predict  $n$  timesteps in a one shot manner) or recursively (use the same model again and again and use predictions as inputs for subsequent predictions).

More on multi-step forecasting here:

<http://machinelearningmastery.com/multi-step-time-series-forecasting/>

Does that help?



**masum** May 7, 2017 at 5:57 am #

Thanks for quick replay

Looks like you sleep very less as you have provided

Thanks for being so kind to your students. May god

By the way do you have any blog or example for setting up a model again and again and use predictions as input to see that I am not using any data from validation data. Model should only from past and current data along with current predictions. I hope you got this confused student.

Thanking you



**Jason Brownlee** May 8, 2017 at 7:41 am #

REPLY ↲

Sorry, I don't have an example of using the model recursively. I believe you could adapt the above example.



**masum** May 9, 2017 at 9:10 am #

REPLY ↲

would you please give me some hint where I have to change the code to make the model recursive as I am not very good at coding.



- Jason Brownlee May 10, 2017 at 8:36 am #

REPLY ↵

# Start Machine Learning

Predictions would be used as history (input) to make predictions on the subsequent time steps.



**Max** May 9, 2017 at 11:10 am #

REPLY ↗

Thanks Jason for the wonderful tutorial!

A little problem here:

```
def invert_scale(scaler, X, value):
new_row = [x for x in X] + [yhat]
...

```

Does it should be

```
def invert_scale(scaler, X, value):
new_row = [x for x in X] + [value]
...

```



**Jason Brownlee** May 10, 2017 at 8:38 am #

Thanks Max.

## Start Machine Learning

X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Charlie** May 17, 2017 at 4:08 am #

REPLY ↗

I'm getting the following error when running the first block of code you provide.

TypeError: strptime() argument 1 must be str, not numpy.ndarray

It seems like this a problem with how Python 3.6 handles byte strings.



**Jason Brownlee** May 17, 2017 at 8:42 am #

REPLY ↗

I have only tested the code on Python 2.7 and Python 3.5, perhaps try one of those?

Otherwise, you may need to adapt the example for your specific platform, if an API change is indeed the cause of your issue.



**Sabih** November 8, 2017 at 6:47 am #

REPLY ↗

The issue was with the text in footer 'Sales of shampoo over a three year period'.

Either delete the footer OR

Start Machine Learning

Change the line:

```
series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True,
date_parser=parser)
```

to:

```
series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True,
date_parser=parser, skipfooter=2)
```



**John Strong** November 23, 2017 at 2:48 am #

REPLY ↗

I think that error is caused by the line  
data file. If you see this comment at the bottom

Sales of shampoo over a three year period

The data file also has the dates enclosed in do  
date\_parser callback is a little fickle.



**Logan** May 18, 2017 at 10:28 am #

Hello, Your examples are great, thanks.

I have a question:

1- What do we do with the array of predicted values  
in the example? It seems like they are not being used. Is it important for the net?



**Jason Brownlee** May 19, 2017 at 8:10 am #

REPLY ↗

Thanks Logan.

What do you mean exactly? Is there a specific section and line of code you are referring to in the tutorial?



**Logan** May 27, 2017 at 2:47 am #

REPLY ↗

Yes, line 101 `lstm_model.predict(train_reshaped, batch_size=1)` on section Complete  
LSTM Example.

By the way, I have another question: When the differentiation is applied to the entire dataset,  
isn't it giving information about the test dataset to the model?

Thanks. You've been helping me a lot.

Start Machine Learning



**Jason Brownlee** June 2, 2017 at 11:58 am #

REPLY ↗

In this case yes, but in practice, we can apply the same method to the test data using data from training step-wise without cheating.



**Long** June 13, 2018 at 10:14 pm #

Hi Jason, I'm also having the same question as Logan mentioned above. Line 101

`Istm_model.predict(train_reshaped)`  
do we need this line.

I comment this line, and I see no difference.  
this line ?

Thank you for the great tutor.

Best, Long

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Brandon** May 19, 2017 at 12:34 pm #

REPLY ↗

Hi Jason,

Thanks so much for this helpful tutorial. Could you clarify which of all the parameters need to be updated to change the lag value? I haven't been able to get it to work.

Thanks!



**Jason Brownlee** May 20, 2017 at 5:34 am #

REPLY ↗

Great question.

Change the parameters to the `timeseries_to_supervised()` function



**John Strong** February 10, 2018 at 9:28 am #

REPLY ↗

I had to change the limits of the `range()` function to get the loop to work for lag values > 1.

Start Machine Learning

```
def timeseries_to_supervised(data, lag=1):
df = DataFrame(data)
columns = [df.shift(i) for i in range(lag, 0, -1)]
columns.append(df)
df = concat(columns, axis=1)
df.fillna(0, inplace=True)
return df
```

**Jason Brownlee** February 11, 2018 at 7:49 am #

REPLY ↗

Nice!

**Nirikshith** May 29, 2017 at 8:23 pm #

Hi Jason,

I have been following you for quite sometime now. trying out sliding window approach in one of the examples involving external features(it can be as simple as implemented).

can you point out where we can add the above extra please for these features please. I will try it out. being reshapes and pre-format required for LSTM s

#aspring data scientist, student

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees.**

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)
**Jason Brownlee** June 2, 2017 at 12:25 pm #

REPLY ↗

You can add them as separate features.

I am working on an example at the moment. I hope to post it when it gives good results.

**Donato** June 1, 2017 at 10:00 pm #

REPLY ↗

Hi Jason, nice example! I have your similar problem, but I don't know how resolve. I have a two different class to predict, in different files. I have to training my neural network with this different files. The problem is that each file has different number of samples, and I have to use the same LSTM for each file. How I can generalize the number of input for each LSTM??  
This is my problem:  
<https://drive.google.com/file/d/0B5hOtU0Xa45RUDJJWHVyeHVNQWM/view?usp=sharing>

**Jason Brownlee** June 2, 2017 at 12:59 pm

## Start Machine Learning



Maybe this post will help you better understand your prediction problem and frame it as supervised learning:

<http://machinelearningmastery.com/how-to-define-your-machine-learning-problem/>



**Donato** June 5, 2017 at 8:24 pm #

REPLY ↗

Thanks for the answer, I have understand which is my problem. It depends from the training, because I have a lot of file for training, and when I insert a new file and remake the fit, the previously information are forgotten. Do you have some examples for train the LSTM with more file? And where each one has the supervisioned learning, apply to time-series and predict information? Thank you!



**Jason Brownlee** June 6, 2017

Yes, this example shows how  
<http://machinelearningmastery.com/update-lstm-time-series-forecasting-python/>



**Pengyuan Shao** June 8, 2017 at 10:02 pm #

Dear Jason:

You did a great job!!

Thank you for sharing the great post with us.

I'm a beginner to machine learning, these days i am busy on doing a project to predict the movement of ship on the sea.

I found this post is very usefull and easy to understand.

I have one question that:

if the number of neurons is the past steps used to predict next step?

Thank you for answer in advance.

Best wishes

## Start Machine Learning

X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** June 9, 2017 at 6:24 am #

REPLY ↗

I'm not sure what you mean.

The number of neurons defines the size/complexity of the model. The model is used to make predictions.

Perhaps you could reframe your question?



**Macarena** June 10, 2017 at 7:19 pm #

REPLY ↗

Start Machine Learning

Hi,

Congrats for the blog, and thanks! It helped me a lot.

I have one question related to the data transformation. I understand the step to make that one input is the last output, so the LSTM can learn and save that info. But I have a different case (classification) and I don't know how could I apply this.

In my case my inputs are images (frames from a video) what have a temporal dependency, and my outputs are numbers (0,1 or 2). First I want to pass the images through a CNN network and then take the features and those would be my LSTM input. In that case, it is necessary to process the data as you have done here? I thought the LSTM would save that info by itself, but now I'm not so sure.

Thank you in advance.



**Jason Brownlee** June 11, 2017 at 8:23 am #

I don't have an example of working wi

Generally, the principles would be the same. A  
the output layer to use one of the log loss activi  
similar, e.g. [cnn]->[lstm]->[dense].

Let me know how you go.



**Phil** June 22, 2021 at 1:51 pm #

Hey did you ever get this working I ha  
value your input!!



**Kim Miller** June 15, 2017 at 6:14 am #

REPLY ↗

In your ARIMA tutorial you also use the shampoo data set, but measure accuracy by MSE.  
Making the change to RMSE:

```
rmse = sqrt(mean_squared_error(test, predictions))
print('Test RMSE: %.3f' % rmse)
```

I get:

> Test RMSE: 83.417

Can we conclude that ARIMA is far better than an LSTM for this problem so far in our research?



**Jason Brownlee** June 15, 2017 at 8:55 am #

REPLY ↗

[Start Machine Learning](#)

No, these are demonstrations. Both algorithms would require tuning to present the best version of themselves (“steel man” vs “straw man”) for model selection.



**Kim Miller** June 15, 2017 at 12:28 pm #

REPLY ↗

We are writing the timeseries\_to\_supervised function to accept a lag:

```
def timeseries_to_supervised(data, lag=1):
```

Is this feature to essentially the Window Method + reshaping described in your July 21, 2016 tutorial?

But we don't seem to use that parameter, always “t”

Using lag values up to around 50 (on the airline da

```
forecast the entire training dataset to
train_reshaped = train_scaled[:, 0:look_back]
train_reshaped = train_scaled[:, 0].res
lstm_model.predict(train_reshaped, batch_
```

seems to make good predictions with very low RMSE

Am I possibly breaking the model now so it's chea

## Start Machine Learning

X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

REPLY ↗



**Kim Miller** June 16, 2017 at 5:27 am #

```
look_back = 60
epochs = 15
```

Month=1, Predicted=321.813110, Expected=315.000000

Month=2, Predicted=310.719757, Expected=301.000000

Month=3, Predicted=363.746643, Expected=356.000000

...

Month=46, Predicted=458.424683, Expected=461.000000

Month=47, Predicted=418.427124, Expected=390.000000

Month=48, Predicted=412.831085, Expected=432.000000

Test RMSE: 18.078

for test harness

```
yhat = y
Test RMSE: 0.000
```

It seems there is contamination of training data in the test forecast in the sense that forecast\_lstm(model, batch\_size, X) is being given X observations, each with 60 historical past observations that overlap into training data. Bu

Start Machine Learning

as it moves into the test data it adds history only seen in the test set. But that's how real life observations work it seems: You always have all of the past to look at.

Finally, you say, "In this tutorial, we will go with the fixed approach for its simplicity, although, we would expect the dynamic approach to result in better model skill." For a small dataset like we have, a large lag seems as if it gives us a semi-dynamic approach in some ways?

Solid model with a seemingly good RMSE?

Full code: <http://tinyurl.com/y74ypvde>

Prediction v. Actual Plot: <http://tinyurl.com/y9ouajdm>



**Kim Miller** June 16, 2017 at 7:26 am #

On the Shampoo dataset:

look\_back = 18 (lag)

epochs = 30

Test RMSE: 110.594

Prediction v. Actual Plot: <http://tinyurl.com/y9ouajdm>



**Jason Brownlee** June 16, 2017 at 8:01 am #

Very sharp!



**Jason Brownlee** June 16, 2017 at 8:10 am #

REPLY ↗

Nice.

Hmmm, I believe by dynamic I was referring to updating the model with new data.

If you have the resources, I would recommend exploring whether updating the model as new obs come in effects model skill. I expect it will. I have an example here:

<http://machinelearningmastery.com/update-lstm-networks-training-time-series-forecasting/>



**Pedro** June 15, 2017 at 5:22 pm #

REPLY ↗

Hi Jason,

You're doing a great job helping people to learn how to apply machine learning. One week ago, I didn't know anything about time series, but now I'm able to play a little bit with it. Thanks you!

Regarding this post, I've a doubt. When we say

'To make the experiment fair, the scaling coefficients (min and max) values must be calculated on the training dataset and applied to scale the test dataset and any forecasts. This is to avoid contaminating the experiment with knowledge from the test dataset. which might give the model a small edge.'

shouldn't we do

Start Machine Learning

```
scaler = scaler.fit(X[:-12])
```

instead of

```
scaler = scaler.fit(X) ?
```

If we use X we are using data from the test set, no?

Thanks for your attention and keep up with the good work!



**Jason Brownlee** June 16, 2017 at 7:51 am #

REPLY ↗

Correct. Ideally, we would use domain values that could ever be observed.



**Eric** June 20, 2017 at 6:57 pm #

Hello Jason,

After using the robust code (last one) I got theses results:

- 1) Test RMSE: 180.438
- 2) Test RMSE: 110.352
- 3) Test RMSE: 119.655
- 4) Test RMSE: 170.720
- 5) Test RMSE: 211.877
- 6) Test RMSE: 101.453
- 7) Test RMSE: 105.532
- 8) Test RMSE: 149.351
- 9) Test RMSE: 88.118
- 10) Test RMSE: 138.013
- 11) Test RMSE: 265.045
- 12) Test RMSE: 135.861
- 13) Test RMSE: 167.766 ... (rest is omitted, it was taking too long).

Being a beginner in machine learning and using your tutorial to learn about it could you tell me if these results are normal (being so different from what you have) while everything before was ok. If these are not, could you tell me what could be the reason of such difference?

Thank you for your attention and please do continue making tutorial like that, it's really helpful!

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** June 21, 2017 at 8:11 am #

REPLY ↗

Hi Eric, see this post to better understand the stochastic nature of neural network algorithms:

<http://machinelearningmastery.com/randomness-in-machine-learning/>

See this post for more details on how to develop a robust estimate of model skill:

<http://machinelearningmastery.com/evaluate-sklearn-model/>

Start Machine Learning

See this post on how to fix the random number seed:

<http://machinelearningmastery.com/reproducible-results-neural-networks-keras/>



**Eric** June 21, 2017 at 1:02 pm #

REPLY ↗

I see, thank you very much for the links and solutions!



**Jason Brownlee** June 22, 2017 at 6:03 am #

REPLY ↗

You're welcome.



**Eric** June 23, 2017 at 5:50 pm #

After reading the links what I can say:  
 – It's pretty hard (impossible) to always get the same result.  
 That can come from a lot of things.  
 – Only by repeating the operation a few times (but doing it like 100 would give a better result).  
 The result can be different and find a range.  
 – Seeding the random number generator.

So to continue my project I thought about what I explained. Like here, repeat the process 30 time (since I have a server I can let it do it overnight and go for 100) and determine the average and decide of a range.

Does it seem ok to begin with?

Am I lacking something?

Thank you for your attention and for the help you provide to everyone.

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** June 24, 2017 at 8:00 am #

See this post on how to estimate the number of repeats:  
<http://machinelearningmastery.com/estimate-number-experiment-repeats-stochastic-machine-learning-algorithms/>

I would not recommend fixing the random number seed for experimental work. Your results will be misleading.



**Dhineshkumar** June 30, 2017 at 7:18 pm #

REPLY ↗

Hi Jason,

Start Machine Learning

Thank you so much for the tutorial.

I have few doubts though.

1. I am working on a problem where the autocorrelation plot of the detrended data show me that the value at time t is significantly correlated to around past 100 values in the series. Is it ideal to take the batch size of 100 to model the series?
2. You mentioned that less than 5 memory units is sufficient for this example. Can you please give me some idea on how to choose the number of memory units for a particular problem like the above? On what other factors does this number depend on?

Kindly clarify.

Thanks



**Jason Brownlee** July 1, 2017 at 6:33 am #

Try passing in 100-200 time steps and

Systematically test a suite of different memory



**Dhineshkumar** July 1, 2017 at 2:42 pm #

Thank you Jason.



**Jason Brownlee** July 2, 2017 at 6:26 am #

You're welcome.



## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**

REPLY ↗



**Sophie** July 4, 2017 at 3:04 pm #

REPLY ↗

Hey Jason! Thanks a lot for this tutorial, it really helped.

I used this tutorial as is for predicting the cost of an item which is of the the range of a dollar and few cents.

My dataset has 262 rows, i.e 0 to 261.

When I run the model, the graph captures even the most intricate trends beautifully, BUT there seems to be a lag of 1 time step in the predicted data.

The predicted values of this month almost exactly matches the expected value of the previous month. And this trend is followed throughout.

The indexing is the only thing I've changed,  
to

**Start Machine Learning**

```
train, test = supervised_values[0:200], supervised_values[200:]
rmse = sqrt(mean_squared_error(raw_values[200:], predictions))
pyplot.plot(raw_values[200:])
```

are the only lines of code I've really changed



**Jason Brownlee** July 6, 2017 at 10:11 am #

REPLY ↗

Consider starting with a simpler linear model:

<http://machinelearningmastery.com/start-here/#timeseries>

I do not believe LSTMs are the best or first choice.

<http://machinelearningmastery.com/suitability-long-term-forecasting/>



**vijay** December 7, 2017 at 8:32 pm #

I think I'm facing the same issue. Were



**Viorel Emilian Teodorescu** July 9, 2017 at 10:56 am #

Hi, Jason

Doesn't LSTM have a squashing gate at input? With outputs in between (-1, 1)? Then why do we need to prepare the input data to be between (-1, 1) if the first input gate will do this for us?

Am I missing something?

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** July 9, 2017 at 10:56 am #

REPLY ↗

Yes, you can rescale input to [0, 1]



**Hans** July 9, 2017 at 2:58 am #

REPLY ↗

Considering

```
unseenPredict = lstm_model.predict(X)
```

...how do we structure X to get a one step forward prediction of unseen data?

Or can we change some offsets in script "Complete LSTM Example" to get the same effect, and if so how?

Start Machine Learning



Paba July 10, 2017 at 8:45 pm #

REPLY ↗

Hi Jason,

Thanks for the excellent tutorial. I tried to modify the above code to include multiple timesteps and multiple lags in the model. I run give these parameters as input and paralleling run the script for different configurations to select the most accurate model. What do you think of the modifications I have done to the following functions? I am especially concerned about the time\_steps included in the model, is that correct?

```
def timeseries_to_supervised(data, lag=1, time_steps=0):
```

```
df = DataFrame(data)
```

```
columns = [df.shift(i) for i in range(1, lag+1)]
```

```
#considers the number of time_steps, t, involved a
```

```
#add next t x columns next to each x columnn
```

```
#question?? if time_steps = 3, does that mean y sh
```

```
#we trim the last 3 values from the dataset?
```

```
if time_steps > 0 :
```

```
columns_df = concat(columns, axis=1)
```

```
#note that I have multiplied i by -1 to perform left sh
```

```
timestep_columns = [columns_df.shift(i*-1) for i in r
```

```
timestep_columns_df =concat(timestep_columns, a
```

```
columns.append(timestep_columns_df)
```

```
columns.append(df)
```

```
df = concat(columns, axis=1)
```

```
df.fillna(0, inplace=True)
```

```
return df
```

```
def fit_lstm(train, batch_size, nb_epoch, neurons, lag, time_steps):
```

```
X, y = train[:, 0:-1], train[:, -1]
```

```
X = X.reshape(X.shape[0], time_steps+1, lag)
```

```
model = Sequential()
```

```
model.add(LSTM(neurons, batch_input_shape=(batch_size, X.shape[1] , X.shape[2]), stateful=True, return_sequences=True))
```

```
model.add(LSTM(neurons, stateful=True))
```

```
model.add(Dense(1))
```

```
model.compile(loss='mean_squared_error' optimizer='adam')
```

```
model.summary()
```

```
for i in range(nb_epoch):
```

```
model.fit(X, y, epochs=1, batch_size=batch_size, verbose=1, shuffle=False)
```

```
model.reset_states()
```

```
return model
```

```
def forecast_lstm(model, batch_size, X, lag,time_steps):
```

```
X = X.reshape(1,time_steps+1,lag)
```

```
pad = np.zeros(shape=(batch_size-1, time_steps+1, lag))
```

```
padded = np.vstack((X, pad))
```

```
yhat = model.predict(padded, batch_size=batch_size)
```

```
return yhat[0,0]
```

## Start Machine Learning X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

[Start Machine Learning](#)



**Jason Brownlee** July 11, 2017 at 10:30 am #

REPLY ↗

Great work!

**Kim Miller** July 13, 2017 at 9:01 am #

REPLY ↗

I want to extend this idea to several features  $x$  the lag values for each  $x$  time observations. Does it seem reasonable to give MinMaxScaler that 3D object? How does the  $y$  truth fit in what I give MinMaxScaler, since it's only 2D?



**Jason Brownlee** July 13, 2017 at 10:04 am #

No, I would recommend scaling each



**Kim Miller** July 13, 2017 at 11:51 am #

Above you seem to scale  $y$  along are not just a time-shifted copy of  $y$ , I assume  $y_{train}$  and  $y_{val}$ ? Then that's actually the

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** July 13, 2017 at 4:58 pm #

REPLY ↗

I would recommend scaling each series before any shifting of the series to make it a supervised learning problem.

I hope that answers your question, please shout if I misunderstood.



**Kim Miller** July 14, 2017 at 2:17 am #

I think the words "each series separately" threw me. I assume we can still scale all incoming values (all  $X$ 's and the time series that will later become  $y$  and shifted, one of the  $X$ 's) using a single scaler object. Then we create the lag values from scaled values. Finally, that single scaler object is used to invert  $y$  predictions. I have that right?



**Jason Brownlee** July 14, 2017 at 8:33 am #

Each series being a different "feature" or "column" or "time series" or whatever we call it.

Start Machine Learning

The coefficients (min/max or mean/stdev) used in the scaling of each series will need to be kept to reverse the operation on predictions and apply the operation to input data later. You can save the coefficients or the objects that wrap them.



**Sasha** July 13, 2017 at 4:15 pm #

REPLY ↗

Jason, great tutorial, thank you!

A question: why do you loop through epochs instead of just setting an appropriate number of epochs within fit() function? Would't it give the same result and be neater?



**Jason Brownlee** July 13, 2017 at 5:02 pm #

So that I can manually manage the re



**Sasha** July 13, 2017 at 5:22 pm #

Ah, I see, it has to do with "stateful".  
here: <http://philipperemy.github.io/keras-stateful-lstm/>



**kotb** July 13, 2017 at 6:59 pm #

Hi, DR jason

thank you very much for this tutorial.

I want to have multiple time steps , but i don't know how to modify function "timeseries\_to\_supervised()". I found another post of you talking about this , but you use function "create\_dataset()" i modify this function as follow:

```
def create_dataset(dataset, look_back=1):
 dataset = np.insert(dataset,[0]*look_back,0)
 dataX, dataY = [], []
 for i in range(len(dataset)-look_back):
 a = dataset[i:(i+look_back)]
 dataX.append(a)
 dataY.append(dataset[i + look_back])
 dataY=numpy.array(dataY)
 dataY = np.reshape(dataY,(dataY.shape[0],1))
 dataset = np.concatenate((dataX,dataY),axis=1)
 return dataset
```

please, check my modification , is it right OR what?

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** July 14, 2017 at 8:25 am #

REPLY ↗

See this post:

<http://machinelearningmastery.com/convert-time-series-supervised-learning-problem-python/>



**Eric** July 20, 2017 at 2:22 pm #

REPLY ↗

Hi Jason,

Sorry if it seem stupid but there is a part that I don't understand.

To predict you use: "yhat = model.predict(X, batch\_size=1)"  
But as we see X is :

```
train, test = supervised_values[0:-12], supervised_values[-12:]
scaler, train_scaled, test_scaled = scale(train, test)
yhat = forecast_lstm(lstm_model, 1, X)
```

So X is the 12 value (after being passed in the scaled variable).  
I'm not sure about them since in normal case we wouldn't know theirs

Once again really thank you for your tutorial, really

## Start Machine Learning

X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** July 21, 2017 at 9:29 am #

In the general case, you can pass in whatever values you want.

For example, if your model was defined by taking 6 days of values and predicting the next day and you want to predict tomorrow, pass in the values for today and 5 days prior.

Does that help?



**Eric** July 21, 2017 at 10:48 am #

REPLY ↗

I think it does help me. In my case I have values for each minute and I have to predict the next week (so more or less 10K of prediction).

I have data from the last year so there isn't any problem with my training, just wondered what I should do at the prediction part (so I can just instead of feeding it test\_scaled send him my training set again?)

Thank you for your help and quick reply!



**Jason Brownlee** July 22, 2017 at 8:29 am #

REPLY ↗

Yes, if your model is setup to predict given some fixed set of lags, you must provide those lags to predict beyond the end of the dataset. These may be part of your training dataset.

[Start Machine Learning](#)



**Eric** July 24, 2017 at 5:55 pm #

I don't think it's in my training dataset, for this part I'm pretty much doing it like you (the lag appearing when turning a sequence to a supervised learning problem). I'm pretty much feeding my model like you do. The problem for me being to know what to feed it when calling the predict command. Sending it "train\_scaled" was a bad idea (got poor result, predicting huge value when it should predict low and predicting low instead of predicting high). I'm working on it but every hint is accepted. Once again thank you and sorry for being a bit slow at learning/understanding.



**Jason Brownlee** July 25,

The argument to the predict function.

The data must be scaled the same way samples may vary, e.g. you may have:

Obviously, you don't need to make sure the data to predict will be the input required. This depends on the framing of your problem.

Does that help?

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**



**Eric** July 25, 2017 at 12:09 pm #

Thank you for your quick reply! I think I'm understanding it better but I have some part I have trouble to understand.

The input is X right?

If I follow your tutorial on input/output (<http://machinelearningmastery.com/time-series-forecasting-supervised-learning/>) and I take my case as an example (the database register the value each 3 minute and we want to predict the next value (so to begin for example it would be 12:24:00)):

Date,congestion

2016-07-08 12:12:00,92

2016-07-08 12:15:00,80

2016-07-08 12:18:00,92

2016-07-08 12:21:00,86

This is (part of) my training data, when turning it into supervised training data (and shifting) I get:

X, y

?, 92

92, 80

80, 92

**Start Machine Learning**

92, 86

86, ?

The problem is that I don't know X for predicting, I only know the X I use for my training (train\_scaled) and the one used to compare my results (test\_scaled).

What is the input I should feed it? I can't feed it my test\_scaled since in real situation I would have no idea of what it would be.

Sorry if my question seem stupid and thank you for taking time to explain it.



**Jason Brownlee** July 26, 2017 at 7:15 am #

It depends how you have

If the input (X) is the observation a predict the next time step.



**Eric** July 26, 2017 at 11:49 am

It is, each of my input X is case used in the tutorial).

Thank you for your answer, you an now!

Thanks for the tutorial too, they ready

## Start Machine Learning



You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** July 26, 2017 at 4:01 pm #

Glad to hear that.



**Eric** July 26, 2017 at 6:49 pm #

Just to be sure that I didn't made any mistake in my reasoning, if I take my example from before:

X, y

?, 92 / T-3

92, 80 / T-2

80, 92 / T-1

92, 86 / T

86, ? / T+1

to predict the next step (T+1) I have to use "yhat = model.predict(X, , batch\_size=batch\_size)" where X is 86 (after scaling/reshaping). Right?

Then I'll get the value predicted for

Start Machine Learning

difference to get a readable value).

If I want to predict farther then I continue (sending the scaled/reshaped value predicted at T+1 to get T+2 and then until I predict as far as wanted).

Thanks for your time and answers!



**Jason Brownlee** July 27, 2017 at 7:58 am #

Correct.



**Eric** August 4, 2017 at 12:05 pm

Thank you Jason with you the next week. I had 2 question tho. First: I removed the code for testing thing I have are the testing set in the When using this code (to train):

```
transform data to be stationary
raw_values = data.values
diff_values = difference(raw_values)

transform data to be supervised
supervised = timeseries_to_supervised
supervised_values = supervised.v
```

```
split data into train and test-sets
train = supervised_values[:-10000]
```

```
transform the scale of the data
scaler, train_scaled = scale(train)
```

and this one (to predict):

```
forecast the entire training dataset to build up state for forecasting
train_reshaped = train_scaled[:, 0].reshape(len(train_scaled), 1, 1)
lstm_model.predict(train_reshaped)
```

# walk-forward validation on the test data

```
predictions = list()
predictionFeeder = list() #Used to feed the model with the value of T-1
```

```
X, y = train_scaled[0, 0:-1], train_scaled[0, -1]
```

```
predictionFeeder.append(X) #Give the last value of training
```

```
for i in range(0, 10000):
```

```
make one-step forecast
```

```
yhat = forecast_lstm(lstm_model, predictionFeeder[i])
```

```
predictionFeeder.append(yhat)
```

```
invert scaling
```

```
yhat2 = invert_scale(scaler, testa[i + 1], yhat)
```

```
yhat3 = inverse_difference(raw_v
```

```
predictions.append(yhat3)
```

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**

and train a model (25 epoch) then predict the results I get result that are way too good (RMSE of 2 or less and prediction that have less than 5% of error).

Being used of thing going wrong for no reasons I decide to remove the testing data from the excel (even though it shouldn't change anything since I'm not using them (I've even set the variable to None at first)). Then when I do that the prediction is way less good and have some lag (though, if you remove the lag you still have a good results, just way less better than before).

Why is that?

My 2nd question is about lag, we can see on the prediction that while the shape of both chart (prediction and reality) look alike the prediction is raising/lowering before the reality, do you have any idea to fix it? Do you think changing the lag or timestep would help?

Once again thank you for your help without your tutorials.



**Jason Brownlee** August 4,

Sorry, I cannot debug you

Perhaps you are accidentally fitting evaluating it on the test (e.g. on da

I would encourage you to evaluate problem.

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Eric** August 4, 2017 at 4:02 pm #

Don't worry, I wouldn't ask you to debug.

Maybe, I don't know, I did remove the variable to be sure to never have affected the testing set and using it but since I'm human I may have made errors.

So changing lag would help me for theses gap between reality and prediction raise. Thank you I'll do that.

Thanks for your answer!



**Josep** July 21, 2017 at 8:08 pm #

REPLY ↗

Can I buy your books physically(not Ebook)?Thanks



**Josep** July 21, 2017 at 8:15 pm #

REPLY ↗

Sorry, now I have read that is not possible. Your books are amazing. Congratulations!

Start Machine Learning



**Jason Brownlee** July 22, 2017 at 8:33 am #

REPLY ↗

Thanks Josep.



**Pawel** July 26, 2017 at 7:59 pm #

REPLY ↗

Hi,

Thanks for very good tutorial. I have one question/doubt

in the following part of the code:

```
invert differencing
yhat = inverse_difference(raw_values, yhat, len(test))
```

should not we rely on predicted value instead of a validation we always refer to the test value(known) have only the predicted values(used as X), and of something?

my proposal:

```
invert differencing – (starting from 2nd loop cycle)
yhat = inverse_difference(predictions, yhat, 1)
```

Thanks in advance

Pawel

## Start Machine Learning

X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** July 27, 2017 at 8:00 am #

REPLY ↗

We could, but in this case the known observations are available and not in the future and it is reasonable to use them for the rescaling (inverting).



**Pawel** July 27, 2017 at 5:03 pm #

REPLY ↗

Hi, Thanks for explanation, got it now. Cause I train the model used e.g. MAY data (15 seconds samples) and then used that model to predict whole JUNE data. Afterwards I compared predicted data VS data that I got from JUNE, and I have to say that model does not work, after few prediction there is huge “off sync”,

In the validation phase as described in your case I got RMSE 0.11 so not bad, but in reality when you use predicted value(t-1) to predict next (t) value there is a problem.

Do you know how to improve the model? should I use multiple step forecasts, or lag features, input time steps?

Start Machine Learning

Thanks a lot.  
Pawel



**Jason Brownlee** July 28, 2017 at 8:29 am #

REPLY ↩

I would recommend brainstroming, then try everything you can think of.

I have a good list here:

<http://machinelearningmastery.com/improve-deep-learning-performance/>

Also, compare results to well-tuned MLPs with a window.



**Surya** July 27, 2017 at 6:49 pm #

Hey Jason, I am not following one point in  
reset\_states() is executed after every epoch. That means it is executed after the current batch. How does it make the network start from scratch?

Thanks



**Jason Brownlee** July 28, 2017 at 8:30 am #

State is maintained between the samples.

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Niklas** August 2, 2017 at 12:03 am #

REPLY ↩

Hi Jason,

thanks for the great tutorial. I have one question. Wouldn't it be better to use a callback for resetting the states? This would enable you to also use for example an EarlyStopping Monitor while training, here is what I changed:

```
class resetStates(Callback):
 def on_epoch_end(self, epoch, logs=None):
 self.model.reset_states()

 model.fit(X, y, epochs=nb_epoch, batch_size=batch_size, verbose=1, shuffle=False, callbacks=[resetStates(),EarlyStopping(monitor='loss', patience=5, verbose=1, mode='min')])
```



**Jason Brownlee** August 2, 2017 at 7:54 am #

REPLY ↩

Yes, that is a cleaner implementation for those problems that need to reset state at the end of each epoch.

Start Machine Learning

**Bharath** August 6, 2017 at 2:38 am #

REPLY ↗

Hello, can we extend this for anomaly detection techniques ?

**Jason Brownlee** August 6, 2017 at 7:40 am #

REPLY ↗

Perhaps, I have not used LSTMs for anomaly detection, I cannot give you good advice.

Perhaps you would frame it as sequence classification?

**Daniel Ruiz** August 7, 2017 at 1:36 pm #

Hi Jason,

In the persistence model plot there is a one time interval that is impossible to overcome this issue? What is causing the weight on time interval  $x[t-1]$ .

Here is an example of the dataset I am analyzing:

iteration: 969

Month=970, Predicted=-7.344685, Expected=280.0

iteration: 970

Month=971, Predicted=73.259611, Expected=212.0

iteration: 971

Month=972, Predicted=137.053028, Expected=0.0

Expected should be 280 and 212 (high magnitudes), and the model captures this more or less with 73 and 137, but this is one time interval behind.

Thanks!

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

**Jason Brownlee** August 8, 2017 at 7:41 am #

REPLY ↗

LSTMs are not great at autoregression problems and often converge to a persistence type model.

**Daniel Ruiz** August 8, 2017 at 8:57 am #

REPLY ↗

Ok thanks. What model would be a good alternative to capture this issue? I ran into the same problem with ARIMA. It could just be a difficult dataset to predict.

**Jason Brownlee** August 8, 2017 at 8:57 am #

Start Machine Learning

I recommend starting with a well-tuned MLP + Window and see if anything can do better.



**Eric** August 9, 2017 at 11:02 am #

REPLY ↗

Hi Jason,

Thanks to you I managed to get a working LSTM network who seem to have a good accuracy (and so a low RMSE)

But I've got a problem, do you know what could be the cause of extreme delay between reality values and predictions (my predictions have the same shape than reality)?

Best regards and please continue what you are doing!



**Jason Brownlee** August 10, 2017 at 6:40 pm #

Hi Eric, it might be a bug in the way you



**Eric** August 10, 2017 at 10:50 am #

It might be the case, I had to make sure that my last predicted value, I may have missed something at this moment. (Also I have to use Tflearn instead of Tensorflow but it shouldn't be a problem since Tflearn is a more transparent way to use tensorflow).

Thank you for your answer!

## Start Machine Learning

X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** August 10, 2017 at 4:40 pm #

REPLY ↗

Hang in there Eric!



**Eric** August 10, 2017 at 5:10 pm #

Thank you!

Well.. I have a gap of 151 (reference to pokemon?).

Just to try I removed these 151 values from my training set, I now have no gap of values (and frankly, the accuracy seem good for a 15 epoch training). I know that this is no way a fix but make me wonder where did I fail..

Start Machine Learning

**Eric** August 10, 2017 at 7:21 pm #

Could it be that while my training set is on 400K of value my prediction start 151 value before the end (so predicting the value for 399849) of the training set (which is strange since the information from training tell me that I'm training on the 400K of data). It would mean that my machine is trying to predict some point of time used for training. Or it would mean that the 151 last data weren't used at all for training (I tried to reduce the number of data but it's the same problem).

**Jason Brownlee** August 11, 2017 at 6:41 am #

The algorithm is trained so the last sample is what is key.

**Eric** August 11, 2017 at 12:24 pm #

Thanks for the reply.  
When I think about it my prediction start at the first value of my training set. So my model start his prediction now at the first value of my training set. This is strange since my training end when I have 151 values (and each line correspond to a month).  
I must have made a mistake somewhere.

Thanks for your answer I think I'm on something!



## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

**Jason Brownlee** August 12, 2017 at 6:45 am #

Hang in there!

REPLY

- srn - **Stone** August 9, 2017 at 12:16 pm #

One thing I don't understand in these models is, how can I predict future.

Say, I have trained my model and let's say we're living August 8, 2017. Now `model.predict()` seems to require some test samples to predict anything, in this case values from `test_scaled`. So what do I give it when I want to get a prediction for August 9, 2017, which we don't know yet? So I want to know `yhat(t+1)`. Can I ask the model to give me forecasts for the next 5 days?

**Eric** August 9, 2017 at 2:11 pm #

REPLY

[Start Machine Learning](#)

I may be wrong but in my case I trained my model on the data I had (let's say one year, until August 8, 2017) and to predict the value T you send the value T-1 (so in this case the value of August 8 2017), you'll get the value of August 9.

Then if you want August 10 (T+1) you send the value of August 9 (T). But this code here show a one step forecast implementation. Maybe if you want to predict more you should look for multi step forecasting? I'm think there is some example of it on this website.

Jason would give you a better answer I think.



**Jason Brownlee** August 10, 2017 at 6:47 am #

REPLY ↗

Nice one Eric!

Yes, I'll add, if the model was trained to predict finalized, the model will need the observation

It all comes down to how you frame the problem you need.



**Jason Brownlee** August 10, 2017 at 6:43 pm #

If you would like to predict one day, the

You have to decide what the model will take as month or year of data. This is how the model is

Then, when it's trained, you can use it to predict future values using perhaps the last few values from the training data. You can use a model trained to predict one day to predict many days in a recursive manner where predictions become inputs for subsequent predictions. More about that here:

<http://machinelearningmastery.com/multi-step-time-series-forecasting/>

You can also make multi-step forecasts directly, here is an example:

<http://machinelearningmastery.com/multi-step-time-series-forecasting-long-short-term-memory-networks-python/>

Does that make things clearer?

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

- srn -

**Stone** August 10, 2017 at 11:08 am #

REPLY ↗

Thanks Eric and Jason. That's exactly how I've done it, but it seems to me, that the prediction that I get is not one timestep forward (T+1), but it's a prediction of T, which doesn't make sense. I'm using close price of a stock as my data. I'll have to check again, if it's really is making a prediction for the future, as you insist. and I will check the Jason's links.

Anyways, thanks Jason for the excellent tutorials! 😊

Start Machine Learning

- srn - **Stone** August 10, 2017 at 12:22 pm #

REPLY ↗

Here's a sample of the close prices:

2017-05-30 5.660  
2017-05-31 5.645  
2017-06-01 5.795  
2017-06-02 5.830

As a matter of fact, it seems to me that the predicted values lags one time step behind from expected ones.

Day=1, Predicted=5.705567, Expected=5.660000

Day=1, Predicted=5.671651, Expected=5.6450

Day=1, Predicted=5.657278, Expected=5.7950

Day=1, Predicted=5.805318, Expected=5.8300

Here I'm going forward one day at a time and was expected one day before than. And anyway

Let's examine the second line:

Day=1, Predicted=5.671651, Expected=5.6450

The expected price (which is the actual price on after the trading day is closed. What I expect this is closer to the actual price the day before!). See changed anything in the framework except the

## Start Machine Learning



You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** August 10, 2017 at 4:41 pm #

REPLY ↗

That is a persistence model and the algorithm will converge to that (e.g. predict the input as the output) if it cannot do better.

Perhaps explore tuning the algorithm.

Also, I believe security prices are a random walk and persistence is the best we can do anyway:

<http://machinelearningmastery.com/gentle-introduction-random-walk-times-series-forecasting-python/>



**Firnas** August 12, 2017 at 1:27 pm #

REPLY ↗

Please do some tutorials RNN, NN with more categorical values in the dataset? please, I could not find many resources using the categorical values.

Thanks



**Jason Brownlee** August 13, 2017 at 9:45 am #

REPLY ↗

## Start Machine Learning

I have a few tutorials on the blog – for example, text data input or output are categorical variables.

What are you looking for exactly?



**nandini** August 16, 2017 at 5:23 pm #

REPLY ↗

is it possible to write the code RNN Regression ,what is the activation function i need to give for regression RNN,how it is difference from classification to regression?



**Jason Brownlee** August 17, 2017 at 6:36 am #

Yes. The output would be 'linear' for regression.



**Maria Jimenez** August 17, 2017 at 1:20 am #

Hi Dr. Brownlee,

I am trying to understand this code. I am a beginner.

My questions about your code:

1. What does "transform data to be stationary" mean?
2. Why do you create diff? What does it mean?
3. If the raw\_values are 36 data, why diff has only 35?

I appreciate your response in advance,

Maria

## Start Machine Learning

X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Maria Jimenez** August 17, 2017 at 1:43 am #

REPLY ↗

Never mind, I already understand. 😊



**Jason Brownlee** August 17, 2017 at 6:46 am #

REPLY ↗

Glad to hear it.



**Jason Brownlee** August 17, 2017 at 6:46 am #

REPLY ↗

Consider using the blog search.

More about stationary data here:

<http://machinelearningmastery.com/time-series-forecasting-long-short-term-memory-network-python/>

Start Machine Learning

More about differencing here:

<http://machinelearningmastery.com/remove-trends-seasonality-difference-transform-python/>

I hope that helps.



**Sourabh** August 18, 2017 at 4:41 am #

REPLY ↗

Thanks, for the awesome blogs,

But, I have a doubt, for time series forecasting, classical methods like autoregression and ARIMA, or this machine learning approach using LSTM RNN model, which is the better approach? Both are good on their own, So, what should be our first choice while forecasting any dataset? How should we choose between them?



**Jason Brownlee** August 18, 2017 at 6:27 pm #

It depends on your problem. Try a few

I recommend starting with linear methods like ARIMA such, then an MLP, then perhaps an RNN.



**Bilal** August 19, 2017 at 10:29 pm #

When I changed train/test ration to 0.8/0.2 with another data that contains around 4000 records, it's taking hours(i still couldn't get result, i don't know how long it'll take). Can you please suggest me how I change settings for very long sequences. Thanks

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** August 20, 2017 at 6:06 am #

REPLY ↗

I have some ideas here:

<http://machinelearningmastery.com/handle-long-sequences-long-short-term-memory-recurrent-neural-networks/>



**Arun** August 22, 2017 at 1:45 am #

REPLY ↗

Hi Jason, for long sequence like 4000 records, can you please help me to understand the changes from code point of view as per your current example, the above link is for classification and i am looking for sequence prediction problem.

**Jason Brownlee** August 22, 2017 at 1:45 pm #

Start Machine Learning



You must split your sequence into subsequences, for example, see this post for ideas:  
<https://machinelearningmastery.com/truncated-backpropagation-through-time-in-keras/>



**Arun Menon** August 23, 2017 at 12:25 am #

Thank you Jason, can you please help us to understand with a complete code example ? The above link just have an idea but not a code implementation, I want to understand from code point of view.

Regards,  
Arun



**hirohi** August 21, 2017 at 12:54 pm #

Is the evaluation method in this post really backtest. Walk-forward corresponds to cross validate models to learn, so these methods are not suitable post.



**Jason Brownlee** August 21, 2017 at 4:26 pm #

Yes, we are using walk-forward validation here.

Walk-forward validation is required for sequence prediction like time series forecasting.

See this post for more about the method and when to use it:

<https://machinelearningmastery.com/backtest-machine-learning-models-time-series-forecasting/>



**hirohi** August 22, 2017 at 11:16 am #

REPLY ↗

Thank you for your reply. Sorry, that's not what I mean it. I've read the post you pasted above, in which you split data into train/test 2820 times and build 2820 models. But, in this LSTM post, you split data and build a LSTM network only once, so I suggest the test in this article is not walk forward. ordered(not shuffle) hold-out test?



**hirohi** August 22, 2017 at 11:17 am #

REPLY ↗

sorry, please delete this comment

[Start Machine Learning](#)



**Jason Brownlee** August 23, 2017 at 6:36 am #

REPLY ↗

Above, the model is evaluated 30 times.

Within the evaluation of one model, we use walk forward validation and samples are not shuffled.

I recommend re-reading the tutorial and code above.



**hirohi** August 23, 2017 at 12:26 pm #

REPLY ↗

Thank you for your reply! Oh, repeat=30! Sorry, I misunderstood. But I think in the back-test article, the size of train data changes, was fixed, right? Sorry for so many questions.



**Gauthier** August 23, 2017 at 10:38 pm #

Thanks for your work, it is seriously awesome!

I wanted to see the chart comparing training loss vs validation loss. I uploaded it here: <https://ibb.co/kqcnBk>. As you mentioned in your other post, "From the plot of loss, we can see that the model has overfit to the training datasets and underfits to the validation datasets (labeled test). If these parallel problems persist, stop training at an earlier epoch."

However, in this case, the test loss never really gets lower than the training loss. This is a problem, and in general, does that mean that no meaningful test-set pattern is learned from the training, as seems to be corroborated by the RMSE being close on average to the persistent/dummy method?

Is this the end of the road? Have NN's failed in their tasks? If not, where to go next?

The main thing I changed in the code was the end of the fit\_lstm function:

```
for i in range(nb_epoch):
 e=model.fit(X, y, epochs=10, batch_size=batch_size, verbose=0, validation_split=0.33, shuffle=False)
 model.reset_states()
return model, e
```



**Jason Brownlee** August 24, 2017 at 6:42 am #

REPLY ↗

Thanks Gauthier.

It may mean that the model under provisioned for the problem.



**Gauthier** August 24, 2017 at 6:37 pm #

REPLY ↗

OK... So in this case would you suggest another ML algorithm?

[Start Machine Learning](#)



**Jason Brownlee** August 25, 2017 at 6:41 am #

REPLY ↗

Both would be a safe bet.



**malcolm** November 10, 2019 at 4:33 pm #

REPLY ↗

Hi, @Gauthier, may i know how you editted ur code to get the loss function?



**Muni** August 28, 2017 at 12:19 pm #

Json,

Thanks for your blogs and tutorials. Being a researcher, your tutorials gave me an excellent start to carry out my research.

I tried to work with the example in this page, But, I can't find the link to datamarket.com. Could you please put the link to the dataset? Please let me know.



**Jason Brownlee** August 29, 2017 at 5:01 pm #

Thanks, I'm glad to hear that.

Here is the full dataset:

```

1 "Month", "Sales"
2 "1-01", 266.0
3 "1-02", 145.9
4 "1-03", 183.1
5 "1-04", 119.3
6 "1-05", 180.3
7 "1-06", 168.5
8 "1-07", 231.8
9 "1-08", 224.5
10 "1-09", 192.8
11 "1-10", 122.9
12 "1-11", 336.5
13 "1-12", 185.9
14 "2-01", 194.3
15 "2-02", 149.5
16 "2-03", 210.1
17 "2-04", 273.3
18 "2-05", 191.4
19 "2-06", 287.0
20 "2-07", 226.0
21 "2-08", 303.6
22 "2-09", 289.9
23 "2-10", 421.6
24 "2-11", 264.5
25 "2-12", 342.3

```

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

```

26 "3-01",339.7
27 "3-02",440.4
28 "3-03",315.9
29 "3-04",439.3
30 "3-05",401.3
31 "3-06",437.4
32 "3-07",575.5
33 "3-08",407.6
34 "3-09",682.0
35 "3-10",475.3
36 "3-11",581.3
37 "3-12",646.9

```



**mikiwate** September 3, 2017 at 3:35 pm #

Hi Jason,

I am getting NAN on the 12 month:::

```

Month=1, Predicted=440.400000, Expected=440.4
Month=2, Predicted=315.900000, Expected=315.9
Month=3, Predicted=439.300000, Expected=439.3
Month=4, Predicted=401.300000, Expected=401.3
Month=5, Predicted=437.400000, Expected=437.4
Month=6, Predicted=575.500000, Expected=575.5
Month=7, Predicted=407.600000, Expected=407.6
Month=8, Predicted=682.000000, Expected=682.0
Month=9, Predicted=475.300000, Expected=475.3
Month=10, Predicted=581.300000, Expected=581.3
Month=11, Predicted=646.900000, Expected=646.900000
Month=12, Predicted=646.900000, Expected=nan

```

from the code below:

```

predictions = list()
for i in range(len(test_scaled)):
 # make one-step forecast
 X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
 #yhat = forecast_lstm(lstm_model, 1, X)
 yhat=y
 # invert scaling
 yhat = invert_scale(scaler, X, yhat)
 # invert differencing
 yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)
 # store forecast
 predictions.append(yhat)
expected = raw_values[len(train) + i + 1]
print('Month=%d, Predicted=%f, Expected=%f' % (i+1, yhat, expected))

```

I do not understand as to why this is happening.

thanks for the feedback.

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)



**Jason Brownlee** September 3, 2017 at 3:49 pm #

REPLY ↗

Interesting.

Does this happen every time?

I wonder if it is platform specific. Are you on a 32-bit machine? What version of Python?



**sms** January 26, 2018 at 10:57 pm #

REPLY ↗

Happened to me too, because after downloading from

<https://datamarket.com/data/set/22r0/sales>

period#!ds=22r0&display=line the end of the

"3-10",475.3

"3-11",581.3

"3-12",646.9

contained

"

Sales of shampoo over a three year period

"

It worked when removed

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** January 27, 2018 at 5:57 am #

REPLY ↗

Nice!



**Eric** September 3, 2017 at 6:30 pm #

REPLY ↗

Dear Jason,

I had a question, what are we supposed to send as first parameter of `inverse_difference` if we go with the idea that we don't know the future?

Regardless what I do, if I give it the testing set then it have perfect accuracy and values. If I send the training set it will look like the last `len(test_scaled)` training values.

Best regards



**Jason Brownlee** September 4, 2017 at 4:26 am #

REPLY ↗

You would pass a list of past observations so that you can reverse the difference operation.

Start Machine Learning

**Eric** September 4, 2017 at 9:42 am #

REPLY ↗

Dear Mr Jason,

Sorry, I have trouble understanding, what do you mean by “list of past observations”? Would that be what we predict (the predictions list)?

It would then follow what “Pawel” asked and said previously?

“my proposal:

```
invert differencing – (starting from 2nd loop cycle (1st would be the starting point
(raw_values[-1]))
```

```
yhat = inverse_difference(predictions, yhat, 1)"
```

Best regards,

Eric

**Jason Brownlee** September 7, 2017 at 10:30 am #

Sorry, I mean it is the true or real value.  
E.g. not predictions.

**Eric** September 7, 2017 at 5:04 pm #

Dear Mr Jason,

Thank you, so let's imagine a 500 000 entry file.

My training would be the 490 000 first line. I want to predict the last 10 000 lines.

For predicting I send the last line of the training set (then for the next prediction I send my first prediction without having unscaled or reversed the difference).

To get the true value I send as argument of “Invert\_difference” my training set, or more specially, my last 10 000 thousands line?

Best regards,

Eric Godard.

## Start Machine Learning

X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**

**Jason Brownlee** September 9, 2017 at 11:40 am #

I would not recommend prediction 10K time steps. I expect skill would be very poor.

Nevertheless, observation 490000 would be used to invert observation 490001, then the 490001 transformed prediction would be used to decode 490002 and so on.

Does that help?

**Start Machine Learning**



**Eric** September 8, 2017 at 10:09 am #

Sorry, I may have badly worded myself here.

My dataset have one entry per minute, one day is 1440 entry, I have years worth of data in my data set. I want to predict the next day, should I send the last day of the training set (trained\_scaled[-1440:]) as argument of inverse\_difference?

Best regards,



**Jason Brownlee** September 8, 2017 at 10:09 am #

Perhaps this post will help:  
<https://machinelearningmastery.com/time-series-forecasting-long-short-term-memory-network-python/>



**Brian** September 16, 2017 at 8:55 pm #

I keep getting inconsistent result of RMSE  
`np.random.seed(1) # for consistent results`  
 to the very top of my code so now every run produces

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**



**Jason Brownlee** September 17, 2017 at 5:27 am #

REPLY ↩

It is hard to get reproducible results, see this post for more information Brian:  
<https://machinelearningmastery.com/reproducible-results-neural-networks-keras/>



**Irene** September 18, 2017 at 8:48 pm #

REPLY ↩

When I try to load the data after executing the following code line with Python 3.6 I get the following error:

```
series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0,
squeeze=True, date_parser=date_parser)
```

`TypeError: ufunc 'add' did not contain a loop with signature matching types dtype('<U32') dtype('<U32') dtype('<U32').`

Can you help me, please?



**Jason Brownlee** September 19, 2017 at 10:09 am #

**Start Machine Learning**

Ouch, I have not seen that error before.

Perhaps confirm that the data is CSV and that you have removed the footer.

Perhaps ensure that your environment was installed correctly:

<http://machinelearningmastery.com/setup-python-environment-machine-learning-deep-learning-anaconda/>

If that does not help, perhaps try posting to stackoverflow?



**Priyank** September 18, 2017 at 10:05 pm #

REPLY ↗

Hello Jason,

Thanks for sharing this post, very helpful as current query.. How can we save a model with best RSME

Thanks



**Jason Brownlee** September 19, 2017 at

This post shows you how to save your  
<https://machinelearningmastery.com/save-load-machine-learning-models-python-keras/>



**Alex** September 21, 2017 at 10:27 am #

REPLY ↗

Thanks Jason for the great tutorial!

Can you please clarify this statement for me: "The batch\_size must be set to 1. This is because it must be a factor of the size of the training and test datasets"

I don't understand how 1 in this case is a factor of the training and test datasets?

Further, for time series data, can we have a batch size greater than 1? If not, what was the relevance of the above statement?

Thank you



**Jason Brownlee** September 21, 2017 at 4:21 pm #

REPLY ↗

Yes, it is a constraint for the chosen framing of the problem so that we can make one-step forecasts with a stateful LSTM.

It is not a constraint generally.

**Nidhi** September 22, 2017 at 9:58 am #

Start Machine Learning



Great post Jason, very granularly explained LSTM example. Thanks a lot. I am trying to train a model to predict scientific notation data like "9.543301358961403E-9", could you suggest a good way to rescale this data that fits LSTM the best?



**Jason Brownlee** September 23, 2017 at 5:34 am #

REPLY ↗

For input data, rescale values to the range 0-1.



**vardhan** October 2, 2017 at 4:22 pm #

REPLY ↗

Hi, I have gone through many of your posts. In this post, is BPTT applied in this example? or just applied when training the current output??

And in this example, the parameters are updated after each step in contrast to stochastic gradient descent (online-learning).



**Jason Brownlee** October 3, 2017 at 5:39 pm #

REPLY ↗

This post provides an intro to BPTT:

<https://machinelearningmastery.com/gentle-introduction-backpropagation-through-time/>

BPTT is applied when you have more than one hidden layer.

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Aljo Jose** October 3, 2017 at 11:42 pm #

REPLY ↗

Hi Jason,

Thank you for great tutorial. Please let me know the difference between below two –

1) `model.fit(trainX, trainY, epochs=3000, batch_size=1, verbose=0)`

2) `for i in range(3000):`

`model.fit(X, y, epochs=1, batch_size=batch_size, verbose=0, shuffle=False)`

`model.reset_states()`



**Jason Brownlee** October 4, 2017 at 5:47 am #

REPLY ↗

Not much.



**Nidhi** October 16, 2017 at 1:11 am #

REPLY ↗

Start Machine Learning

Hi Jason,

You have reshaped value of X and then predicted the value of yhat, after that invert scaling of yhat is done and then this value is added to the raw value of the previous month. My question is what if I add the reshaped value of yhat and reshaped value of previous month and then call invert scale\_function and not call inverse\_difference function at all. Will both give the same result?

Best Regards



**Jason Brownlee** October 16, 2017 at 5:44 am #

REPLY ↗

I believe so, we need to group yhat with anything, even 0 values to ensure that we meet the shape requirements when inverting the transform.



**Nidhi** October 16, 2017 at 1:43 am #

I ran your code multiple times with different random seeds and the results are plotting both "pyplot.plot(raw\_values[-12:])" and "pyplot.plot(yhat[-12:]). They coincide at certain points (like in your example graph). It's good to at least show a similar trend, but why in most of the results do the blue line (raw values) more closely associate with the blue line for time steps 24-36? I am plotting raw values of month 24-36 and predictions for month 37-48. I am plotting both prediction values (predicted from month 24-36) and raw values (month 37-48). Please clarify, why is that the case?. I am definitely not using a persistence model. Thanks again.

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** October 16, 2017 at 5:45 am #

REPLY ↗

This suggests a persistence model (e.g. no skill).



**Rubens** October 16, 2017 at 4:50 am #

REPLY ↗

I'm using lag>1

Shouldn't the function "timeseries\_to\_supervised" revert the range range(1, lag+1) ?

```
columns = [df.shift(i) for i in reversed(range(1, lag+1))]
```



**Thong Bui** October 16, 2017 at 5:00 am #

REPLY ↗

Hi Jason,

Thanks so much for the in-depth tutorial. I have a question about this block of code in fit\_lstm():

Start Machine Learning

```
for i in range(nb_epoch):
 model.fit(X, y, epochs=1, batch_size=batch_size, verbose=0, shuffle=False)
 model.reset_states()
```

Can it simply be done by this line?

```
model.fit(X, y, epochs=nb_epoch, batch_size=batch_size, verbose=0, shuffle=False)
```

Thanks for your clarification



**Jason Brownlee** October 16, 2017 at 5:46 am #

REPLY ↗

Yes, but I wanted to explicitly reset state after each batch.



**Thong Bui** October 16, 2017 at 7:37 am #

what's the difference between them?



**Jason Brownlee** October 16, 2017 at 7:40 am #

You can learn more about LSTM in my course:  
<https://machinelearningmastery.com/univariate-time-series-forecasting-with-python-keras/>

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Danier** October 19, 2017 at 7:14 pm #

REPLY ↗

Hi Jason,

I am a college student. I am particularly interested in machine learning and are working hard on it. But my English is not very good, so my problems may be a little rough.

1. The rmse in this article is somewhat large, and you have some other articles that describe how to adjust the parameters to reduce the rmse? What are the reasons for regulation?
2. Would you have a better article about using LSTM to predict time series problems? Ps: in Uni-variate time series. Because I did the experiment myself, if I use the model of this article, my experiment accuracy is 85% to 90%. I want to improve it.
3. Is there any article to solve the multivariate time series forecasting problem? If so, where is the latest?
4. I have been running the model several times with my own macbookpro, which is too slow. If I want to be more efficient, can I buy some GPU cloud services? Like amazon? But it seems that these cloud services are especially expensive. Do you have any good solutions?

Thank you so much! I want to say that you are my machine learning primer teacher.



**Jason Brownlee** October 20, 2017 at 5:30 pm #

Start Machine Learning

REPLY ↗

This post has a list of things to try to lift model skill:

<http://machinelearningmastery.com/improve-deep-learning-performance/>

Generally, I would not recommend LSTMs for autoregression problems, see this post:

<https://machinelearningmastery.com/suitability-long-short-term-memory-networks-time-series-forecasting/>

Here is a multivariate time series example:

<https://machinelearningmastery.com/multivariate-time-series-forecasting-lstms-keras/>

In practice, AWS is very cheap:

<https://machinelearningmastery.com/develop-evaluate-large-deep-learning-models-keras-amazon-web-services/>



**Danier** October 24, 2017 at 3:14 pm #

Hello, Jason

Thank you for your last reply. I have improved the `bias_regularizer =L1L2 (0.01, 0.01)` to reduce the variance. I would like to ask about Tutorial Extensions.

I want to know how to complete multi-step Forecasts. I have 24 sets of data, I would like to predict the six sets coming after. I have no expected value and are empty. I tried 10 or 20 steps. Is it bad? Or is my understanding of LSTM not enough? I am depressed.....

I hope you can give me some Suggestions. Thank you for your help as my reference?

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**



**Jason Brownlee** October 24, 2017 at 4:03 pm #

REPLY ↗

Here is an example:

<https://machinelearningmastery.com/multi-step-time-series-forecasting-long-short-term-memory-networks-python/>



**argyn** October 25, 2017 at 1:23 am #

REPLY ↗

Hello

I have an issue with your terminology regarding the dynamic forecasting, which you defined as follows: "re-fit the model or update the model each time step of the test data as new observations from the test data are made available (we'll call this the dynamic approach)."

That's not what's called dynamic forecasting in time series and econometric analysis. The dynamic forecasting involves using the previous forecast value as input into the next forecast. In your case, you are using the lagged value of the dependent variable as a regressor in the estimation of the model:  $X(t) = Y(t-1)$ . For instance, see this paper

<https://files.stlouisfed.org/files/htdocs/publications/>

**Start Machine Learning**

So, the dynamic multi-step forecast would involve something like this:  $\text{Yhat}(t) = \text{function}(\text{Yhat}(t-1))$ . It doesn't need to re-estimate the model at each step to be dynamic. All it needs is to use the previous forecast value to come up with the next.

Particularly this step would change to use  $\text{yhat}$  as  $X$ :

```
X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
yhat = forecast_lstm(lstm_model, 1, X)
```



**Jason Brownlee** October 25, 2017 at 6:49 am #

REPLY ↗

Thanks for the clarification.



**Sulgi** October 25, 2017 at 7:40 pm #

Hi I like to input time difference between time steps. I manually put this term into the weight, so weight now includes time difference in input,  $X$ . Do you have any advice?



**Jason Brownlee** October 26, 2017 at 5:20 pm #

Sorry, I have not done this, perhaps take a look at this post:

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Eren** October 26, 2017 at 10:55 am #

REPLY ↗

My observation is the model tries to predict the previous value so that it is able to pick a moderate precision. (I used another dataset)



**tony** October 26, 2017 at 5:01 pm #

REPLY ↗

Hi, Jason:

Use the shampoo sales data and my own test data, the results are good. But when the time series with seasonal trends, LSTM result is very bad. So I want to ask is there any way to solve include seasonal time series forecasting problem?  
thank you.



**Jason Brownlee** October 27, 2017 at 5:15 am #

REPLY ↗

You can seasonally adjust your data:

<https://machinelearningmastery.com/time-series-forecasting-long-short-term-memory-network-python/>

Start Machine Learning



**Frank** November 7, 2017 at 11:49 am #

REPLY ↗

While the visuals and numbers in this tutorial are compelling, there are several things seriously wrong that aren't apparent until you go over it with a very fine-toothed comb. Among all the arrays, lists, and indices flying all over the place, there are errors that make the results non-reproducible for one.

At the detail level you've got this prediction output:

Month=1, Predicted=339.700000, Expected=339.700000

Month=2, Predicted=440.400000, Expected=440.400000

Month=3, Predicted=315.900000, Expected=315.900000

....

However running the code block that is supposed to calculate "Expected" is supposed to come from:

```
expected = raw_values[len(train) + i + 1]
```

Index i starts out at 0 because you are looping through the first three months. The first be raw\_values[len(train) + 0 + 1] which is raw\_values[3] (Month 1).

This kind of thing seems minor, but it's actually very non-reproducible output, but instead they end up having

At a more general level, the purpose of building a regression model is to fit metrics on test data and leave it at that. Instead, it's common to fit the model on training data that you haven't encountered yet. From the beginning, you take raw\_values, remove the trend, then you forecast the residuals, and then add the trend back in with an inverse difference function. Right there is a major violation as you are literally imputing known information into what is supposed to be the forecast. The inverse difference function takes raw\_values as a parameter, however in a real out-of-sample scenario, you wouldn't have access to those raw\_values, so you couldn't simply adjust the forecasted residuals to reflect the trend in that manner. How did no one catch this?

I'm doing a research project on LSTMS for time-series data and trying to extract something useful from this tutorial that I can apply to this problem, but I feel there are too many land-mines hidden throughout.

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Luciano** November 7, 2017 at 10:37 pm #

REPLY ↗

Hi Jason!

I have a dataset of 200 thousand time-series. Each with 150 timestamps.

Let's say they are sales of the same product but for different stores.

What is the best way to design the problem?

I can't figure out what would be the ideal way to structure it. Just feed each sequence of 149 values to predict the 150th; Or should I do a rolling-window? If I give the whole sequence, I'm giving it a lot of context to work on, but I'm afraid the series is too big

[Start Machine Learning](#)

Any thoughts?

There is a lack of literature on how to design time-series problems with LSTMs. You are the only one that is talking about it. Thanks!



**Jason Brownlee** November 8, 2017 at 9:23 am #

REPLY ↗

I would recommend exploring many different framings of the problem and see what works best.

200K series is a lot. Perhaps you can group series and use sub-models and combine results in an ensemble of some sort.



**Ashima** November 9, 2017 at 6:02 am #

Hi Jason,

Would like to request your help in updating the above code so that it updates every observation and accordingly predicts the future values to suit that purpose.



**Jason Brownlee** November 9, 2017 at 10:02 am #

See this post:

<https://machinelearningmastery.com/update-lstm-networks-training-time-series-forecasting/>



**Ashima** November 9, 2017 at 6:16 pm #

REPLY ↗

Also if I need to create a lookback of a number other than 1, say 7 what all do I need to update in this code above to get it running for a look back of 7.



**Bahar** November 14, 2017 at 4:24 pm #

REPLY ↗

Hi Jason,

Thanks for the useful tutorial.

```
I have a question about X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
yhat = forecast_lstm(lstm_model, 1, X)
```

(line110):

I didn't understand how the prediction will use timestamps to predict on test data set.

As far as I know, eventually we want to predict based on future time. So we should test based on timestamps (not values) and compare the prediction results with real test values to see the accuracy of the prediction.

[Start Machine Learning](#)

But in the code, it seems that you are predicting based on the values and comparing again with those values!

Would you please explain it to me? Maybe I misunderstood some part.

Best Regards,  
Bahar



**Jason Brownlee** November 15, 2017 at 9:48 am #

REPLY ↗

We are predicting t+1.

Perhaps this post will help to better understand

<https://machinelearningmastery.com/index-slicing-time-series-data/>



**Pradeep** November 19, 2017 at 5:57 pm #

Jason awesome tutorial which helps in great way.  
I had 2 quick questions.

1. Is there a good example of adding other “features” such as whether a particular hour is rush hour or not, whether a particular day of the week is a weekend etc?
2. I have hourly data for 3 years and wish to forecast the next year. I have approximately 14600 training rows and have to forecast 8760 rows. This seems to be the right approach. However will be able to do this in advance?

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** November 20, 2017 at 10:12 am #

REPLY ↗

This post shows you how to use multiple input variables:

<https://machinelearningmastery.com/multivariate-time-series-forecasting-lstms-keras/>

Multistep forecasting may result in poor results. I'd recommend testing a suite of approaches to see which works best for your data:

<https://machinelearningmastery.com/multi-step-time-series-forecasting/>



**Pradeep** November 20, 2017 at 4:17 pm #

REPLY ↗

Awesome Thanks Jason.I will go through the links suggested.  
Are there any other references which you can suggest for multi step forecasting?



**Jason Brownlee** November 22, 2017 at 10:38 am #

REPLY ↗

Start Machine Learning

See this post:

<https://machinelearningmastery.com/multi-step-time-series-forecasting-long-short-term-memory-networks-python/>



**Harini** November 20, 2017 at 1:45 am #

REPLY ↗

Hi Jason, I have a quarterly sales dataset(below) and tried to use the same code on this post:  
Month, Sales

5-09,11  
5-12,20  
6-03,66  
6-06,50  
6-09,65  
6-12,63  
7-06,25  
7-12,34

I trained the model with data from all preceding quarters except last quarter: predicted=-1.329951, actual=34.000000

The sales number can't be negative but the model predicted a negative value. Is there a way to prevent the model from predicting

## Start Machine Learning X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** November 20, 2017 at 1:46 pm #

REPLY ↗

I would recommend starting with some simple linear models to better understand your data and how time series works:

<https://machinelearningmastery.com/start-here/#timeseries>



**Harini** November 25, 2017 at 6:27 am #

REPLY ↗

Thanks Jason, I tried linear regression but it didn't seem to give me expected results.

Then I tuned the LSTM in your code to predict positive values by

- \* changing the scaler range from (-1,1) to (0,1)
- \* adding the activation function 'relu' to the model



**SRIRAM P** November 21, 2017 at 5:59 pm #

REPLY ↗

Hi Jason,

Need a small clarification,

Does the for loop started at 100th line ends at 102?

Start Machine Learning

```

for r in range(repeats):
 # fit the model
 lstm_model = fit_lstm(train_scaled, 1, 3000, 4)
 # forecast the entire training dataset to build up state for forecasting
 train_reshaped = train_scaled[:, 0].reshape(len(train_scaled), 1, 1)
 lstm_model.predict(train_reshaped, batch_size=1)

```

here in the above, does the lines

```

train_reshaped = train_scaled[:, 0].reshape(len(train_scaled), 1, 1)
lstm_model.predict(train_reshaped, batch_size=1)

```

also belongs to for loop?



**AB** November 28, 2017 at 5:13 am #

hey Jason,

How to Predict next one month for above example'



**Jason Brownlee** November 28, 2017 at 8:15 pm #

Perhaps this post will help you make predictions.

<https://machinelearningmastery.com/make-predictions-time-series-data/>



**Vijay** December 6, 2017 at 11:32 pm #

REPLY ↗

Hey Jason,

Thanks for all your wonderful articles. I've trying this model to predict a spending pattern of a customer. I have mean of the weekly expenditure of a customer. Once I finished training, i'm seeing this pattern where the predicted output for current is matching the expected output in the last week. I changed yhat = y to see if all the conversions are taking place correctly. But even yhat = y is showing the same behaviour.

```

Week=1, Predicted=975.556027, Expected=989.100588
Week=2, Predicted=989.100588, Expected=928.604400
Week=3, Predicted=928.604400, Expected=921.209794
Week=4, Predicted=921.209794, Expected=919.813532
Week=5, Predicted=919.813532, Expected=904.214533
Week=6, Predicted=904.214533, Expected=913.695025
Week=7, Predicted=913.695025, Expected=971.727005
Week=8, Predicted=971.727005, Expected=1095.774093
Week=9, Predicted=1095.774093, Expected=1159.611429
Week=10, Predicted=1159.611429, Expected=732.743950
Week=11, Predicted=732.743950, Expected=963.615794
Week=12, Predicted=963.615794, Expected=896.842055
Week=13, Predicted=896.842055, Expected=960.615794

```

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Week=14, Predicted=960.630000, Expected=905.481048  
 Week=15, Predicted=905.481048, Expected=994.167244  
 Week=16, Predicted=994.167244, Expected=975.601111  
 Week=17, Predicted=975.601111, Expected=1210.307037  
 Week=18, Predicted=1210.307037, Expected=9293.130000  
 Week=19, Predicted=9293.130000, Expected=367.780000

I'm not differencing the data since my data is stationary, I don't see an increasing/decreasing pattern in the spending patterns. Can you help me out?



**Jason Brownlee** December 7, 2017 at 7:50 PM

REPLY ↗

It may suggest that the model is only learning from the input).

I would recommend starting with a simple MLP.



**Anton** December 17, 2017 at 2:04 am #

First of all, thanks for your great work and help.  
 Let's look at my changed csv data:

```
"1-01",0.1
"1-02",-0.2
"1-03",0.3
"1-04",-0.4
"1-05",0.5
"1-06",-0.6
"1-07",0.7
"1-08",-0.8
"1-09",0.9
"1-10",-1.0
"1-11",1.1
"1-12",-1.2
"2-01",1.3
"2-02",-1.4
"2-03",1.5
"2-04",-1.6
"2-05",1.7
"2-06",-1.8
"2-07",1.9
"2-08",-2.0
"2-09",2.1
"2-10",-2.2
"2-11",2.3
"2-12",-2.4
"3-01",2.5
```

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

[Start Machine Learning](#)

"3-02",-2.6  
 "3-03",2.7  
 "3-04",-2.8  
 "3-05",2.9  
 "3-06",-3.0  
 "3-07",3.1  
 "3-08",-3.2  
 "3-09",3.3  
 "3-10",-3.4  
 "3-11",3.5  
 "3-12",-3.6

Its a very simple sequence of “resonant” points, increasing every step over time. I didn’t change any changes in the program code, and here is the result:

Why LSTM can not predict so simple and logical series?

Waiting for your reply, and once again, thanks a lot!



**Jason Brownlee** December 17, 2017 at 8:30pm #

Nothing, LSTMS are poor at autoregression.

<https://machinelearningmastery.com/suitability-forecasting/>

Use an MLP instead.



## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

REPLY ↩



**Anton** December 17, 2017 at 2:21 am #

One interesting thing: I made the time serie 2 years longer, but with the same increasing resonant coefficient (0.1):

"1-01",0.1  
 "1-02",-0.2  
 "1-03",0.3  
 "1-04",-0.4  
 "1-05",0.5  
 "1-06",-0.6  
 "1-07",0.7  
 "1-08",-0.8  
 "1-09",0.9  
 "1-10",-1.0  
 "1-11",1.1  
 "1-12",-1.2  
 "2-01",1.3  
 "2-02",-1.4  
 "2-03",1.5  
 "2-04",-1.6  
 "2-05",1.7

Start Machine Learning

"2-06",-1.8  
"2-07",1.9  
"2-08",-2.0  
"2-09",2.1  
"2-10",-2.2  
"2-11",2.3  
"2-12",-2.4  
"3-01",2.5  
"3-02",-2.6  
"3-03",2.7  
"3-04",-2.8  
"3-05",2.9  
"3-06",-3.0  
"3-07",3.1  
"3-08",-3.2  
"3-09",3.3  
"3-10",-3.4  
"3-11",3.5  
"3-12",-3.6  
"4-01",3.7  
"4-02",-3.8  
"4-03",3.9  
"4-04",-4.0  
"4-05",4.1  
"4-06",-4.2  
"4-07",4.3  
"4-08",-4.4  
"4-09",4.5  
"4-10",-4.6  
"4-11",4.7  
"4-12",-4.8  
"5-01",4.9  
"5-02",-5.0  
"5-03",5.1  
"5-04",-5.2  
"5-05",5.3  
"5-06",-5.4  
"5-07",5.5  
"5-08",-5.6  
"5-09",5.7  
"5-10",-5.8  
"5-11",5.9  
"5-12",-6.0

Now, we see that LSTM made a predictions more closely, but still don't see the increment over Y axis (the red lines and arrows are my own painting) <http://prntscr.com/hob7eo>

Any ideas?

## Start Machine Learning X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

Start Machine Learning



**Jason Brownlee** December 17, 2017 at 8:53 am #

REPLY ↗

Not enough data, model needs tuning, lstms are poor at autoregression.



**Anton** December 17, 2017 at 3:30 am #

REPLY ↗

One more question about the code:

Why you performs 1 prediction before the predictions cycle:

```
train_reshaped = train_scaled[:, 0].reshape(len(train_scaled), 1, 1)
lstm_model.predict(train_reshaped, batch_size=1)
```

what is this for?

And at the line

```
make one-step forecast
X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
```

what is "y" for?



**Jason Brownlee** December 17, 2017 at 8:53 am #

It's unused.

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

REPLY ↗



**Krasimir** December 18, 2017 at 8:28 am #

REPLY ↗

Hi Jason,

Thanks for the great posts, very helpful and informative.

I am using some of your example for stock market values predictions. The models work well. Could you please provide more detail on why there is a one sample delay (lag) in the response of the model. It can be seen on all your plots and mine too.



**Jason Brownlee** December 18, 2017 at 3:26 pm #

REPLY ↗

Because the model is not very good, it's a sign that it is outputting the input as an output, called persistence:

<https://machinelearningmastery.com/persistence-time-series-forecasting-with-python/>



**Anton** December 20, 2017 at 5:53 am #

REPLY ↗

Is the model state updated after each

```
lstm_model.predict(train_reshaped, batch_size=1)
```

## Start Machine Learning

?

So can we be sure that after 100 predictions the model will have actual state for 101 prediction? Or should we use some additional model updates technique after each prediction step? It is not clear for me...



**Jason Brownlee** December 20, 2017 at 5:55 am #

REPLY ↗

State does not seem to matter that much in this example.



**Anton** December 20, 2017 at 6:13 am #

Agree, but in general? Does `lstm_model.predict`



**Jason Brownlee** December 20, 2017 at 3:34 pm #

It will update internal state variables, but not the input.

Does that help?



**Anton** December 21, 2017 at 8:34 am #

Yes, thank you very much

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** December 21, 2017 at 3:34 pm #

REPLY ↗

You're welcome.



**Novin** January 7, 2018 at 1:37 am #

REPLY ↗

Thanks Jason for the tutorial.

What's your suggestion for adapting this network for different stores shampoo sales? Things that I could think of are:

- 1- run this model for different stores separately (maybe we can learn something from sales of all stores together, but with this approach, we can't benefit that)
- 2- giving in the number of sales of each store as features

Start Machine Learning



**Jason Brownlee** January 7, 2018 at 5:10 am #

REPLY ↗

Some ideas:

- model each store separately
- model each store separately and add data from all stores
- model each store separately and add data from related stores (regionally? profile? etc.)
- create a regional store model and ensemble with separate store model.
- create all store model and ensemble with separate store model
- so on...

Think up 100 ideas and try them all to see what works best for your specific data.



**mostafa kotb** January 8, 2018 at 12:47 am #

thank you Dr. Jason for this tutorial.  
If i want to perform a walkforward validation on train

```
for i in range(len(train_scaled)):
 # make one-step forecast
 X, y = train_scaled[i, 0:-1], train_scaled[i, -1]
 yhat = forecast_lstm(lstm_model, 1, X)
 # invert scaling
 yhat = invert_scale(scaler, X, yhat)
 # invert differencing
 yhat = inverse_difference(raw_values, yhat, len(raw_values)-i)
 # store forecast
 predictions_train.append(yhat)
 expected = raw_values[i+1]
 print('Month=%d, Predicted=%f, Expected=%f' % (i+1, yhat, expected))

report performance
rmse_train = sqrt(mean_squared_error(raw_values[1:len(train_scaled)+1], predictions_train))

please, verify if it is right or not
```

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**  
Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** January 8, 2018 at 5:44 am #

REPLY ↗

Sorry, I cannot debug your code. You can learn more about walk-forward validation here:  
<https://machinelearningmastery.com/backtest-machine-learning-models-time-series-forecasting/>



**mostafa kotb** January 8, 2018 at 6:16 am #

REPLY ↗

thanks for your reply...

let me rephrase my question....i want to know where point in the raw values?

Start Machine Learning

```
rmse_train = sqrt(mean_squared_error(raw_values[1:len(train_scaled)+1], predictions_train))
```

if you notice in the last equation i started from the second point raw\_values[1:len(train\_scaled)+1]

because we don't predict the first point, actually we start by first point as x and predict the second point as y.



**Swapnil** January 8, 2018 at 3:53 pm #

REPLY ↗

Hi, Thanks for the great article.

I have some log data from my honeypot. I have encoded the actions taking place with label encoder. Is this data good enough for training the lstm? I finally to vector instead?



**Jason Brownlee** January 8, 2018 at 3:55 pm #

Perhaps try a few representations to see what works best.



**Jacques** January 9, 2018 at 11:37 am #

Hey, I'm trying to understand why you implemented the problem as a supervised learning task compared to your other post (<https://machinelearningmastery.com/time-series-forecasting-long-short-term-memory-network-python/>). In your other post, you framed the problem as a time series forecasting problem. Why didn't you do the same here? You also talk about t-1 and t+1 on that post, not t-1 and t. Could you clarify why you didn't convert the problem in this post into supervised until the lstm model?



**Jason Brownlee** January 9, 2018 at 3:18 pm #

REPLY ↗

I was trying to make the example simpler.



**Swapnil** January 9, 2018 at 6:24 pm #

REPLY ↗

Hi how do u judge efficacy with rmse? Less is better or more?



**Jason Brownlee** January 10, 2018 at 5:22 am #

REPLY ↗

Small RMSE is better.

[Start Machine Learning](#)



Ian January 23, 2018 at 7:02 pm #

REPLY ↗

Hi Jason, nice tutorial. I am using XBT data for 2017, but just trying daily. When I follow this suggestion "Update: Consider trying 1500 epochs and 1 neuron, the performance may be better!" the RMSE and the graph look very similar to the persistent model. Was that your finding?



Jason Brownlee January 24, 2018 at 9:52 am #

REPLY ↗

I don't recall sorry.



andr January 24, 2018 at 12:44 am #

Hi,  
is it possible to append to training set some data after  
with the fit function and the updated training set as



Jason Brownlee January 24, 2018 at 9:55 am #

You can. Try it and see.

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about  
services and special offers by email. For more  
information, see the [Privacy Policy](#).

START MY EMAIL COURSE



andr January 26, 2018 at 1:30 am #

REPLY ↗

I've tried two experiments:

- 1) update the training set with new data as shown in <https://machinelearningmastery.com/update-lstm-networks-training-time-series-forecasting/> and update the model with 1 epoch
- 2) same as point one but when I append the new data at the end of training set I remove the first element in it (a sort of sliding window. Does it have sense?).

The rmse in the two experiments is always worse than the version with no update in the training set. Any suggestion?



Jason Brownlee January 26, 2018 at 5:45 am #

REPLY ↗

Interesting, nice work.

If you are looking for further improvements, I would suggest brainstorming 10 different approaches and compare them.

After that you can then think about more general things:

<http://machinelearningmastery.com/improve-deep-learning-performance/>

Start Machine Learning



**Antonio Bautista** February 1, 2018 at 4:27 am #

REPLY ↗

Hello, thanks for this tutorial it is great and gave me a lot of insights. I just have one question.

Is the historial data still needed after the training of the net ? or just the last observations are needed?

Thank you!



**Jason Brownlee** February 1, 2018 at 7:26 am #

REPLY ↗

You model requires inputs to make predictions. It depends on how you have defined your model as to what you will need to retain.

Generally, training data is only used to fit the model to new obs.



**Shotaro** February 2, 2018 at 7:36 pm #

Hi, Jason

Which type LSTM method do you use in this sample?

1:[Hochreiter & Schmidhuber, 95;97] Simple LSTM

2:[Gers & Schmidhuber, 99] LSTM with Forget Gate

3:[Gers & Schmidhuber, 00] LSTM with Peephole Connection

4:[Graves & Schmidhuber, 05] LSTM Full Gradient

5: Other

I think that Simpe LSTM. Will you show other model in the future?

## Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this free and practical course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Clock ZHONG** February 11, 2018 at 4:51 pm #

REPLY ↗

Jason,

Thanks again for you this post. I've also followed your instructions in this post, and finished the test.

But I've a question:

Comparing with other models, training this RNN model is very slow, I checked the code and found that you set the batch\_size as 1, I believe one solution to accelerate the training is increasing batch\_size to a bigger number, but I fear it will cause some other errors, as you explain as following:

"The batch\_size must be set to 1. This is because it must be a factor of the size of the training and test datasets."

Does i mean if we set the batch\_size as a common factor of both the size of training and the testing datasets, then it'll be OK? E.g. if my testset size is 1000, and my training dataset is 100000, then I could set the batch\_size as 1000, is my understanding right?

If my understanding is wrong, what's other restriction?

Start Machine Learning

Thanks in advance.

Clock ZHONG



**Clock ZHONG** February 11, 2018 at 5:07 pm #

REPLY ↗

Jason,

I read explanation in Keras website:

<https://keras.io/getting-started/faq/#how-can-i-use-stateful-rnns>

It seems if we could ensure the samples sequence in a training batch, then the batch size could be set as any number.

And if the test dataset follows the training data state reset, I believe so.

Thanks!

Clock ZHONG



**Jason Brownlee** February 12, 2018 at 8:28 am #

Sure, try it.



**Ridhima Kumar** February 11, 2018 at 6:44 pm ..

Hi Jason,

Thanks for the wonderful explanation on LSTM.

I have tried running the code. However, I am getting an error – Keyword Error 1 when I difference the series.

The error comes after running the following piece:

```
transform to be stationary
differenced = difference(series, 1)
```

It would be great if you could help me out with this.

Thanks!

Ridhima Kumar



**Jason Brownlee** February 12, 2018 at 8:28 am #

REPLY ↗

Perhaps double check that you have copied all of the code?

[Start Machine Learning](#)



**Peter** February 17, 2018 at 3:50 pm #

REPLY ↗

Hi Jason,

Thanks for the Awesome tutorial. Sorry if I'm stupid asking this here, New to application level ML. Only have amateur knowledge in the area, trying to learn by building a web app.

I'm using a sample data of three years to forecast the price for next 3 years- made up my dataset exactly like the sample used here, I followed your tutorial and everything worked fine till the last steps (lol I know how to copy-paste)

But I'm now stuck at writing the code for forecasting the next 3 years price.



**Jason Brownlee** February 18, 2018 at 6:

What is the error for the last step?



**Peter** February 19, 2018 at 3:22 pm #

I'm not getting any errors, but I have values for the year 2030.



**Peter** February 19, 2018 at 3:54 pm #

Simply: I want to forecast the 4th 5th and 6th years using this model, so how do I modify my code?



**Jason Brownlee** February 21, 2018 at 6:22 am #

REPLY ↗

Ho Peter, this post will show you how to do a multi-step forecast:

<https://machinelearningmastery.com/multi-step-time-series-forecasting-long-short-term-memory-networks-python/>



**Julian** February 21, 2018 at 3:54 am #

REPLY ↗

Hi Jason,

thank you for your tutorial. I have a question for you. I copied your source code and changed the parameters for the number of epochs and units (LSTM) to 10 and 1. With this configuration I achieved very good results. If I now comment out the lines that are responsible for training the neural network, I get the same results. How is that possible?

Is there a bug in your program or do I have problem

Start Machine Learning



**Jason Brownlee** February 21, 2018 at 6:41 am #

REPLY ↗

Perhaps you did not comment out enough?



**Julian** February 21, 2018 at 8:16 am #

REPLY ↗

I have commented out the following lines:

```
for i in range(nb_epoch):
model.fit(X, y, epochs=1, batch_size=batch_size)
model.reset_states()
```

and

```
train_reshaped = train_scaled[:, 0].reshape(-1, 1)
lstm_model.predict(train_reshaped, batch_size=1)
```

I mean, that should be enough. Or did I for-



**Jason Brownlee** February 22,

Looks reasonable.

## Start Machine Learning

X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Paula** February 22, 2018 at 7:07 am #

REPLY ↗

Hi Jason,

I'm new in LSTM nn. I copied your example, run it and it works but I don't understand why I get different results each time I run the code. I didn't find any random number that could make to get all time different results.

Moreover, I would like to know the reason why you create and fit the model 30 times and not only one. Could I take the best result of those 30?

Thanks!



**Jason Brownlee** February 22, 2018 at 11:23 am #

REPLY ↗

Great question, see this post:

<https://machinelearningmastery.com/randomness-in-machine-learning/>

**Ian Zhang** March 6, 2018 at 9:47 pm #

## Start Machine Learning



Hi Jason,

i'm trying to deal with my time prediction problem using LSTM-based architecture. There are almost 30 id in my dataset. Each data of id is similar with your sample.(time related data.)Here is my question:

1.Can this model train multiple sets of data at the same time?(e.g. you want to deal with sales of different brands of shampoo. And you want to predict each of sale at same time.Maybe each set of data is relevant, but you dont care.)If i can, how to decide input shape and output shape?Is there any influence in the network of each set data

2.If my data is stationary ,do i need to process my data to train my data without difference using this architecture(looks like a horizontal line).I don't know why it happens designed for this method(difference)?

Thanks!



**Jason Brownlee** March 7, 2018 at 6:13 am #

Sure, try it and see how you go. Perhaps

If the data is stationary, no need to use a differ-



## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address



I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**



**Jamie** March 10, 2018 at 12:32 am #

REPLY ↗

Hi Jason! I have been reading your blog as a reference for a project at my internship and your content is awesome!!

In my project I need to both do series forecasting and also sequence classification. So I am playing with 2 models as my first approach to the problem: one LSTM network for forecasting and a second model for which I am considering a sliding window MLP to take small patches of my time series and do sequence classification (I have labels for all time steps).

So the general idea is to classify forecast sequences (classification over my predictions). I hope it makes sense.

In case it does, I wanted to ask for advice on the sliding window MLP since I really can't find good references on the web. What would I do in general? Would I stack the inputs from the previous time steps and give it to my network as input? So an example for simplicity: if I have 5 variables, and I want to have a 5-step sliding window, would my input for time t be  $[[t-4], [t-3], [t-2], [t-1], [t]] \rightarrow y[t]$  and let my network figure out how to learn from it? or would I have to instead take the 5 vectors and aggregate them somehow (e.g. average) and then give my network only 1 resulting vector and its label? I am sorry if my question seems dumb but I really have not been able to find useful information on this topic, if you have anything on this on any of your posts or resources please let me know.

Thanks!

**Start Machine Learning**



**Jason Brownlee** March 10, 2018 at 6:32 am #

REPLY ↗

Yes, in general the lag obs become a window of features as input for the MLP.

Also, consider LSTM for sequence classification.

Also, explore a multiple output model to do both predictions from one model to see if it offers a lift in skill (if I understand your data correctly).

Let me know how you go.



**Jamie** March 12, 2018 at 8:27 pm #

Thank you Jason, I will definitely explore the multiple output model and also put it simply to see if the third option is better. I have an illustration that is not my real problem but it may help with the illustration.

Imagine you had weather data and you had multiple sensors. All the data so all that matters to you is simply saying Can I predict the weather? idea. You have historical data and multiple sensors (temperature, humidity, etc.), and you want to give recommendations (rain, hours, etc.). At every time step, you also have a recommendation.

So all that matters would be recommending rain or sun. I am not sure if I would be able to do this if not with 2 individual models, one that forecasts the values for temperature, humidity, etc and a second model that does sequence classification and can classify the sequences I predicted with my first model?

Is there a way I could simply do my recommendation forecasts with only 1 model? What type of model would this be (if any)? Thank you Jason!

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** March 13, 2018 at 6:26 am #

REPLY ↗

Yes, you could do a multi-step forecast. Here are some examples of the different ways you could structure this model:

<https://machinelearningmastery.com/multi-step-time-series-forecasting/>



**Nikhil Singh** March 12, 2018 at 10:06 am #

REPLY ↗

Hi Jason,

Thanks for the tutorial.

I have used your code and made some modifications like I have increased the number of layers to 4 and add neurons (6,12, 18 and 1) respectively at each layer.

[Start Machine Learning](#)

function and regularization also but I when I plotted the graph, the output is lag by 1 to input, like persistence model.

When I shift down the output value and again plot it nearly sometimes almost matches the input.

For example, if the input has 10 value and output also have 10 value. I got better graph when I plot input value from 2 to 10 and the output value from 1 to 9. So I ignored the 1st value of the input and last value of the output.

I am trying to match the graph but I feel that I am wrong here because I don't find any reason.

Could you please help me why I am getting this result? Means, even after using complex model, my model behaves like persistence model.

Please help !!!

I will be eagerly waiting for your reply

Thanks



**Jason Brownlee** March 12, 2018 at 2:25

Perhaps the model is a bad fit for the



**Atreya Bandyopadhyay** March 17, 2018 at

"The predict() function on the model is also mean? If i set a batch size of 4 then will I have to p

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** March 18, 2018 at 6:00 am #

REPLY ↗

Correct.



**Hervaldo** March 20, 2018 at 1:14 am #

REPLY ↗

Hi Jason. I got this error and I could not discover what is wrong

ValueError: time data '1901-Jan' does not match format '%Y-%m'



**Jason Brownlee** March 20, 2018 at 6:24 am #

REPLY ↗

Looks like there is an issue with your data.

Try downloading it from here:

<https://raw.githubusercontent.com/jbrownlee/Datasets/master/shampoo.csv>

Start Machine Learning



**Hervaldo** March 20, 2018 at 1:17 am #

REPLY ↗

I also looked at the dataset and the value of the first index is 1/1/2018

Thanks,

Hervaldo



**Kevin** March 27, 2018 at 9:39 am #

REPLY ↗

I think I am missing something basic here, I am new at this. How would you use the model to predict what would occur past the test data, i.e. Mo



**Jason Brownlee** March 27, 2018 at 4:18 pm #

Great question, yes you are missing a key part of evaluating an LSTM. See this post:

<https://machinelearningmastery.com/make-predictions-lstm-time-series-data/>



**Don Lilly** March 29, 2018 at 5:41 am #

Hi Jason,

Thanks for all the great articles. I have a situation with 18 input variables and 1 output variable (which can be -1, 0 or 1) and that I convert with numpy's to\_categorical function to generate 3 outputs of (1, 0, 0), (0, 1, 0), and (0, 0, 1). I trained the model and it reports high accuracy (it starts out low and gets higher and higher as the epochs go by). So far so good.

All 18 input variables are floating point and scaled to be between -1 and 1.

The problem enters when I use the predict function on my test data – my answers are all values around 0 (floating point, like (0.0001, 0.002, 0.00123). The shape of the training and test data is the same, as reported by model.summary(). It is as if the model hasn't learned anything.

Are the outputs from the predict function scaled or transformed in some way? I don't need them to be, but it could be I am misunderstanding something here (an understatement!).

For reference, my code is roughly (there are 5766 rows of data):

```
model = Sequential()
model.add(LSTM(18, batch_input_shape=(1, 5766, 18), stateful=True, return_sequences=True))
model.add(Dense(10, input_shape=(18,), kernel_initializer=initializers.random_normal(seed=51)))
model.add(Dense(3, kernel_initializer='uniform'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
X_train = np.reshape(X, (1, X.shape[0], X.shape[1]))
Y_train = np.reshape(Y_train, (1, Y_train.shape[0], Y_train.shape[1]))

for i in range(25):
 model.fit(X_train, Y_train, epochs=1, batch_size=1)
```

[Start Machine Learning](#)

```
model.reset_states()
```

... then reshape the x test array in similar fashion and call:

```
y_pred = model.predict(X_test)
```

Thanks,

Don



**Jason Brownlee** March 29, 2018 at 6:41 am #

REPLY ↗

Perhaps there is a bug? Perhaps you need to tune the model some more?

I give some ideas here:

<http://machinelearningmastery.com/improve-deep-learning-models/>



**Don Lilly** March 29, 2018 at 8:13 am #

Thanks for your response. I will continue to

Don



**Fredrik Nilsson** April 3, 2018 at 1:41 am #

Hi

Very good writing Jason!

I just have one question about the prediction which is a bit unclear.

I tried your exemplel of using `yhat = lstm_model.predict(X)` but it only give me the error:

`ValueError: Error when checking : expected lstm_2_input to have 3 dimensions, but got array with shape (1, 1)`

For exampel how can we make an prediction for the period "4-01" or even "4-02" that does not exist in the trainingdata?

Since so many are asking for this maybe its time to put it in the main text?

Thanks for great job!

//Fredrik Nilsson

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** April 3, 2018 at 6:38 am #

REPLY ↗

Sorry to hear that, did you copy all of the code from the example?

Are you using the same dataset?

Is your environment (Python/Keras) up to date?

Start Machine Learning



**frank gallagher** April 3, 2018 at 5:41 pm #

REPLY ↗

the examples is great. I wonder if the neutral network is too simple with only one lstm layer and a Dense. Would it do better changing the model? Like adding more layers.



**Jason Brownlee** April 4, 2018 at 6:09 am #

REPLY ↗

Try it and see.



**Akira** April 4, 2018 at 6:32 pm #

Sir , i got a question in the codes:

```
make one-step forecast
X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
yhat = forecast_lstm(lstm_model, 1, X)
invert scaling
yhat = invert_scale(scaler, X, yhat)

which means:
yhat = yhat + raw_values [-(len+1-i)]
```

We use observation to add yhat each time? How for?

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)



**Jason Brownlee** April 5, 2018 at 5:54 am #

REPLY ↗

This is called walk-forward validation, you can learn more here:

<https://machinelearningmastery.com/backtest-machine-learning-models-time-series-forecasting/>



**VanDarkholme** April 11, 2018 at 7:45 pm #

REPLY ↗

Hello Jason

Is there any difference between:

```
for i in range(nb_epoch):
 model.fit(X, y, epochs=1, batch_size=batch_size, verbose=0, shuffle=False)
 model.reset_states()
```

and

```
model.fit(X, y, epochs=nb_epoch, batch_size=batch_size, verbose=0, shuffle=False) ?
```

[Start Machine Learning](#)



**Jason Brownlee** April 12, 2018 at 8:39 am #

REPLY ↗

Yes, we have control over when the state is reset.

In the first case, it is reset at the end of each epoch. In the second case it is reset at the end of each batch.



**Gustav** April 18, 2018 at 10:46 pm #

REPLY ↗

How would this example work with a dataset that is stockprices.

I have a csv file that is basically:

"DateTime","Price [USD]"  
 "2015-08-07 00:00:00",2.77  
 "2015-08-08 00:00:00",0.8077  
 .  
 .  
 .

Could I use the base from your code above for exp



**Jason Brownlee** April 19, 2018 at 6:32 am #

REPLY ↗

I had advice for preparing data for LSTM

<https://machinelearningmastery.com/faq/single-faq/how-do-i-prepare-my-data-for-an-lstm>

I have advice on forecasting the stock market here:

<https://machinelearningmastery.com/faq/single-faq/can-you-help-me-with-machine-learning-for-finance-or-the-stock-market>



**Jason** April 29, 2018 at 6:39 pm #

REPLY ↗

I have a general question about predict function.

Say I have trained a LSTM model (one step prediction e.g. X is 1D, Y is 1D), and I want to predict some Ys for given future Xs (  $X=(X_1, X_2, X_3, X_4, X_5, X_6, X_7) = (0.1, 0.2, 0.1, 0.5, 0.5, 0.9, 0.3)$  )

I know I will get a set of Y (  $Y = (Y_1, Y_2, Y_3, Y_4, Y_5, Y_6, Y_7)$  )

Since  $X_1=X_3=0.1$ ,  $X_4=X_5=0.5$ , does it mean  $Y_1=Y_3$  and  $Y_4=Y_5$  ?

My thought is for a given model, the coefficients are fixed after trained, when input a same value (e.g. value=0.1), it will output(predict) a fixed value.

If not the case, could you briefly tell me how does the predict function work in the above example?

Thanks!

**Jason Brownlee** April 30, 2018 at 5:33 am #

[Start Machine Learning](#)



The LSTM weights are fixed after training, but the LSTM also has an internal state, like a local variable that is accumulated over the input sequence.

You can learn more about how LSTMs work here:

<https://machinelearningmastery.com/start-here/#lstm>



**Ting** May 10, 2018 at 6:32 pm #

REPLY ↗

Thanks for your sharing.

It is very helpful for a rookie try to handle time series data.

Below please find question I faced when reproducing

ValueError: Error when checking : expected lstm\_2  
(1, 0)

(more info on url)

My environment is:

Windows 10

Python 3



**Jason Brownlee** May 11, 2018 at 6:35 am #

X

I'm sorry to hear that, here are some suggestions:  
<https://machinelearningmastery.com/faq/single-predictions/>

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Sanjoy Datta** May 14, 2018 at 1:41 am #

REPLY ↗

Hi Jason,

Struggling to get shampoo data in. Getting the following errors:

TypeError: strptime() argument 1 must be str, not numpy.ndarray

ValueError: time data '190'1-01",266.0' does not match format '%Y-%m'

Input data is a comma delimited csv file with no footer:

```
"Month","Sales"
"1-01",266.0
"1-02",145.9
"1-03",183.1
"1-04",119.3
"1-05",180.3
"1-06",168.5
"1-07",231.8
"1-08",224.5
```

Start Machine Learning

"1-09",192.8  
 "1-10",122.9  
 "1-11",336.5  
 "1-12",185.9  
 "2-01",194.3  
 "2-02",149.5  
 "2-03",210.1  
 "2-04",273.3  
 "2-05",191.4  
 "2-06",287.0  
 "2-07",226.0  
 "2-08",303.6  
 "2-09",289.9  
 "2-10",421.6  
 "2-11",264.5  
 "2-12",342.3  
 "3-01",339.7  
 "3-02",440.4  
 "3-03",315.9  
 "3-04",439.3  
 "3-05",401.3  
 "3-06",437.4  
 "3-07",575.5  
 "3-08",407.6  
 "3-09",682.0  
 "3-10",475.3  
 "3-11",581.3  
 "3-12",646.9

Would be grateful for your help

## Start Machine Learning X

You can master applied Machine Learning **without math or fancy degrees.**  
 Find out how in this *free* and *practical* course.

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)



**Jason Brownlee** May 14, 2018 at 6:38 am #

REPLY ↗

Here is the data file:

<https://raw.githubusercontent.com/brownlee/Datasets/master/shampoo.csv>

If you are still having trouble, perhaps double check you copied the code exactly, with white space.

Here are some more suggestions:

<https://machinelearningmastery.com/faq/single-faq/why-does-the-code-in-the-tutorial-not-work-for-me>



**Kevin** May 22, 2018 at 6:50 pm #

REPLY ↗

You can check your x.csv file, It's best to change the format of column month to short date, just like this:

Month Sales

Start Machine Learning

2018/1/1 266  
 2018/1/2 145.9  
 2018/1/3 183.1  
 2018/1/4 119.3  
 2018/1/5 180.3



**Hirotaka Nakagame** May 17, 2018 at 6:33 am #

REPLY ↗

It looks like my result is shifted by 1 timestep into the future but I am not sure where it came from.



**Jason Brownlee** May 17, 2018 at 6:41 am #

This is common, see my explanation [here](https://machinelearningmastery.com/faq/single-faq/why-is-my-forecasted-time-series-right-behind-the-actual-time-series)



**Hirotaka Nakagame** May 23, 2018 at 6:18 pm #

do you have links to further tuning good 'shifted' forecast? Any math reasons

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** May 23, 2018 at 6:31 am #

REPLY ↗

Yes, see here:

<https://machinelearningmastery.com/faq/single-faq/why-is-my-forecasted-time-series-right-behind-the-actual-time-series>



**kevin** May 22, 2018 at 6:46 pm #

REPLY ↗

Hi Jason, Thanks for your great article, May I ask some questions?

What if I wanna predict different products' sales?  
 For example, shampoo , cellphone , T-shirt, etc.  
 Should I predict each items independently or together?



**Jason Brownlee** May 23, 2018 at 6:25 am #

REPLY ↗

Great question.

Start Machine Learning

It depends on the data. It might be better to model each separately, to model them in groups or to model all products together. Some experimentation is required.



**Chirag mandot** May 24, 2018 at 1:08 am #

REPLY ↗

Is it a good idea to build 10k different models for all products? Is there a better approach that you have used before for such scenarios?

Thank you

Great article Jason



**Jason Brownlee** May 24, 2018

Yes, models for categories of all products.



**Hiro** May 25, 2018 at 7:20 am #

If I need to take a second difference, is thi

```
first_diff = difference(raw_values, 1)
second_diff = difference(first_diff, 1)
first_diff.index = raw_values.index[1:]
```

```
predictions = []
temp = []
for i in range(len(test_scaled)):
 X, y = test_scaled[i, 0:-1], test_scaled[i, -1]

 yhat = forecast_Istm(Istm_model, 1, X)
 yhat = invert_scale(scaler, X, yhat)
 temp.append(yhat)
 yhat = inverse_difference(first_diff, yhat, len(test_scaled)+1-i)
 yhat = inverse_difference(raw_values, yhat, len(test_scaled)+12-i)
 predictions.append(yhat)
```

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** May 25, 2018 at 9:35 am #

REPLY ↗

You can difference the differenced series or use an ARIMA model that can do differencing for you at different orders.

Start Machine Learning



WEIKE May 31, 2018 at 3:55 am #

REPLY ↗

Dear Jason:

I am a little confuse about why `yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)` instead of `yhat = inverse_difference(raw_values, yhat, len(test_scaled)-i)`. I am not sure why we need the 1.

I am a new man to access the LSTM.



Leon June 18, 2020 at 3:49 am #

REPLY ↗

Hi Jason,

In the section where you show us how to difference this line code:

```
inverse_difference(series, differenced[i], len(series) - i)
```

And in the example you use this line of code:

```
yhat = inverse_difference(raw_values, yhat, len(raw_values) - 1)
```

What would be the reason to add "1" to the code?



Jason Brownlee June 18, 2020 at 6:14 am #

REPLY ↗

I believe it is required to align the `[1,2,3,4,5,6,7,8,9,10]` and inspect the results.



Ger June 7, 2018 at 11:03 pm #

REPLY ↗

Hey Jason!

I am working with your example of the LSTM.

Since the data is transformed to a supervised learning problem ( $t-1$ ) shouldn't we shift the predictions by -1 as well? Right now my graph seems to line up almost exactly with the real data but shifted one day ahead.

Thanks!



Jason Brownlee June 8, 2018 at 6:14 am #

REPLY ↗

Perhaps this post will help better explain how to prepare time series data for supervised learning:

<https://machinelearningmastery.com/convert-time-series-supervised-learning-problem-python/>

[Start Machine Learning](#)



Ari June 9, 2018 at 1:24 am #

REPLY ↗

Hi,

I was getting errors about expecting 2D but being 1D so I changed scaler to be:  
 scaler = scaler.fit([train]) # <—Different from tutorial  
 i.e. adding brackets

That fixed the problem but now I get "IndexError: tuple index out of range" referencing this line:  
 train = train.reshape(train.shape[0], train.shape[1])

The code other than that one change is the same. Do you know why this might be? I'm using python 3.5.5

Thank you for any help!

```
Shampoo sales dataset
LSTM
1. Transform the time series into a supervised learning problem
2. Transform the time series data so that it is stationary
3. Transform the observations to have a specific scale

Import packages
from pandas import DataFrame
from pandas import Series
from pandas import concat
from pandas import read_csv
from pandas import datetime
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import MinMaxScaler
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import LSTM
from math import sqrt
from matplotlib import pyplot
import numpy

date-time parsing function for loading the dataset
def parser(x):
 return datetime.strptime('190'+x, '%Y-%m')

Transform time series to supervised learning
Shift all values down to create t-1 input and t output
Helper function time_series_to_supervised()

Test the function with loaded sales data & convert to supervised learning problem
Frame a sequence as a supervised learning problem
def timeseries_to_supervised(data, lag=1):
 df = DataFrame(data)
 columns = [df.shift(i) for i in range(1, lag+1)]
 columns.append(df)
 df = concat(columns, axis=1)
 df.fillna(0, inplace=True)
 return df
```

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees.**

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

[Start Machine Learning](#)

```

Transform Time Series to Stationary (remove trend)
Stationary = Not dependent on time
Method = Differencing the data; observation from t-1 is subtracted from t, only leaves
Create a differenced series
def difference(dataset, interval=1):
 diff = list()
 for i in range(interval, len(dataset)):
 value = dataset[i] - dataset[i - interval]
 diff.append(value)
 return Series(diff)

Invert differenced value
Necessary to take forecasts made on the differenced series back into their original scale
def inverse_difference(history, yhat, interval=1):
 return yhat + history[-interval]

Scale train and test data to [-1, 1]
def scaler(train, test):
 # Fit scaler
 scaler = MinMaxScaler(feature_range=(-1, 1))
 scaler = scaler.fit([train]) # Supervised learning form
 def fit_lstm(train, batch_size, nb_epoch, neurons):
 # Reshape data from X, y to Samples/TimeSteps/F
 X, y = train[:, 0:-1], train[:, -1]
 X = X.reshape(X.shape[0], 1, X.shape[1])
 # Specify network -> Compile in efficient symbolic
 # Specify loss function and optimization algorithm
 model = Sequential()
 model.add(LSTM(neurons, batch_input_shape=(batch_size, X.shape[1], X.shape[2]), stateful=True))
 model.add(Dense(1))
 model.compile(loss='mean_squared_error', optimizer='adam')
 # Once compiled, fit to training data
 # Default in epoch = shuffled = not ideal for LSTM- SET TO FALSE!
 # Verbose = Reporting of debug info
 # Loop manually fits the network to training data
 for i in range(nb_epoch):
 model.fit(X, y, epochs=1, batch_size=batch_size, verbose=0, shuffle=False)
 model.reset_states()
 return model

 # Fixed approach = Fit the model once on all of the training data
 # then predict each new time step one at a time from the test data (used here).
 # Dynamic approach = Re-fit the model or update the model each time
 # step of the rest data as new obs from the test data are made available.
 # Given a fit model, a batch size used when fitting the model (1), and a row
 # from the test data, the function will separate out the input data from the
 # test row, reshape it, and return the prediction as a single floating point value.

 # Make a one-step forecast
 def forecast(model, batch_size, row):
 X = row[0: -1]

```

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

```

X = X.reshape(1, 1, len(X))
yhat = model.predict(X, batch_size=batch_size)
return yhat[0,0]

#Load dataset w/ earlier parser function
series = read_csv('/Users/arianalemadriscoll/Documents/SeasonalityData/shampoo-sales.csv',
header=0, parse_dates=[0], index_col=0, squeeze=True, date_parser=parser, engine='python',
skipfooter=3)

Transform to be stationary
raw_values = series.values
diff_values = difference(raw_values, 1)

Transform to supervised learning
supervised = timeseries_to_supervised(diff_values)
supervised_values = supervised.values

Split data into train and test sets
train, test = supervised_values[0:-12], supervised_values[-12:]

Transform the scale of the data
scaler, train_scaled, test_scaled = scaler(train, test)

Fit the model
lstm_model = fit_lstm(train_scaled, 1, 3000, 4)
Also try 1, 1500, 2?

Forecast the entire training dataset to build up state
train_reshaped = train_scaled[:, 0].reshape(len(train_scaled), 1, 1)
lstm_model.predict(train_reshaped, batch_size=1)

Walk-forward validation on the test data (rolling forecast scenario)
predictions = list()
for i in range(len(test_scaled)):
 # Make one-step forecast
 X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
 yhat = forecast_lstm(lstm_model, 1, X)
 # Invert scaling
 yhat = invert_scale(scaler, X, yhat)
 # Invert differencing
 yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)
 # Store forecast
 predictions.append(yhat)
 expected = raw_values[len(train) + i + 1]
 print('Month=%d, Predicted=%f, Expected=%f' % (i+1, yhat, expected))

Report performance
RMSE punishes large errors and the score reports in the same units as the forecast data (monthly
shampoo sales)

rmse = sqrt(mean_squared_error(raw_values[-12:], predictions))
print('Test RMSE: %.3f' % rmse)

#Line plot of observed vs. predicted
pyplot.plot(raw_values[-12:])

```

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

[Start Machine Learning](#)

```
pyplot.plot(predictions)
pyplot.show()

batch_input_shape = tuple that specifies the expected # obs to read each batch
+ # time steps, # features
Batch size and epochs define how quickly the network learns the data
i.e. how often the weights are updated
Number of neurons/memory units/blocks (1-5 for simple problem is sufficient)

#layer = LSTM(neurons, batch_input_shape=(batch_size, X.shape[1], X.shape[2]), stateful=True)
```



**Jason Brownlee** June 9, 2018 at 6:55 am

This is a common question that I answer in my [FAQ](https://machinelearningmastery.com/faq/single/machine-learning-faq/).



**Rafael Caballero** June 12, 2018 at 2:15 am #

I think there is a bug in your proposal. Just find that your model still predicts the right values!! A

The error, I think is in the inverse\_difference function. You are incrementing the raw value by the previous raw value. But you are ne

That is, the last point you can 'see' is the last value you predict an increment  $y_{hat0}$  the predicted value  $y_{hat1}$  from  $y_{hat0}$ . Then the predicted value is  $v + y_{hat0} + y_{hat1}$ , but you are using  $raw\_values[train\_size+1]+y_{hat1}$ , which is wrong.

I think the same happens in other of your posts about LSTM.

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** June 12, 2018 at 6:46 am #

REPLY ↗

The model is not skilful. Don't use LSTMs for time series forecasting:

<https://machinelearningmastery.com/suitability-long-short-term-memory-networks-time-series-forecasting/>



**Rafael Caballero** June 12, 2018 at 8:23 am #

REPLY ↗

I see, thanks for the reference. However, I have found that LSTM can beat ARIMA in very short time predictions, just in the first units of time.

**Jason Brownlee** June 12, 2018 at 9:00 am #

Start Machine Learning



Really!?

I have not seen that myself.

I find that a grid search'ed SARIMA or Holt-Winters model outperforms a neural net on every univariate time series dataset I try.



**Aditya** June 20, 2018 at 9:25 am #

REPLY ↗

while\_loop() got an unexpected keyword argument 'maximum\_iterations' . Has anyone encountered this error? I believe this is an error due to Keras/ Tensorflow version?



**Jason Brownlee** June 21, 2018 at 5:58 a

I have not seen this error sorry. Perha



**Mario** June 27, 2018 at 1:39 am #

I have the same error, i thought the problem persist.

```
TypeError Traceback (most recent call last)
in ()
1 # fit the model
--> 2 lstm_model = fit_lstm(train_scaled, 1, 3000, 4)
3 # forecast the entire training dataset to build up state for forecasting
4 train_reshaped = train_scaled[:, 0].reshape(len(train_scaled), 1, 1)
5 lstm_model.predict(train_reshaped, batch_size=1)

in fit_lstm(train, batch_size, nb_epoch, neurons)
50 X = X.reshape(X.shape[0], 1, X.shape[1])
51 model = Sequential()
--> 52 model.add(LSTM(neurons, batch_input_shape=(batch_size, X.shape[1], X.shape[2]),
stateful=True))
53 model.add(Dense(1))
54 model.compile(loss='mean_squared_error', optimizer='adam')

C:\Users\mmariscal\Anaconda3\lib\site-packages\keras\engine\sequential.py in add(self, layer)
C:\Users\mmariscal\Anaconda3\lib\site-packages\keras\layers\recurrent.py in __call__(self, inputs,
initial_state, constants, **kwargs)
498
499 if initial_state is None and constants is None:
-> 500 return super(RNN, self).__call__(inputs, **kwargs)
501
502 # If any of initial_state or constants are specified and are Keras
```

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

```
C:\Users\mmariscal\Anaconda3\lib\site-packages\keras\engine\base_layer.py in __call__(self, inputs, **kwargs)
```

```
C:\Users\mmariscal\Anaconda3\lib\site-packages\keras\layers\recurrent.py in call(self, inputs, mask, training, initial_state)
```

```
2110 mask=mask,
```

```
2111 training=training,
```

```
-> 2112 initial_state=initial_state)
```

```
2113
```

```
2114 @property
```

```
C:\Users\mmariscal\Anaconda3\lib\site-packages\keras\layers\recurrent.py in call(self, inputs, mask, training, initial_state, constants)
```

```
607 mask=mask,
```

```
608 unroll=self.unroll,
```

```
-> 609 input_length=timesteps)
```

```
610 if self.stateful:
```

```
611 updates = []
```

```
C:\Users\mmariscal\Anaconda3\lib\site-packages\keras\l
```

```
rnn(step_function, inputs, initial_states, go_bac
```

```
2955 # Arguments
```

```
2956 x: A tensor or variable.
```

```
-> 2957 axis: The dimension softmax would be
```

```
2958 The default is -1 which indicates the last
```

```
2959
```

TypeError: while\_loop() got an unexpected key

## Start Machine Learning



You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)



**Jason Brownlee** June 27, 2018 at 8:20 am #

REPLY ↗

Are you able to confirm that you are using the latest version of Keras and TensorFlow?

For example:

tensorflow: 1.8.0

keras: 2.1.6



**Pei** June 28, 2018 at 6:59 am #

REPLY ↗

I got the same error with tensorflow=1.3.1, keras=2.2.0.

After I upgraded to tensorflow=1.8.0, (keras=2.2.0), the error was gone!

So the error is caused by lower versions of tensorflow.



**Jason Brownlee** June 28, 2018 at 2:02 pm #

[Start Machine Learning](#)

I'm glad to hear that!



**Sonya** June 22, 2018 at 1:33 am #

REPLY ↩

Dear Jason,

I have a multidimensional data with 24 features.

I would like to convert my data into batches of lets say size 100, so that's 100 rows 24 columns: (100,24). Then I would like to predict the next batch based on the lets say previous 20. Is this possible, how can I achieve this?

Thanks for your help.



**Jason Brownlee** June 22, 2018 at 6:14 am #

This is a common question that I answer in my course:

<https://machinelearningmastery.com/faq/single-faq/>



**Fatima** July 5, 2018 at 9:05 pm #

Hi Jason,

Can you please guide me if i can do this process in

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** July 6, 2018 at 6:40 am #

REPLY ↩

Sorry, I don't have material on deep learning in R.



**ibrahim** July 7, 2018 at 6:04 am #

REPLY ↩

Hi Jason,

Thanks for posting such a useful post. I would like to ask what does neurons mean? What does it mean when it is 4 and 1?



**Jason Brownlee** July 7, 2018 at 6:20 am #

REPLY ↩

It is the number of nodes within a layer.

Maybe this will help:

<https://machinelearningmastery.com/faq/single-faq/how-many-layers-and-nodes-do-i-need-in-my-neural-network>

Start Machine Learning



**ibrahim** July 13, 2018 at 3:52 am #

REPLY ↗

Thanx again Jason. I have been searching an issue for a long time. There is some revelations about my question in this site however, there is not the exact one. I have an electricity time series consumption with hourly loads and temperatures. I implemented NARXnet to the time series(3 years hourly loads) and predicted next 24 hrs' values. There is a cycle that repeats every week so that i have chosen 168 time delay for NARXnet. The predictions is satisfactory. Now, i would like to implement LSTM for the time series, however, i cannot conceptualize some LSTM inputs for modelling. For example inputs in the LSTM layer must be 3d format [samples, timesteps, features]. What is samples, timesteps and features for my 26304 observations, 2-11 prediction attributes(some dummies for days) and 168 cycles? The second question is can i adopt autoreg stateful as True enable autoregression adaption to say 168 for my time series? Fourth one is how can has dominant patterns in data? The last one is can to catch correlations in the days?

Thanks from now for your responses.



**Jason Brownlee** July 13, 2018 at 7:45 am #

Samples are sub-sequences, time steps features are what was recorded or observed at

You can use LSTMs for autoregressive type models

Stateful may not be required, I find it is often not required.

This may help you prepare your data:

<https://machinelearningmastery.com/faq/single-faq/how-do-i-prepare-my-data-for-an-lstm>

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**ibrahim** July 25, 2018 at 8:08 pm #

REPLY ↗

Dear Jason, hello again

I tried to do sth to prepare my data that you directed. As mentioned earlier, my data has 26304 rows (with hourly observations) and 1(features)+9 (day and holidays' dummies) columns. I wanted to add autoregression to the problem so that I prepared my data with the shape (this part contained validation and train sets) 25969x168x10. Here 25969 refers observation numbers, 168 refers 168 previous look back of to be predicted value (namely 168 back step AR values for 169th value that will be predicted) and 10 refers data dimensions. In implementation I have some challenges. Train and val sets have 24000 and 1969 values with 168x10 matrices. When I determine timesteps different from 168 with `input_shape=(time_steps,data_dim)`, it causes error such that expected `lstm_1_input` to have shape (None, 24, 10) but got array with shape (24000, 168, 10). Does it mean my sequence will be chain in length of 168 or 24 by determining timesteps as if 168 or 24? Namely 168 timesteps refer 168 lstm units will be in chain with horizontal row?

When I take timesteps 168, it is OK.

Start Machine Learning

Another question, return sequences cause problem in the dataset when i use it? The problem is actually in data dimension mismatch.

My aim was that when i prepared my data with 3d matrix, by specifying timesteps 24, each of my  $168 \times 10$  values will be fed to chain with length of 24. When i specified batch size say 48, my weights will be updated at then end of  $48 \times 168 \times 10$  th round in the chain of 24 horizontal lstm units. I tried to use first tensor as days sequence as you expect. However, i failed in these implementations. I am struggling with finding where am i wrong ? I will be glad if you can help me also?

Thanks from now.



**Jason Brownlee** July 26, 2018

I have posts on how to prepare  
<https://machinelearningmastery.com/forecasting-time-series/>

You can output a sequence either as a



**Finn** July 20, 2018 at 8:19 pm #

Hi Jason:

I am doing the univariate lstm, but the result is very tutorial:

Comment out the line that fits the LSTM model in w

```
yhat = forecast_lstm(lstm_model, 1, X)
```

And replace it with the following:

```
yhat = y
```

This should produce a model with perfect skill (e.g. a model that predicts the expected outcome as the model output).

The results should look as follows, showing that if the LSTM model could predict the series perfectly, the inverse transforms and error calculation would show it correctly.

What is this part meaning, why use `yhat=y`, if use this, the prediction result is the testset, right? which is not the lstm predicted result.

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** July 21, 2018 at 6:34 am #

REPLY ↗

LSTMs are poor at time series forecasting, I recommend trying classical methods first:

<https://machinelearningmastery.com/start-here/#timeseries>



**Finn** July 23, 2018 at 3:11 am #

Cheers Jason

Start Machine Learning



**Awes** July 31, 2018 at 12:15 am #

REPLY ↗

Hello,

So i followed few tutos of yours and i'm trying to build an LSTM that takes 3 last inputs  
Let's say my data is as (1,2,3,4,5,6,7,8) so i want that in the prediction i input (1,2,3) and the prediction  
should be 4 (in the best case).

so here what i did

for the training which i assume is ok

```
train_reshaped = train_scaled[:, 0:3].reshape(len(train_scaled), 1, 3)
lstm_model.predict(train_reshaped, batch_size=2)
```

and for the test

```
X,y= test_scaled[i, 0:3].reshape(1, 1, 3), test_scale
yhat = forecast_lstm(lstm_model, 1, X)
```

here is my test input [[[-0.94936709 -0.97468354 -0.94936709]]]

```
def forecast_lstm(model, batch_size, X):
 yhat = model.predict(X, batch_size=batch_size)
 return yhat[0,0]
```

i get this ValueError: could not broadcast input array

i could not figure it out why its coming

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about  
services and special offers by email. For more  
information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** July 31, 2018 at 6:04 am #

REPLY ↗

I think the input shape would be (1,3,1) not (1,1,3).

Other than that, I'm not sure I follow your problem, sorry. Perhaps you can explain the issue another way or give more context?



**Awes** August 9, 2018 at 8:02 pm #

REPLY ↗

So what I'm trying to do is the following.

i have a sequence of timestamp say [1,2,3,4,5,6,7,8,9,10,11]

i want to input the networking [[1],[2],[3]] and it predicts [4] then i slide the window by one [[2],[3],[4]] and it predicts [5] etc

here is the modified code

```
repeats = 30
error_scores = list()
for r in range(repeats):
 # fit the model
 lstm_model = fit_lstm(train_scaled, 1, 500, 3)
 print('trained',r,'/30\n')
 train_reshaped = train_scaled[:, 0:3].reshape
```

## Start Machine Learning

```
#print(train_reshaped[0].shape)
lstm_model.predict(train_reshaped, batch_size=1)
walk-forward validation on the test data
predictions = list()
for i in range(len(test_scaled)):
 # make one-step forecast
 # X, y = test_scaled[i, 0:3].reshape(1, 1, 3), test_scaled[i, -1]
 X, y = test_scaled[i, 0:3], test_scaled[i, -1]
 yhat = forecast_lstm(lstm_model, 1, X)
 # invert scaling
 yhat = invert_scale(scaler, X, yhat)
 # invert differencing
 #yhat = inverse_difference(raw_values[-10])
 yhat = inverse_difference(raw_values, yhat)
 #print(convertToDate(yhat),convertToDate(raw_values[-10]))
 print(day_diff(yhat,raw_values[-(len(test_scaled)+1)]))
 #yhat = inverse_difference(raw_values[-10])
 # store forecast
 predictions.append(yhat/86400)
 # report performance
 ytrue=raw_values[-(len(test_scaled))+1:-1]
 rmse = sqrt(mean_squared_error(ytrue, predictions[-1]))
 mae = numpy.sum(numpy.absolute((ytrue-predictions[-1])))
 print('%d) Test RMSE: %.3f' % (r+1, rmse))
 print('%d) Test MAE: %.3f' % (r+1, mae))
 error_scores.append(rmse)
 # summarize results
results = DataFrame()
results['rmse'] = error_scores
print(results.describe())
```

## Start Machine Learning X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**



**Jason Brownlee** August 10, 2018 at 6:13 am #

REPLY ↗

Sounds great, this function will help to prepare the data:

<https://machinelearningmastery.com/convert-time-series-supervised-learning-problem-python/>



**Awes** August 10, 2018 at 8:48 pm #

Hello,

I already did that. So my what I did here is taking timestamps series, make the difference between them (in days), scale those differences and input it to the network. I just wanted to make sure that I input a sequence 3 data points and the network should predict the 4th data point.

Another question is: is it really important to repeat the states during the training?

**Start Machine Learning**



**Jason Brownlee** August 11, 2018 at 6:10 am #

Depends on the problem, generally, I don't find the state adding much.



**Jiapeng Yuan** July 31, 2018 at 11:33 am #

REPLY ↗

hello Jason, could you please explain why the LSTM seems lagged (based on blue and orange line)



**Jason Brownlee** July 31, 2018 at 2:57 pm #

Bad performance. E.g. it has learned :



**Jingfei** August 3, 2018 at 10:14 am #

Hi Jason, While I am checking your multi-step prediction, it seems that it is using the predicted value back for next step prediction, it might be causing the whole result inaccurate. Is there any way to fix this?



**Jason Brownlee** August 3, 2018 at 2:23 pm #

REPLY ↗

Yes, this is called an iterative or recursive model.

An alternative with neural networks is to forecast the multiple steps directly as a vector.

Perhaps this post will help to understand the difference between recursive and direct multi-step forecasting:

<https://machinelearningmastery.com/multi-step-time-series-forecasting/>



**Jingfei** August 4, 2018 at 10:24 am #

REPLY ↗

Thank you very much for your hint. It's been very helpful. Do you have more tutorials about modeling building strategies like using multiple datasets to build a model and make use of different activation and regularization to find the best way to improve the model behavior? or you have any books on this topic?



**Jason Brownlee** August 5, 2018 at 5:22 am #

REPLY ↗

[Start Machine Learning](#)

This general process might help:

<https://machinelearningmastery.com/start-here/#process>



Oleg legorov August 29, 2018 at 5:26 am #

REPLY ↗

Hi,

This phrase seems to be false “By default, an LSTM layer in Keras maintains state between data within one batch”. According to Keras documentation, samples within a batch are processed independently.



Jason Brownlee August 29, 2018 at 8:15 pm #

Why do you say that? Can you point to a reference?



Oleg legorov August 30, 2018 at 12:30 pm #

Sure, here <https://keras.io/getting-started/sequential-model-guide/>.  
 “Batch: a set of N samples. The samples in a batch are processed independently of each other, which means that a given sample is not dependent on the previous one. This is important if you are trying to learn arbitrarily long dependencies, because it allows the network to have a long ‘memory’ without having to store all the information in memory.”  
 I’m new to neural networks and was curious about how to learn arbitrarily long dependencies in time series. In my current understanding, the way to learn arbitrarily long dependencies is to have a recurrent neural network and specify stateful=True

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



Jason Brownlee August 30, 2018 at 6:31 am #

REPLY ↗

Nice find. You could be right, but I’m not convinced that applies to an RNN.



Andrew Jidw September 5, 2018 at 6:53 pm #

REPLY ↗

Hello Jason, I would like to ask, if the time series itself is stable, I need to establish a persistence model. After the establishment, I found that the performance is better than LSTM.



Jason Brownlee September 6, 2018 at 5:33 am #

REPLY ↗

Yes, this is often the case. Try a suite of other methods like MLPs, CNNs and hybrids. Also try classical methods like SARIMA and ETS.

Start Machine Learning



**Geeta** September 18, 2018 at 11:09 am #

REPLY ↗

Hi,

If we want to apply this to log analytics, how can we transform the data and apply to Time Series and NN

Thanks,

Geeta



**Jason Brownlee** September 18, 2018 at 2:25 pm #

REPLY ↗

It depends on the logs.

You will likely have to write code to make your



**Suker** September 18, 2018 at 6:22 pm #

Hi,

I'm not clear that is this problem itself a destined to solve such problems?

Thanks,

Suker

## Start Machine Learning

X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** September 19, 2018 at 6:16 am #

REPLY ↗

No, LSTMs are terrible at univariate time series forecasting and are easily out-performed by SARIMA and ETS.



**via** September 22, 2018 at 12:54 am #

REPLY ↗

Hello!

Why did you reverse the order in this line:

X, y = train[:, 0:-1], train[:, -1]

Thanks!



**Jason Brownlee** September 22, 2018 at 6:30 am #

REPLY ↗

I don't follow, where exactly?

Start Machine Learning



**Melania** September 25, 2018 at 1:27 am #

REPLY ↗

thank you for your explanation,

My question, at this level and after evaluating the model, is to know how to predict and visualize some values in the future.

Otherwise, How visualize the value  $t + 1$  which follows the last value of the data base

thanks



**Jason Brownlee** September 25, 2018 at 6:30 am #

REPLY ↗

The tutorial shows how to plot  $y$  vs  $y_{pred}$

What problem are you having exactly?



**Tyra Nguyen** October 8, 2018 at 3:25 am #

Hi Jason,

Thank for your tutorial, I love it!

I have some questions:

1. I wonder whether LSTM can forecast prices of crypto currency?
2. May I use the historical data to forecast? I plan to use ARIMA.
3. I see many people use ARIMA to do this, so can't we use ARIMA.

Sorry for my lack of knowledge. I hope to receive your answer.

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** October 8, 2018 at 9:28 am #

REPLY ↗

You can develop a model to forecast anything you like. I don't know if crypto prices are predictable, I suspect that are a random walk.

By definition, a skillful forecasting model will make use of historical data to predict the future.

Perhaps this will help:

<https://machinelearningmastery.com/how-to-develop-a-skilful-time-series-forecasting-model/>



**Theekshana** October 18, 2018 at 2:57 am #

REPLY ↗

Hi Jenson,

Can these LTSM neural networks preserve multiple time series variation patterns? In my project, I have a set of different time series variation patterns. So, the thing I want to know is that,

1. Are LTSM networks capable of identifying the current variation pattern and predicting the next value accordingly (given the previous variation)?

## Start Machine Learning

2. If not, what will be a solution to this type of problem (Time series pattern identification and prediction)

Can you please give me some guidelines.

Thank you for the tutorials. They were really helpful. !! 😊



**Jason Brownlee** October 18, 2018 at 6:36 am #

REPLY ↗

Yes, but LSTMs are terrible at time series forecasting. Try MLPs, CNNs and hybrid models as well.

Checkout this tutorial:

<https://machinelearningmastery.com/how-to-develop-forecasting-of-household-power-consumption/>



**Charles Yu** October 26, 2018 at 3:19 am #

Hi Jason,

Is there any tutorial about single (categorical) variable LSTM based forecasting?

Thanks



**Jason Brownlee** October 26, 2018 at 4:01 am #

I believe you are referring to time series classification.

I have an example of time series classification with LSTMs here:

<https://machinelearningmastery.com/how-to-develop-rnn-models-for-human-activity-recognition-time-series-classification/>



**Charles Yu** October 26, 2018 at 3:26 am #

REPLY ↗

Hi Jason,

For categorical variable LSTM based forecasting, the requirement is as follows:

I want to forecast what a person's doing @ 10:00AM in some day (weekday) so I collected his doing from day 1 to day n. All of them are in the weekdays. Then, if I want to forecast his doing @ day (n+1) @ 10:00AM. How to do that? Is there any related tutorial in your posting?



**Jason Brownlee** October 26, 2018 at 5:41 am #

REPLY ↗

Yes, I believe the time series activity recognition is a close fit for what you want.

<https://machinelearningmastery.com/how-to-develop-forecasting-of-household-power-consumption/>

Start Machine Learning

[series-classification/](#)**Devakar Kumar Verma** November 15, 2018 at 12:26 am #[REPLY ↗](#)

Hi Jason,

At the time of training the model, we are resetting the state after each epoch

```
for i in range(nb_epoch):
 model.fit(X, y, epochs=1, batch_size=batch_size, verbose=0, shuffle=False)
 model.reset_states()
```

But, if we want to reduce the learning rate at the time of training, we can use “ReduceLROnPlateau” defined in keras. The learning rate will decrease at a certain point in time that keras provides due to the resetting of all the parameters. This will help us to avoid getting stuck in less optimal minima.

What are your views on this?

**Jason Brownlee** November 15, 2018 at 5:45 pm #[X](#)

↗

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)
**Devakar Kumar Verma** November 15, 2018 at 7:37 pm #[REPLY ↗](#)

Hi Jason,

As per your implementation learning rate is constant at 0.001 after every epoch, due to initialization of Adam optimizer in keras

```
keras.optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=None,
decay=0.0, amsgrad=False)
```

From your second point, you mean to have the learning rate scheduler which operates on begin\_of\_epoch.

**Jason Brownlee** November 16, 2018 at 6:14 am #[REPLY ↗](#)

Not sure I follow, can you elaborate?

**mk** December 11, 2018 at 7:49 pm #
[Start Machine Learning](#)

Hi Jason,

I ran your program ,but,epoch is 23 during trainning.I set epoch 100 , epoch is still 23 during trainning.I'm a little confused.



**Jason Brownlee** December 12, 2018 at 5:51 am #

REPLY ↩

I have some suggestions here that might help:

<https://machinelearningmastery.com/faq/single-faq/why-does-the-code-in-the-tutorial-not-work-for-me>

And some suggestions related to debugging here:

<https://machinelearningmastery.com/faq/single>



**Muhammad Akbar Almuttaqin** December 12, 2018 at 5:51 am #

Hello,

I am now working in a project that involves a time series. I am trying to predict the price of milk. In my mind, how could LSTM performs better than another model like linear regression or decision tree? LSTM only consider data alone and neglect any other attributes? For example, in this case we are dealing with the price of milk. How could it outperform linear regression, that considers many other factors too?

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** December 16, 2018 at 5:21 am #

REPLY ↩

Generally LSTMs do not perform well on time series forecasting problems.



**Anthony** December 21, 2018 at 3:18 am #

REPLY ↩

Hi Jason, can I edit the function 'def forecast\_lstm(model, batch\_size, X)', in the line 75, in Complete lstm Example, before 'yhat = model.predict(X, batch\_size=batch\_size)' to use a previously saved model with model.save('Lstm\_model.h5'), or where can I edit anywhere else? Thank you very much



**Jason Brownlee** December 21, 2018 at 5:31 am #

REPLY ↩

Sure.

## Start Machine Learning



**Anthony** December 21, 2018 at 10:03 am #

REPLY ↗

Excuse me, and then in which step need I edit in Complete Lstm Example to use a previously saved model with `model.save('Lstm_model.h5')`? is there an example? Thank you.



**Jason Brownlee** December 21, 2018 at 3:18 pm #

REPLY ↗

Sorry, I'd don't have the capacity to customize the example for you.

Perhaps try getting familiar with the API first:

<https://machinelearningmastery.com/save-load-keras-lstm-networks/>



**Anthony** December 27, 2018

Yes, of course, you are right. I will replace the line 98 `lstm_model = fit_lstm(...)` with `lstm_model = load_model('lstm_model.h5')`, as you suggested. I will also change <https://machinelearningmastery.com/forecasting-time-series-usage-baltimore/> in the line 15, at the end of the code. Thank you very much

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** December 28, 2018 at 5:45 am #

You can load instead of re-fit a model.

I don't have the capacity to advise you on how to best modify the code example for your requirements – likely a one line change would be insufficient.



**venkat chana** January 21, 2019 at 11:51 am #

REPLY ↗

Hi, Jason at the section "Complete LSTM Example", what should I edit in the code to change the `batch_size`? Please could you give me an example?



**Jason Brownlee** January 21, 2019 at 12:02 pm #

REPLY ↗

Specify the `batch_size` argument to the `fit()` function.



**venkat chana** January 22, 2019 at 2:19 pm #

REPLY ↗

Start Machine Learning

Jeson , sorry to disturb,

In the code “Complete LSTM Example” I used another dataset and 5 is a factor both of the size of the training and test datasets.

so I changed

the line 98: `Istm_model = fit_Istm(train_scaled, 1, 3000, 4)`  
to: `Istm_model = fit_Istm(train_scaled, 5, 3000, 4)`

the line 101: `Istm_model.predict(train_reshaped, batch_size=1)`  
to: `Istm_model.predict(train_reshaped, batch_size=5)`

the line 108: `yhat = forecast_Istm(Istm_model, 1, X)`  
to: `yhat = forecast_Istm(Istm_model, 5, X)`

but I got this error:

`ValueError: could not broadcast input array`  
at the line 77 `yhat = model.predict(X, batch`

might you help me to fix this bug? Thank you



**Jason Brownlee** January 22, 2019

I have some suggestions here  
<https://machinelearningmastery.com/forecasting-time-series-long-short-term-memory-networks-python/>



**venkat chana** January 22, 2019

Yes, of course, you don't have to write the code for me or fix my bugs, I just wanted to change the `batch_size` from 1 to another number, I just wanted a clarification that I have not understood well by reading your code. Thank you 😊



**Julius** January 28, 2019 at 10:27 pm #

REPLY ↗

Hi Jason,

I'm a bit confused about the input shape.

Let's say I've got 2500 time series of class A, and 2500 time series of class B, both 1000 time steps long. There is only one numerical feature per time series. I want to train a LSTM using Keras to classify the time series into one of the two classes. They can be easily distinguished using features which require the LSTM to understand the time correlation between the data, like one class has a seasonal pattern, while the other hasn't. Therefore I guess I need to use a stateful LSTM, and therefore call `model.reset_states()` manually right?

But I'm totally confused about how to construct the `input_dimensions` for Keras.

As far as I understand it the input for the LSTM is defined as (samples, time steps, features).

Is my X now (5000,1000,1) and I call `model.fit()` exactly once? When do I call `model.reset_states()`?

[Start Machine Learning](#)

Or do I show the time series each time step at a time, so my input is (1,1,1) and I just show it then 5000 times for each time series 1000 times per each time step the single feature, so call `model.fit()` 5,000,000 times?! I would then call after I'm finished with one time series the `model.reset_states()` method to forget everything for the new time series? That is the version you've been using in this tutorial?

Or is it (1,1000,1), and I call `model.reset_states()` after calling `fit()` 5000 times (per time series)?



**Jason Brownlee** January 29, 2019 at 6:11 am #

REPLY ↗

Perhaps this will help:

<https://machinelearningmastery.com/faq/single-faq/what-is-the-difference-between-samples-timesteps-and-features-for-lstm-input>



**Julius** January 29, 2019 at 9:07 pm #

Thanks, that really helped a lot!



**Julius** January 30, 2019 at 12:57 am #

Ok, I'm lost again. Let's say I've got 25, 1000 time steps and 2 features. To combine labels and put them together in one batch. time series. But because I've specified a batch size of 25 I now always need to show it 25 timeseries?!

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** January 30, 2019 at 8:13 am #

REPLY ↗

No, you can make a prediction for one sample [1,1000,2]



**Jack** February 10, 2019 at 11:10 pm #

REPLY ↗

Hello, Jason, thanks for your excellent tutorial! Now I am confused about some places of LSTM in Keras. Hope you have time to answer those questions! Thanks in advance!

Assumed that there is a time series data set with two variables and one output. The data set is not continuous so it can be actually seen as a group of two independent sequences. Every sequence is not the same length. One has 100 data points and another one has only 50 data points. I want to use the sequence with 100 data points as the training data set and the rest sequence as the test data set.

The first one: Should I use stateful LSTM or stateless LSTM? By the way, it seems that when having only one continuous sequence, we should select stateful

Start Machine Learning

The second one: How should I deal with the inputs with different length?

Thank you again!



**Jason Brownlee** February 11, 2019 at 8:00 am #

REPLY ↗

Start with stateless and try stateful to see if it makes a difference.

Use padding to make all input the same length:

<https://machinelearningmastery.com/data-preparation-variable-length-input-sequences-sequence-prediction/>



**Jack** February 11, 2019 at 12:19 pm #

Thanks for your prompt reply! You



**Jack** February 11, 2019 at 6:13 pm #

Hello, Jason. Now I am doing some exper

Let's say that I want to use the Many-to-Many architecture. I have 50 sequences of data every same time step. So both the input and output have 50 sequences of data for input and output data. The length of input data or output data is fixed to 100 and padding is used to make all sequences the same length (100).

The question is: the input shape required by Keras (samples, time steps) should be (50, 100), right?

Thanks for your reply in advance!

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** February 12, 2019 at 7:54 am #

REPLY ↗

If you want to predict one output for one input, then this would be one-to-one.

I have tons of help on how to shape data, start here:

<https://machinelearningmastery.com/faq/single-faq/what-is-the-difference-between-samples-timesteps-and-features-for-lstm-input>



**Jack** February 12, 2019 at 3:52 pm #

REPLY ↗

Thanks for your reply.

Actually I have read the website before. It is very helpful. However, I still feel confused about the below paragraph.

Start Machine Learning

"As another example, if you have 20 time steps of data for 2 features, you might represent it as 20 samples with 1 time step [20, 1, 2] or as 1 sample with 20 time steps [1, 20, 2]. In the first case, you will have 20 samples, meaning 20 input and output examples and your model will make 20 predictions. In the latter case, you have 1 sample, with 1 input and 1 output and the model will make 1 prediction that is accumulated over 20 time steps."

I understand the first case. However, in the second case, you say that we have only one output to be predicted. Which time step will the output be calculated and output at? It is at the last time step (20th time step)?

As for the architecture of LSTM, the below website shows it. X has two features, and Y is a scalar. I want to predict one x for one y at every same time step. I think it is many-to-many, isn't it?

<https://www.dropbox.com/s/e8zp520iz6lpk8>

Thanks for your help! I really appreciate it!



**Jason Brownlee** February 13,

Correct, the output is predicted (after last time step).

If you want to make a prediction for each observation, one forecast per observation, in case I was confusing.



**dyy** February 12, 2019 at 2:17 pm #

REPLY ↗

Hi Jason!

what if the data contains seasonality which is also anomaly that carry significant information to the domain? Do i need to difference it as well?



**Jason Brownlee** February 13, 2019 at 7:47 am #

REPLY ↗

If the series contains a seasonal structure, remove it from the series and model the residual.

This is called a seasonal adjustment:

<https://machinelearningmastery.com/remove-trends-seasonality-difference-transform-python/>



**YN** February 15, 2019 at 8:46 am #

REPLY ↗

Hi Jason,

[Start Machine Learning](#)

I was following your steps and testing some synthetic data with straightforward patterns such as (0,0,0,...,0,1) repeated n times. I was trying to see if LSTM network can learn such patterns with long range dependency and make correct predictions moving forward. It turns out that the optimizer often gets stuck at poor local minimums, even though I tried different learning rates. Instead, if I use the training sequence as a time-series and set the batch\_input\_shape to be (1, seq\_length, n\_lags), the training becomes much easier and converges much faster. Do you expect treating training data as one sample multiple steps to be better than as multiple samples one step in general? Thanks.



**Jason Brownlee** February 15, 2019 at 2:18 pm #

REPLY ↗

It really depends on the problem.

For example, I show how using 1 time step samples works:  
<https://machinelearningmastery.com/memory-in-lstm-networks/>



**Rushikesh Khot** March 2, 2019 at 11:56 pm #

X, y = train[:,0:-1], train[:, -1]  
X = X.reshape(X.shape[0], 1, X.shape[1])

!Error:

IndexError Traceback (most recent call last)  
in

```
—> 1 X, y = train[:,0:-1], train[:, -1]
2 X = X.reshape(X.shape[0], 1, X.shape[1])
```

IndexError: too many indices for array

I am getting an error as above. Can you help me with this Jason??

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** March 3, 2019 at 8:03 am #

REPLY ↗

Sorry to hear that, I have some suggestions here:

<https://machinelearningmastery.com/faq/single-faq/why-does-the-code-in-the-tutorial-not-work-for-me>



**Joshua** March 5, 2019 at 7:26 pm #

REPLY ↗

Hi Jason!

I have to say, really interesting blog. I'm recently starting a new project and I'm new with LSTM (actually still kind of new with NNs). This is why I'd like to ask you for advice: My problem is similar to the one on this blog, but instead of having a dataset which is a

Start Machine Learning

length 8, therefore, instead of training with only 1 vector, I'd like to train with many of them and then test with many others.

Particularly, I'd like to train a model to predict the last 3-4 entries of each vector. For this, I'm thinking of using ~300 vectors for training, and then use the other 200 to test the predictions of the last 3-4 entries of each vector.

Could you give me any recommendation/link of how to proceed?

Thanks!



**Jason Brownlee** March 6, 2019 at 7:47 am #

Sounds reasonable.

Yes I have tens of tutorials that will help you with  
<https://machinelearningmastery.com/start-here/>

REPLY ↗

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Luan** March 6, 2019 at 8:25 pm #

Jason great job –

One question, do you have any similar tutorial to use times and so they are not uniform distributed?

Best!

REPLY ↗



**Jason Brownlee** March 7, 2019 at 6:48 am #

No. But perhaps the model does not care. You can try modeling with variable input time steps, or try padding with zeros, or try resampling so they are consistent – then compare the results of each approach?



**Luan** March 7, 2019 at 10:06 pm #

REPLY ↗

What do you mean by 'variable input time steps'?

As far as I know, the time steps need to be fix when constructing the network, since it is required for the batch\_input\_shape parameter of the LSTM layer. Or am I wrong?

Thx!



**Jason Brownlee** March 8, 2019 at 7:50 am #

REPLY ↗

## Start Machine Learning

They can vary, but you can fill the gaps with zeros, you can pretend the time steps are evenly spaced when they're not, you can truncate or resample the time series, you have many options.



**Duc** March 11, 2019 at 11:12 pm #

REPLY ↗

Hi Jason,

Thanks for very useful implementation. I would have a following question.

I trained the LSTM model on train data with batch\_size = 32. Then I tried do prediction with the test data using the learned model with batch\_size = 1 but it doesn't work.

Is it required to set the batch\_size = 1 (in both training and prediction)?

Thanks.



**Jason Brownlee** March 12, 2019 at 6:53 am #

This is a common problem when using an LSTM model to make predictions. You can find more information about this in my post on how to use different batch sizes when making predictions with an LSTM model: <https://machinelearningmastery.com/use-different-batch-sizes-predictions-lstm/>



**Sai** March 12, 2019 at 9:37 pm #

How to forecast for the future values .. Suppose I have data till December then how can we predict for the month January.



**Jason Brownlee** March 13, 2019 at 7:55 am #

REPLY ↗

This post will help:

<https://machinelearningmastery.com/make-predictions-long-short-term-memory-models-keras/>



**sai** March 13, 2019 at 7:55 pm #

REPLY ↗

After saving the model load\_model is used bt what is the input value of X.

ie. Suppose I have a dataset(data.csv) how can I put into it.

code:

# snip...

# later, perhaps run from another script

```
load model from single file
```

```
model = load_model('lstm_model.h5')
```

Start Machine Learning

```
make predictions
yhat = model.predict(X, verbose=0)
print(yhat)
```

How many time steps it gives the result.



**Jason Brownlee** March 14, 2019 at 9:20 am #

REPLY ↗

X is whatever your model expects as input in order to make a prediction. E.g. the input components of a given sample in order to get the output component of the sample.



**sai** March 14, 2019 at 12:44 pm #

I input the dataset which contains the values. I don't understand what does those values considered as predicted values? Doesn't it work like ARIMA where we have to input previous values?



**Jason Brownlee** March 14, 2019 at 1:01 pm #

You model it like any supervised learning problem. Except here, the input samples are

## Start Machine Learning

×

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Michael** March 16, 2019 at 8:32 pm #

REPLY ↗

Thank you for this great tutorial.

I was going through it, step by step, but this ends before 'the magic happens' (unfortunately):

```
X, y = train[:, 0:-1], train[:, -1]
X = X.reshape(X.shape[0], 1, X.shape[1])
```

Up to this point, you perfectly advised me and I could check all the lines by myself.

Now, when using the complete example, this is not that easy ;/



**Jason Brownlee** March 17, 2019 at 6:21 am #

REPLY ↗

You can learn more about reshaping arrays here:

<https://machinelearningmastery.com/index-slice-reshape-numpy-arrays-machine-learning-python/>



**Athiran** March 28, 2019 at 3:52 pm #

Start Machine Learning

sir,

how to use LSTM Networks to Predict Engine Condition on Large Scale Data Processing



**Jason Brownlee** March 29, 2019 at 8:20 am #

REPLY ↗

My best advice is to start here and follow the outlined process:

[https://machinelearningmastery.com/start-here/#deep\\_learning\\_time\\_series](https://machinelearningmastery.com/start-here/#deep_learning_time_series)



**Kunal** April 3, 2019 at 3:57 am #

REPLY ↗

Why rmse ? Why not other types of errors dataset or something?



**Jason Brownlee** April 3, 2019 at 6:49 am #

RMSE is widely used and deeply under-

MAE is also popular.



**Brain** April 3, 2019 at 4:06 am #

Hi Jason,

I'm not sure if this is the best tutorial to post this question but here I go!

I've gone through multiple of your tutorials of RNN for prediction and classification, and I've been applying them to my data with limited success so far. So I'd like to go one step further. Here is my problem:

I have 2 features and I want to predict both of them. However, one is a continuous variable and the other one is a class. So far I've been predicting them separately, but now I'd like to merge the prediction because, of course, both features are related to each other. This why I'd like to build a network with keras based on LSTM cells that for an input with two features is able to provide two different outputs, one should go behind a binary\_crossentropy activation function to obtain the class, and the other one after no-activation function ( $f(x)=x$ ).

I've seen a tutorial where you include attention in a enc-dec network that you show how to merge upper layer of the network. However, do you know how could we split the network to obtain two different results?

Thanks

Brain

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** April 3, 2019 at 6:50 am #

REPLY ↗

Start Machine Learning

Perhaps you can try a model with two outputs, each with its own loss function?



**Brain** April 3, 2019 at 5:31 pm #

REPLY ↗

Thanks! do you have any example doing such thing?



**Jason Brownlee** April 4, 2019 at 7:39 am #

REPLY ↗

Yes, this tutorial will help:

<https://machinelearningmastery.com/ke>



**lingya** April 5, 2019 at 1:44 am #

Hi

ine 25-30

```
1 # invert transform
2 inverted = list()
3 for i in range(len(differenced)):
4 value = inverse_difference(series, d)
5 inverted.append(value)
6 inverted = Series(inverted)
7 print(inverted.head())
```

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**. Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).
 [START MY EMAIL COURSE](#)

I got key error since the index is negative for history[-interval], which is series[-interval] here.

Should I use history = series.values instead of series(line 28

value = inverse\_difference(series.values, differenced[i], len(series)-i)?



**Jason Brownlee** April 5, 2019 at 6:19 am #

REPLY ↗

Perhaps try it and see how you go.



**JZ** April 8, 2019 at 2:30 pm #

REPLY ↗

Hi,

You have many posts on LSTM for time series and some seem to be dated. Is this the most current one. For the dated one's can you clearly mention that in the title.

Thanks!

Start Machine Learning



**Jason Brownlee** April 9, 2019 at 6:19 am #

REPLY ↗

You can see my most recent / best posts on LSTMs for time series here:

[https://machinelearningmastery.com/start-here/#deep\\_learning\\_time\\_series](https://machinelearningmastery.com/start-here/#deep_learning_time_series)



**mbelahcen** April 11, 2019 at 12:00 am #

REPLY ↗

Hello Jason,

So i'm trying to use LSTM to forecast a timeseries but I have an issue with an important feature. Let's say that X is what i'm trying to predict.

Let's consider  $X = a+b+c$ . The values of a,b,c depend represents for example a color. What's the best way  
 1- Use LSTM to predict a,b,c depending on time or  
 2- sum up the values a,b,c and then apply LSTM

To make more sens out of this, let's say that a,b,c are total profit.

I also would like to understand why there is no validation

Thanks



**Jason Brownlee** April 11, 2019 at 6:42 am #

REPLY ↗

Using two models one after the other is a good start.

Try predicting all elements and see if the model can keep the values consistent?

Get creative, perhaps brainstorm alternate framings of the problem and test them out?



**Alex Teng** April 19, 2019 at 2:31 pm #

REPLY ↗

Hi Jason,

Thank you so much for the tutorial in detail! After training the data and validating it with test dataset, I would like to produce predictions on new data, say for the future time steps without real values y (i.e., trained data for 10 years, and predict for future years). Could you please tell me how to make a proper X to use the trained model to predict?



**Jason Brownlee** April 19, 2019 at 3:08 pm #

REPLY ↗

Yes, you can call `model.predict()`.

Perhaps this post will help:

<https://machinelearningmastery.com/make-predictions-long-short-term-memory-networks/>

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**belahcenm** April 25, 2019 at 1:42 am #

REPLY ↗

Hello Jason! Thank you for your tutorials.

I have tuned my model through epochs, number of neurons and batch size and I ran repeated experiments. The rmse in many cases is very satisfying for my problem. However, in some cases the error is not. Is it normal? Should I consider my model ready if the majority of experiments are satisfying or a model is ready only if the rmse is almost stable?

Also, what's the difference between forecasting before and within the for loop? (See below)

case 1: `lstm_model.predict()` before for loop

```
output = lstm_model.predict(test_reshaped, batch_size=1)
predictions = list()
for i in range(len(output)):
 yhat = output[i,0]
 X = test_scaled[i, 0:-1]
 # invert scaling
 yhat = invert_scale(scaler, X, yhat)
 # invert differencing
 yhat = inverse_difference(raw_values, yhat, len(test_scaled)-i)
 # store forecast
 predictions.append(yhat)
```

case 2: `lstm_model.predict()` inside for loop

```
for i in range(len(test_scaled)):
 # make one-step forecast
 X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
 yhat = forecast_lstm(lstm_model, 1, X)
 # invert scaling
 yhat = invert_scale(scaler, X, yhat)
 # invert differencing
 yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)
 # store forecast
 predictions.append(yhat)
```

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** April 25, 2019 at 8:22 am #

REPLY ↗

It might be a good idea to counter the variance of the final model by making ensemble predictions:

<https://machinelearningmastery.com/ensemble-methods-for-deep-learning-neural-networks/>



**Neil z** April 28, 2019 at 6:07 pm #

REPLY ↗

Excellent blog !

Start Machine Learning



**Jason Brownlee** April 29, 2019 at 8:17 am #

REPLY ↗

Thanks!



**help** May 4, 2019 at 8:05 am #

REPLY ↗

Hi , please would you mindwriting us the function for future prediction (values we don't know) and where we put it in the code .i had a problem dealing with it and thank you so much



**Jason Brownlee** May 5, 2019 at 6:17 am #

Yes, I explain how here:

<https://machinelearningmastery.com/make-predictions-long-short-term-memory-network-python/>



**JCR** May 7, 2019 at 2:57 am #

Hi,

I am having an issue with the walk forward validation.

Running:

```
predictions = list()
for i in range(len(test_scaled)):
 # make one-step forecast
 X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
 yhat = forecast_lstm(lstm_model, 1, X)
 # invert scaling
 yhat = invert_scale(scaler, X, yhat)
 # invert differencing
 yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)
 # store forecast
 predictions.append(yhat)
expected = raw_values[len(train) + i + 1]
print('Month=%d, Predicted=%f, Expected=%f' % (i+1, yhat, expected))

returns "KeyError: -76"
```

The length of my testing data is 75 days, so it seems to be related to that. There's nothing in the chunk above that needs to be specified for my specific data, right?

Interestingly, editing the `inverse_difference` function to return `"yhat + history[interval]"` rather than `"history[-interval]"` seems to rectify the issue, but I'm not sure why...

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** May 7, 2019 at 6:19 am #

REPLY ↗

Perhaps this post will help:

<https://machinelearningmastery.com/backtest-machine-learning-models-time-series-forecasting/>

Also try removing or data scaling until you get the model/test harness working the way you want, then add it back in.



**Berta** May 23, 2019 at 4:21 am #

REPLY ↗

Hi!

Really nice blog!

I was reading this section and I don't really understand what is shown here (<https://machinelearningmastery.com/time-series-forecasting-with-python-keras/>). Both cases are predicting one time step.

I am doing a project where the dataset is the speed of a vehicle trying to predict the future speed. Right now I get a good result but would like to know the main differences between the two approaches.



**Jason Brownlee** May 23, 2019 at 6:09 am #

The linked example is older and this video is more recent:

I have even better examples here that may help:

[https://machinelearningmastery.com/start-here/#deep\\_learning\\_time\\_series](https://machinelearningmastery.com/start-here/#deep_learning_time_series)



**Amin** May 28, 2019 at 5:32 pm #

REPLY ↗

Hi Jason,

in the line:

```
forecast the entire training dataset to build up state for forecasting
train_reshaped = train_scaled[:, 0].reshape(len(train_scaled), 1, 1)
```

you are going to reshape the first column of train\_scaled. If lag=1 is used then it is ok it reshapes the feature column (1 column). But if you use for example 3 previous time step (timeseries\_to\_supervised(data, lag=3) it generates features with 3 columns. So, in this case, what does "train\_scaled[:, 0]" mean? should it be train\_scaled[0:-1, 0] ?

Thanks



**Amin** May 31, 2019 at 9:41 am #

REPLY ↗

[Start Machine Learning](#)

Hi Jason,

Have you had time to look into this? Thanks



**Jason Brownlee** May 31, 2019 at 2:44 pm #

REPLY ↗

Ideally the data should be scaled prior to being transformed into a supervised learning problem.



**mbelahcen** May 29, 2019 at 2:01 am #

REPLY ↗

Hello Jason,

say I have after differencing the data I have history with: `model.predict(history, batch_size=1)`. Then i get history data. I'm confused with what should I revert predicted June 2019 with the month before or with



**Jason Brownlee** May 29, 2019 at 8:55 am #

X

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**guzuomuse** May 29, 2019 at 10:44 pm #

REPLY ↗

thanks for this great article. i have a question: is it necessary (important) to make the data

stationary when using neural network (deep learning).

i have searched some articles. most of them said that the answer should be "NO"

here they are:

<https://www.oreilly.com/ideas/3-reasons-to-add-deep-learning-to-your-time-series-toolkit>

[https://www.researchgate.net/post/in\\_financial\\_time\\_series\\_we\\_have\\_non-stationary\\_time\\_series\\_is\\_it\\_needed\\_to\\_transfer\\_the\\_data\\_to\\_stationary\\_time\\_series\\_for\\_forecasting\\_or\\_not](https://www.researchgate.net/post/in_financial_time_series_we_have_non-stationary_time_series_is_it_needed_to_transfer_the_data_to_stationary_time_series_for_forecasting_or_not)

can you give me your opinion? thank you!



**Jason Brownlee** May 30, 2019

REPLY ↗

## Start Machine Learning

It can help, but is not required.



**guzuomuse** May 30, 2019 at 2:14 pm #

thanks Man!



**Amin** June 1, 2019 at 6:59 pm #

REPLY ↗

Hi Jason, Could you please check the below issue?

I used r2\_Score for test set:

```
from sklearn.metrics import r2_score,
results['r2train'] = r2_score(train_scaled[:, -1], lstm_results)
results['r2test'] = r2_score(raw_values[-12:], predictions)
```

Results:

1) Test RMSE: 113.155

rmse r2test r2train

```
count 1.000000 1.000000 1.000000
mean 113.155189 -0.016695 -0.389893
min 113.155189 -0.016695 -0.389893
25% 113.155189 -0.016695 -0.389893
50% 113.155189 -0.016695 -0.389893
75% 113.155189 -0.016695 -0.389893
max 113.155189 -0.016695 -0.389893
```

The R2 value is negative while RMSE is 113.

## Start Machine Learning

X

You can master applied Machine Learning  
without math or fancy degrees.

Find out how in this free and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** June 2, 2019 at 6:39 am #

REPLY ↗

Yes.



**Amin** June 5, 2019 at 1:43 pm #

REPLY ↗

I used your original code and data (Section: Develop a Robust Result) and just added R2 SCORE in addition to RMSE: `r2=r2_score(raw_values[-12:], predictions)` which inputs are the same as `mean_squared_error` in the original code.

Results:

- 1) Test RMSE: 94.286 R2: 0.294
- 2) Test RMSE: 86.854 R2: 0.401
- 3) Test RMSE: 189.482 R2: -1.851
- 4) Test RMSE: 147.910 R2: -0.737
- 5) Test RMSE: 204.352 R2: -2.316
- 6) Test RMSE: 219.807 R2: -2.836

Start Machine Learning

- 7) Test RMSE: 106.188 R2: 0.105
  - 8) Test RMSE: 108.560 R2: 0.064
  - 9) Test RMSE: 199.271 R2: -2.153
  - 10) Test RMSE: 129.355 R2: -0.329
- ...

As you can see, R2 values are low or even negative.

If there is no error in my code, my question is that should we check multi-criteria like RMSE, R2, etc. simultaneously? or is there any rule for selecting between R2 or RMSE based on the data.

Another example is that I used Grid Search for wide ranges of hyperparameters based on AIC to generate SARIMA Model (Best AIC:4779) and the test R2 is 70 %. When I check the Training R2, it is around 45%. Is the model acceptable or I should use Training R2 instead of AIC? Many Thanks



**Jason Brownlee** June 5, 2019

I recommend choosing a metric or domain experts for your specific data.

If you're unsure, RMSE is a good start.

Model selection is challenging. It is important to have a model evaluation strategy that captures the true performance of the model.

That is the best that I can say without knowing more about your specific data and what you want to do with it. I don't know if you have the capacity to do.

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**



**Tayyab** June 11, 2019 at 8:59 pm #

REPLY ↩

If the amount of data is in thousands of lines, in my case 40K plus lines, the algorithm takes forever to train and evaluate or tune. One approach is to reduce the data but then the evaluation and training results are not for the whole data. How can I evaluate my algorithm in such a scenario?

I am using ARIMA and LSTM and both take forever, like sometimes a whole day. I shifted to google colab GPU but still the same issue. I want to get a solution for the above issue so that I can tune and evaluate both the algorithms in less time.



**Jason Brownlee** June 12, 2019 at 8:00 am #

REPLY ↩

Use a small sample to estimate skill and choose a model, then use all data and a large computer over night to compute the final skill estimate.



**Lopa** June 13, 2019 at 3:28 am #

REPLY ↩

Hi Jason,

## Start Machine Learning

I have gone through all your tutorials on LSTM forecasting. I have trained the model using 549 data points & in order to make predictionI have used 90 data points .

```
X2 = scaler.fit_transform(X2)
X3 = X2[0:-1]
X3=np.reshape(X3, (X3.shape[0], 1, X3.shape[1]))

yhat = model.predict(X3, 1)
```

This generates 89 forecasts. But what if I want to predict only 31 data points using 89 data points how should I modify my code. It would be great if I can get some help.



**Jason Brownlee** June 13, 2019 at 6:22 am #

You can provide 31 input samples to run the model.

Perhaps I don't understand your question?



**Lopa** June 13, 2019 at 7:06 pm #

Hi Jason,

Thanks for your prompt reply. It is so that I have data points. I was trying to forecast Jan 2019 using demonstrated above I can understand that in doing so as the number of inputs.

My questions are :

1. that can these predictions which are generated using Oct 2018-Dec 2018 data be considered to be the predictions for Jan 2019 – Mar 2019 . Is my understanding correct ?
2. Is there a way that I input 90 data points to predict 31 future data points ?

In case of vanilla LSTM I have seen that we can input an array of 3 elements [70,80,90] & predict [100] as shown in one of your tutorials:

<https://machinelearningmastery.com/how-to-develop-lstm-models-for-time-series-forecasting/>

Thanks for helping out by designing these useful tutorials.

## Start Machine Learning



You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**



**Jason Brownlee** June 14, 2019 at 6:40 am #

REPLY ↗

You can frame the prediction problem any way that you wish. It is up to you.

Perhaps I don't understand the problem?



**Lopa** June 14, 2019 at 2:14 am #

**Start Machine Learning**

Also do I have to convert the new series to supervised format in order to get the predictions?



**Jason Brownlee** June 14, 2019 at 6:50 am #

REPLY ↩

Typically, yes.



**skyrim4ever** July 5, 2019 at 9:22 pm #

REPLY ↩

Hello. I am trying to understand the code implementation. I find few of these defined functions hard to grasp but maybe I need take more time understanding these reshape lines that you used for variables 'train' and 'test'.

```
train = train.reshape(train.shape[0], train.shape[1])
test = test.reshape(test.shape[0], test.shape[1])
```

The variables do not change their shape at all am I right?



**Jason Brownlee** July 6, 2019 at 8:35 am #

Perhaps start with a much simpler implementation:  
<https://machinelearningmastery.com/how-to-develop-a-lstm-time-series-forecasting-model-python/>



**Roma** July 12, 2019 at 10:34 am #

REPLY ↩

We are using the last month's shampoo sales to predict the next month's sales.

1. Why do we need an LSTM to do this? I don't understand how the time component of LSTM plays into this.
2. When we set statefulness to True, I think the network is accounting for what came before with when it reads in a new month's sales. But then how is that info used in predict, when we have just 1 datapoint (and not a series as predictors)?



**Jason Brownlee** July 13, 2019 at 6:49 am #

REPLY ↩

An LSTM is not needed, in fact it is worse than a linear method. It is an example you might want to use as a starting point.

Statefulness means we control when the internal state is reset, instead of at the end of every batch/weight update.

[Start Machine Learning](#)



**psychologie** August 20, 2019 at 2:29 am #

REPLY ↗

What's up, after reading this awesome article i am as well glad to share my knowledge here with friends.



**Jason Brownlee** August 20, 2019 at 6:27 am #

REPLY ↗

Thanks.



**Dillan Muthanna** September 10, 2019 at 10:15 pm #

How to apply the same method when we I am trying to solve a forecasting problem at my interview on a day basis. I would really really appreciate you



**Jason Brownlee** September 11, 2019 at 10:15 pm #

Great question, you can adapt the suggestion at <https://machinelearningmastery.com/faq/single-faq/how-do-i-calculate-accuracy-for-regression>



**Vivaka Nand** October 11, 2019 at 5:56 pm #

REPLY ↗

Hi Sir,

your blogs are very informative.

my question is how can i find the accuracy of this model?

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** October 12, 2019 at 6:50 am #

REPLY ↗

You cannot calculate accuracy for regression, you must calculate error, more details here:

<https://machinelearningmastery.com/faq/single-faq/how-do-i-calculate-accuracy-for-regression>



**Ricardo** October 21, 2019 at 12:49 pm #

REPLY ↗

Jason, thank you for your post, I enjoyed it, very clear and to the point. I have a question about GPU vs CPU performance. I initially run the example on my CPU, when I noticed it was taking too long I installed tensorflow-gpu and used a Nvidia RTX2070 card instead, but still didn't see much improvement. As far as I know Im using the GPU but it still takes

[Start Machine Learning](#)

causes or how can I improved the performance? I have worked before with CNN and Pytorch and noticed a great difference in the performances of CPU and GPU, thus, I don't think is an issue with the card.

Thanks again,



**Jason Brownlee** October 21, 2019 at 1:43 pm #

REPLY ↗

Yes. Performance of LSTMs is typically not improved with GPUs.

You can try exploring some of unrolling parameters that will take my more memory but improve execution.



**Mark Hendrix** November 6, 2019 at 11:29 pm #

Hello Jason,

I have a question, How can i predict the next 12 months input?

I hope you can help me!



**Jason Brownlee** November 7, 2019 at 6:10 pm #

This is a common question that I answer here:

<https://machinelearningmastery.com/faq/single-faq/how-do-you-use-lstms-for-multi-step-time-series-forecasting>

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Malcolm** November 10, 2019 at 4:25 pm #

REPLY ↗

Hi, have been trying to plot loss function; ive added these lines of code at line 126. The Code is as follows:

```
training_loss = lstm_model.history.history['loss']
epoch_count = range(1, len(training_loss) + 1)

pyplot.plot(epoch_count, training_loss, 'r-')
pyplot.legend(['Training Loss', 'Test Loss'])
pyplot.xlabel('Epoch')
pyplot.ylabel('Loss')
pyplot.show()
```

However, i keep getting a blank loss function; How can i edit the previous/ these lines of code to plot the loss function?

## Start Machine Learning



**Jason Brownlee** November 11, 2019 at 6:05 am #

REPLY ↗

Perhaps confirm that the history object contains the data you expect?



**vusal** November 21, 2019 at 1:05 pm #

REPLY ↗

Hi Jason,

I attempted to modify the lag size (e.g., window size to 3) in the above code but I faced the ValueError: cannot reshape array of size of...

The code seems to be pretty adaptable. Do you kn...

Vusal



**Jason Brownlee** November 21, 2019 at 1:05 pm #

Not sure off hand, you may need to de...



**Tanushree Banerjee** November 27, 2019 at 11:42 am #

Hi, can I use this kind of LSTM model for...

Training Phase: I'll supply input and output signal, I...

Test case: I'll provide input signal and I want the model to predict the output signal for me based on what it has learnt in training phase



**Jason Brownlee** November 28, 2019 at 6:35 am #

REPLY ↗

That sounds like a great project.

Perhaps try the LSTM and compare the results?



**Sushanth** December 12, 2019 at 11:42 am #

REPLY ↗

Hi Jason

I have gone through many of your articles especially LSTM. There are predictions based on testing set where we have values till t-1 & we predict 't' value. But I would like to understand how do we predict for 1 month or 10 days of future?

I have tried feeding the predicted value as input to the model, but this is leading to a flat output, variance is getting lost.

[Start Machine Learning](#)



**Jason Brownlee** December 12, 2019 at 1:44 pm #

REPLY ↗

See these examples:

<https://machinelearningmastery.com/faq/single-faq/how-do-you-use-lstms-for-multi-step-time-series-forecasting>



**Femi** December 13, 2019 at 1:06 pm #

REPLY ↗

Thanks greatly for this article. What about if I want to predict a future occurrence that I don't have data for based on this model.



**Femi** December 13, 2019 at 1:08 pm #

I mean, forecast for future dates not ir



**Jason Brownlee** December 13, 2019 at 1:10 pm #

Call the predict() function with the input



**Femi** December 13, 2019 at 4:59 pm #

There are no inputs (because it is univariate). I only have the date index with "nan" values in place of my target values since the events have not occurred. To make predictions by splitting data into train and test is easy because there are available for that period.

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** December 14, 2019 at 6:06 am #

REPLY ↗

Not quite.

Time series framed as a supervised learning problems has inputs and outputs, e.g. the model takes some number of prior obs and uses it to predict future obs.

This applies to all forecasting models.

More on this here:

<https://machinelearningmastery.com/time-series-forecasting-supervised-learning/>



**Femi** December 13, 2019 at 5:00 pm #

REPLY ↗

projecting beyond my train and test data is

Start Machine Learning



**Jason Brownlee** December 14, 2019 at 6:06 am #

REPLY ↗

Yes, I understood and answered accordingly.

Perhaps start with this to understand the basics:

<https://machinelearningmastery.com/start-here/#timeseries>



**Femi** December 25, 2019 at 7:29 am #

REPLY ↗

I understand this procedure. Using want to predict in the and this X is actually time-series explains this. This means if I am the results that I am trying to predict.



**Jason Brownlee** December 25, 2019 at 10:00 am #

No. Outcomes are only needed to make a prediction, even a multiple step prediction.  
You have choice over how to frame the available at prediction time.

## Start Machine Learning



You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Femi** December 25, 2019 at 12:04 pm #

REPLY ↗

I forgot to mention I'm dealing with a univariate dataset. So only inputs I have are dates, which will be my index. This why I'm asking though because with univariate I can't seem to hack my way around what my input (X) will be for non-sample-dataset.



**Jason Brownlee** December 26, 2019 at 7:32 am #

REPLY ↗

The index/dates/times are dropped and you only work with the series of observations.



**Thefenix** January 12, 2020 at 8:46 pm #

REPLY ↗

Hello, I just follow your instruction and code in <https://machinelearningmastery.com/how-to-develop-lstm-models-for-time-series-forecasting/> (Multivariate LSTM Models- Multi input). And I set step size =3, feature =3, predict 1 output in 3 day forward. From the result, the predicted\_output was plotted look like(a little bit different) actual\_output, but it was shifted forward(delay) about 1-3 days. Is there anything wrong? or I have to add reset\_state like the code (Multivariate multi input vs. this post)? (In LSTM)

Start Machine Learning



**Jason Brownlee** January 13, 2020 at 8:20 am #

REPLY ↗

I believe the model learned a persistence forecast:

<https://machinelearningmastery.com/faq/single-faq/why-is-my-forecasted-time-series-right-behind-the-actual-time-series>



**KHADIJA AIT DERHEM** February 23, 2020 at 5:32 am #

REPLY ↗

le premier exemple comment je peut l'executer  
parser(x)  
j'ai pas arrivé à trouver malgré qu'ai enregistrer les



**Jason Brownlee** February 23, 2020 at 7:00 am #

You may not need it if your date-times parse the dates for you.



**KHADIJA AIT DERHEM** February 24, 2020 #

thank you for your answer now I have another problem is that if I put for the executed it gave me 'Data\_Digital\_Currencies .csv' does not exist: b'Data\_Digital\_Currencies .csv'

hat I save all the data in the same folder

```
from pandas import read_csv
from pandas import datetime
from matplotlib import pyplot
load dataset
def parser(x):
 return datetime.strptime('%Y-%m')
series = read_csv('Data_Digital_Currencies .csv', header=0, parse_dates=[0], index_col=0,
squeeze=True)
summarize first few rows
print(series.head())
line plot
series.plot()
pyplot.show()
```



**Jason Brownlee** February 25, 2020 at 7:46 am #

REPLY ↗

[Start Machine Learning](#)

Ensure the data file is in the same directory as your code file, or specify the full path to the file.



**Kavya** March 11, 2020 at 7:00 am #

REPLY ↗

Hello Jason,

I would like to know is it possible to predict multi outputs time series from multi input time series data using LSTM( like if i have 10 input attributes and 4 output attributes in time series). Is there any of your blog that explains about this kind of modeling.



**Jason Brownlee** March 11, 2020 at 8:04

Yes, you can start here for simple examples:  
<https://machinelearningmastery.com/how-to-develop-a-skilful-time-series-forecasting-model/>

More advanced examples here:

<https://machinelearningmastery.com/start-here/>



**Luca** March 15, 2020 at 2:40 am #

Hi Jason. Congratulations for your great site on stackoverflow (<https://stackoverflow.com/questions/50983797/time-series-forecasting-with-lstm>). Why when I predict on a test set the predictions are not stable? They oscillate between predictions or converged to a unique value or diverged? I tried many models (MLP, LSTM, GRU, ecc) . All of them cause the same problem. I think is because the models can't forecasting for many steps. What do you think? Thank you.

Luca.



**Jason Brownlee** March 15, 2020 at 6:20 am #

REPLY ↗

Yes, perhaps your model is overfit the training or does not have skill.

Perhaps try other model configs, other model types, etc. and discover what works best for your dataset. See this:

<https://machinelearningmastery.com/how-to-develop-a-skilful-time-series-forecasting-model/>



**Maunish** April 2, 2020 at 11:42 pm #

REPLY ↗

Hey Jason, excellent article

I have one doubt that can we scale the data before performing the shift() function because after shifting data gets converted to shape of (len(data), 1)

[Start Machine Learning](#)

and then at the time of prediction data is of shape (len(data),1)  
so we have to write a complicated invert\_scale function,

so can we scale before shifting so that it saves us from writing complicated code and just perform simple  
scaler.inverse\_transform(prediction)?



**Jason Brownlee** April 3, 2020 at 6:54 am #

REPLY ↗

Yes, things get very sticky!



**sarah** April 7, 2020 at 12:14 pm #

Thanks for your great site,  
Can I use the model,( with the same coding structure )  
Could you help me what should I consider to use a  
Another question, do you have any courses about learning methods?



**Jason Brownlee** April 7, 2020 at 1:30 pm #

Perhaps.

I recommend this process:

<https://machinelearningmastery.com/how-to-develop-a-skilful-time-series-forecasting-model/>

The tutorials here will help you to get started:

[https://machinelearningmastery.com/start-here/#deep\\_learning\\_time\\_series](https://machinelearningmastery.com/start-here/#deep_learning_time_series)

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**  
Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Erik** April 11, 2020 at 12:34 am #

REPLY ↗

Hey, thanks for the great help you've been!

I was wonder how to plot the train loss versus test loss for this particular example. I have looked at how you usually do it from other examples but I can't get it to work.

I am still kind of new to LSTM so it would be of great help if you could assist me.

Thanks  
Erik



**Jason Brownlee** April 11, 2020 at 6:22 am #

REPLY ↗

Start Machine Learning

It can be challenging for LSTMs, I expect you will either have to collect the history yourself with custom code or change the way the model is fit.



**Erik** April 13, 2020 at 12:11 am #

REPLY ↗

Thanks for the fast reply!

Yes you were right, I changed the way the model is fit and it works.  
Awesome work, and thanks for providing these guides!

Br  
Erik



**Jason Brownlee** April 13, 2020 at 6:19 am #

Well done!



**Meisam** April 14, 2020 at 7:23 am #

Hello,

Thanks for this great article.. I'm new in time series  
When I fit the model and test that, it will not predict  
price did change, it will adopt to change not predict it. For example, if day 1 expected price was 201, the  
day 2 price was 220 and the third day was 200, it will not predict the price rotation in day 3 and it will  
show it in day4... It's like the model copy the move or something... Please help me...I stuck in this and  
don't know what's the problem. Thanks



**Jason Brownlee** April 14, 2020 at 10:36 am #

REPLY ↗

I believe your model has learned a persistence model:

<https://machinelearningmastery.com/faq/single-faq/why-is-my-forecasted-time-series-right-behind-the-actual-time-series>



**Meisam** April 15, 2020 at 12:12 am #

REPLY ↗

Thank's so much... You made my day 😊



**Jason Brownlee** April 15, 2020 at 8:00 am #

REPLY ↗

You're welcome.

Start Machine Learning

**Meisam** April 15, 2020 at 8:07 am #

REPLY ↩

Hi

Sorry for bothering again... But I'm still stuck on that... I couldn't get rid of persistent forecasting even with changing data or the hyper parameters... I don't know what to do...please help me...

( Also another question... Can we use GRU to predict stock market, because I didn't run into any article that include GRU for stock prediction )

**Jason Brownlee** April 15, 2020 at 1:18 pm #

Perhaps explore using an alternate m

**Randall Kakaire** April 25, 2020 at 8:09 pm #

Hello,

could you provide example on how to increase the example instead of just using the previous value, i values.

## Start Machine Learning

X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE

REPLY ↩

**Jason Brownlee** April 26, 2020 at 6:09 am #

This requires a change to the model, called multi-step forecasting:

<https://machinelearningmastery.com/faq/single-faq/how-do-you-use-lstms-for-multi-step-time-series-forecasting>

**Jagadeesh** May 19, 2020 at 11:48 am #

REPLY ↩

Great tutorial Jason.

I was planning to predict the values beyond the test data and got stuck in a confused state. It will be really helpful if you can provide some code snippets on predicting future unseen values, thanks

**Jason Brownlee** May 19, 2020 at 1:26 pm #

REPLY ↩

Yes, see this:

<https://machinelearningmastery.com/make-sample-forecasts-arima-python/>

## Start Machine Learning

**Jagadeesh** May 19, 2020 at 11:03 pm #

REPLY ↗

That's great Jason. Thank you for the link.

Can you please provide any links related to LSTM unseen data forecast. Thank you.

**Jason Brownlee** May 20, 2020 at 6:26 am #

REPLY ↗

Yes, here:

<https://machinelearningmastery.com/make-predictions-long-short-term-memory-models-keras/>

**younes** May 25, 2020 at 10:38 pm #

Great tutorial Jason.  
I have a question.

do we apply the difference() on all the data then we  
or first split the data into trainset and testset and then

**Jason Brownlee** May 26, 2020 at 6:22 am #

There's no right answer, perhaps choosing  
and does not lead to data leakage.

**RAYENNE AMMAR** June 1, 2020 at 11:52 am #

REPLY ↗

Hi Jason, Thanks a lot for this great article. I have a question on the Persistence Model if it  
works like the naif model as they both require the previous observation to predict the next one?

**Jason Brownlee** June 1, 2020 at 1:42 pm #

REPLY ↗

You're welcome.

What is "naif"?

**RAYENNE AMMAR** June 1, 2020 at 11:14 pm #

REPLY ↗

I'm sorry, I meant naive model

[Start Machine Learning](#)



**Pablo** June 23, 2020 at 11:01 am #

REPLY ↗

Hi Jason, thanks for the great post. My question is, does it make sense to use LSTM when looking only one step back? Or is just for example purpose.

Thanks



**Jason Brownlee** June 23, 2020 at 1:28 pm #

REPLY ↗

Probably not. Generally, I recommend testing a suite of models and discover what works best for your dataset.



**Marco** July 7, 2020 at 6:18 am #

Hi Jason, many thanks for your brilliant post. I have a few questions I might ask:

1. In feature scaling, when you do `scaler.fit(train)`, are we fitting the scaler on the training data?
2. You mentioned that to achieve a better level of accuracy, we can use an LSTM model. How do you find the best configuration for an LSTM model?
3. How would you add a Dropout layer (in coding terms)?

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** July 7, 2020 at 6:48 am #

REPLY ↗

Thanks!

When we fit the scaler on training data, we estimate the properties required by the scaler from the training dataset. We can then use the scaler to transform the training dataset, test dataset, and other datasets:

<https://machinelearningmastery.com/standardscaler-and-minmaxscaler-transforms-in-python/>

The only way to find a good/best config for a model is to use controlled experiments and test many hyperparameter configurations.

Here's an example of adding dropout for LSTMs:

<https://machinelearningmastery.com/use-dropout-lstm-networks-time-series-forecasting/>



**Marco** July 7, 2020 at 7:26 am #

REPLY ↗

Many thanks for your reply.

I get the testing many hyperparameters, but perhaps do you have an example on how to conduct controlled experiments?

[Start Machine Learning](#)

Thanks!



**Jason Brownlee** July 7, 2020 at 1:56 pm #

REPLY ↩

Yes, you can see many examples on the blog, perhaps start here:

<https://machinelearningmastery.com/how-to-grid-search-deep-learning-models-for-time-series-forecasting/>



**Thomas** July 7, 2020 at 9:08 am #

REPLY ↩

Hey Jason, your website is dope!

I have tried this model with a larger data (3774 rows) predictions, they are just like the raw data! Does this in regards? cheers!



**Jason Brownlee** July 7, 2020 at 2:00 pm #

Thanks.

I recommend testing a suite of models, model selection in order to discover what works best for your specific problem.

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jacob** July 7, 2020 at 6:03 pm #

REPLY ↩

Hello Jason, My question is, if I want to try different learning rates, do I just have to put it next to adam? e.g., optimizer='adam(lr=n)'

Thank you!



**Jason Brownlee** July 8, 2020 at 6:28 am #

REPLY ↩

You must use SGD to test different learning rates.



**James** July 19, 2020 at 8:35 pm #

REPLY ↩

Hi Jason, great tutorial!

I am struggling to implement your differencing (and inverse difference) method on multivariate data. It works when I have only one column of data, but how do I modify it when I have more than one column?

Thanks and keep doing this amazing work!

Start Machine Learning



**Jason Brownlee** July 20, 2020 at 6:11 am #

REPLY ↗

You could try the transforms one variable at a time, then once you're comfortable try simplifying the code to do it all at once.



**James** July 20, 2020 at 11:16 pm #

REPLY ↗

Thank you for your reply Jason.

I have figured out how to differentiate all variables.

Then I know how to inverse it:

e.g.,

```
data_differentiated.iloc[0] = normal_data.iloc[0]
data_inverted_back = data_diff.cumsum()
```

However, I am struggling to implement this having scaled back. Do you think you could help? Thanks



**Jason Brownlee** July 21, 2020

Maybe this will help:

<https://machinelearningmastery.com/machine-learning-data-transforms-for-time-series-forecasting/>

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**James** July 21, 2020 at 10:22 pm #

Thank you for your response.

I have tried to follow this post, but when I use my multivariate data I get , which relates to the interval but I do not know how to correctly modify it. Any advise ?

Thanks!



**Jason Brownlee** July 22, 2020 at 5:31 am #

Sorry, I don't have the capacity to prepare custom code for you. Perhaps you can learn from other examples on the blog, there are hundreds on this topic, perhaps start here:

[https://machinelearningmastery.com/start-here/#deep\\_learning\\_time\\_series](https://machinelearningmastery.com/start-here/#deep_learning_time_series)

Start Machine Learning

**James** July 22, 2020 at 5:25 am #

# I get

**Sanat** July 23, 2020 at 4:33 pm #

REPLY ↩

Hi Jason,

I have a dataset of customers\_number ,sales and date. I have used LSTM and ARIMA both worked for whole dataset but I want to find it customer wise so how can I do it? Is there another way? Can you please enlighten me in this issue?

**Jason Brownlee** July 24, 2020 at 6:23 am #

Sorry, I don't follow your question, per

**dis** November 16, 2020 at 8:18 am #

Hi Jason,

Thanks for the detailed and well-explained post. I would like to know if it's possible to predict help prediction in the LSTM model? Or is it possible to predict help prediction in the LSTM model? Thank you.

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)
**Jason Brownlee** November 16, 2020 at 8:51 am #

REPLY ↩

Yes, see examples here:

<https://machinelearningmastery.com/how-to-develop-lstm-models-for-time-series-forecasting/>

**Dr D Niyogi** December 5, 2020 at 7:14 am #

REPLY ↩

Hi Jason,

Excellent post! Thank you!!!

I have a question: suppose I have the following code:

```
fit the model
lstm_model = fit_lstm(train_scaled, 1, 3000, 4)
forecast the entire training dataset to build up state for forecasting
train_reshaped = train_scaled[:, 0].reshape(len(train_scaled), 1, 1)
train_predicted = lstm_model.predict(train_reshaped)
```

[Start Machine Learning](#)

where `train_predicted` will have the predictions for the training data. Now we have to invert scale and invert differentiation. The routines you have provided (`invert_scale` and `inverse_difference`) cannot be applied for “`train_predicted`”. Could you please show me how to invert scale and invert differentiate “`train_predicted`”?

Many thanks in advance!



**Jason Brownlee** December 5, 2020 at 8:15 am #

REPLY ↗

You're welcome.

Yes, see this:

<https://machinelearningmastery.com/machine-learning-forecasting-stateful-lstm/>



**Joseph** December 25, 2020 at 2:52 am #

Hey,

I've seen examples where you don't reset the state.

Thanks.



**Jason Brownlee** December 25, 2020 at 5:15 am #

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Jason Brownlee** December 25, 2020 at 5:15 am #

I recommend testing a suite of algorithms and algorithm configurations and discover what works best for your specific dataset.

In practice, I have not seen much difference between stateful and stateless LSTMs. Perhaps I'm too heavy-handed with my model training.



**Joseph** December 25, 2020 at 5:37 am #

REPLY ↗

Hey Jason,

You used a long loop (walk-forward) to predict each input, when you could've just given the predict function all the samples, and it seems to loop through them and give the predictions. Whether loop was used to go through each input of the test or not, the results seem to be the same. Correct me if I'm wrong.



**Jason Brownlee** December 25, 2020 at 5:46 am #

REPLY ↗

Perhaps, as long as the data is not shuffled.

Start Machine Learning



SRan January 23, 2021 at 8:44 am #

REPLY ↗

Hey Jason,

Thanks a lot for your post. I want to make predictions for the next 12 months, so in my test data, I only have months, but no sales value. It can only prediction for the first month, and the rest shows NAN. So seems like it also use expected value as input. So how to predict for the next 12 steps instead of only 1 step?



Jason Brownlee January 23, 2021 at 12:49 pm #

REPLY ↗

Good question.

One approach is to use the same model and use the output from one step as input for the next step.

Another approach is to change the model to predict multiple steps at once.

Perhaps try both approaches and see what works best.

This will help you get started:

<https://machinelearningmastery.com/faq/single-step-time-series-forecasting/>

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



SRan January 25, 2021 at 4:05 am #

Hi Jason,

Thanks a lot for your reply. I went through all the comments from other people. Seems like many people have the same question as I do. I wonder if you can post an example to predict next many steps (not the test data). There are at least two problems I am suffering with (many other people as well): 1. how to keep using the output as input to continuously predict 2. how to scale the predicted data and inverse scale it (what if the future data is not in the scale of the training data? For example, I would expect the total Covid number keeps growing, the number I want to predict might be hundred times larger than the training data. How to do the scale?) I would really appreciate if you can post such an example. That will help to answer many people's questions. Thanks!



Jason Brownlee January 25, 2021 at 5:53 am #

REPLY ↗

Thanks for the suggestion.

There are tens of examples of LSTMs for multi-step prediction on the blog. I recommend starting with them.

If you are having trouble with scaling the data, this will help:

<https://machinelearningmastery.com/machine-learning-data-transforms-for-time-series-forecasting/>

Start Machine Learning

**SRan** January 25, 2021 at 7:36 am #

Thanks, Jason. I am aware of your nice posts on the multi-step predictions. But they all work on test data. I think what we are looking for is an example to predict for a period of time in the future, either one step or multi step methods.

**Jason Brownlee** January 25, 2021 at 7:48 am #

It is exactly the same code to predict on a test set vs out of sample data.

If making a prediction on data is a new idea for you, this may help:

<https://machinelearningmastery.com/predictions-for-deep-learning-models/>

**Shervin** February 26, 2021 at 2:51 am #

Hi Jason.

I have rather simple question.

Where do we actually define number of cells of each

I haven't found the answer anywhere in your tutorial.

Thanks in advance.

## Start Machine Learning

You can master applied Machine Learning without math or fancy degrees.

Find out how in this free and practical course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)

[REPLY](#) ↩

**Jason Brownlee** February 26, 2021 at 5:03 am #

An LSTM layer has a number of units, like "neurons" or "nodes". Cells and blocks are no longer discussed.

**Shervin** February 26, 2021 at 8:08 am #

[REPLY](#) ↩

Thanks a lot for reply.

Do you have any tutorial on the mentioned difference or any useful links discussed?

If not, so how do these nodes interact in this case?

I fully appreciate your assistance.

**Shervin** February 26, 2021 at 8:13 am #

[REPLY](#) ↩

In this case, do we act each LSTM layer as single block consisting cells with given number of neurons (in terms of blocks and cells)?

[Start Machine Learning](#)



**Jason Brownlee** February 26, 2021 at 1:26 pm #

REPLY ↗

The nodes do not interact.

Perhaps this will help:

<https://machinelearningmastery.com/faq/single-faq/how-is-data-processed-by-an-lstm>

And this:

<https://machinelearningmastery.com/faq/single-faq/what-is-the-difference-between-samples-timesteps-and-features-for-lstm-input>

And these:

<https://machinelearningmastery.com/start-here/#lstm>

## Start Machine Learning

X



**Shervin** February 26, 2021 at 1:26 pm #

Thanks so much Jason.



**Jason Brownlee** February 26, 2021 at 1:26 pm #

You're welcome.

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address



I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Belle** March 3, 2021 at 1:44 pm #

Hi Jason, I have a question...

How to fix seed ?

REPLY ↗



**Jason Brownlee** March 3, 2021 at 1:57 pm #

Tough question, see this:

<https://machinelearningmastery.com/reproducible-results-neural-networks-keras/>



**ronak** March 10, 2021 at 8:04 pm #

REPLY ↗

Hi Jason, i have a question... I am new on Machin learning and Python also Anaconda...

I tried your first example in machin learning (iris flowers) and it was easy to follow and apply the commands on the Anaconda powershell prompt.

But now for the time-series forecasting, at the very beginig stage i have a problem with opening .CSV file and define it as the variable (series). i have saved it as you said in the my current diroctory and i see it clearly when i search in the prompt. i have no error `print(series.head())=>` here i get this error:

Start Machine Learning

```
>>> # summarize first few rows
>>> print(series.head())
Traceback (most recent call last):
File "", line 1, in
NameError: name 'series' is not defined
```

I do not know why it did not define the series variable containing shampoosales.csv values!

It is good to mention that I have tried to enter my CSV values using URL and it worked to show the head of dataset.

Thank you for your response in advance,  
Ronak



**Jason Brownlee** March 11, 2021 at 5:10

The error suggests you might have skipped a step.

Perhaps save the code in a .py script file and run it from the command line.  
<https://machinelearningmastery.com/faq/single-threaded-code-examples/>



**Mark** July 28, 2021 at 8:36 pm #

Hi Jason,

Thank you for the tutorial.

Is there any way you can explain (with code) on how to use this example to predict the sales value on the x future days instead of data you already know (test set)?

Mark.

## Start Machine Learning

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**



**Jason Brownlee** July 29, 2021 at 5:11 am #

REPLY ↗

Yes, call model.predict().

See this:

<https://machinelearningmastery.com/make-predictions-long-short-term-memory-models-keras/>

And this:

<https://machinelearningmastery.com/how-to-make-classification-and-regression-predictions-for-deep-learning-models-in-keras/>



**Mark** July 29, 2021 at 7:45 am #

REPLY ↗

Hi Jason,

Thank you for the extra links and sorry for my inexp

## Start Machine Learning

What I did was this:

1) I split the dataset in half (train + test)

2) I trained the model with train data and used the test data to obtain a result:

```
testPredict = model.predict(testX)
```

This was the result: <https://ibb.co/7pGQXXL>

The model looked ok. So I went to re-do the code to be able to predict future values.

3) Predicting new values: I built a function that would grab the history (half of the dataset) and predict a new value, via the model trained above, than grab that value, add it to the history and repeat, where steps\_ahead is the total prediction I wanted to do.

```
for i in range(steps_ahead):
```

```
Predict next step
```

```
prediction = self.perform_prediction(history)
```

```
Store prediction
```

```
model_predictions.append(prediction)
```

```
Add prediction to history
```

```
next_index = history.tail(1).index.item() + timedelta
```

```
history.loc[next_index, 'Series name'] = round(pred
```

Unfortunately this was the result: <https://ibb.co/2hX>

Q1) I am at the point of total confusion. In 2) when exactly? The model will run once (internally) for each

Q2) The above question takes me to a new question: a future value based on the past/present. From this provides a result which doesn't really tell anything

Any help is appreciated.

Sorry for the big text.

Mark



**Jason Brownlee** July 30, 2021 at 6:24 am #

REPLY ↗

Yes, the model runs once for each sample. If you're using an LSTM, it will process a batch of samples at a time and share internal state across samples in the batch.

Generally we evaluate a forecast model using walk-forward validation:

<https://machinelearningmastery.com/backtest-machine-learning-models-time-series-forecasting/>



**Manuel** August 6, 2021 at 6:26 am #

REPLY ↗

Hi Jason, I have to predict taking a sample from the training set as parameter, how can I perform the inverse difference? I don't know how to do the inverse difference to only one sample.

Thank you.

[Start Machine Learning](#)



**Jason Brownlee** August 7, 2021 at 5:35 am #

REPLY ↗

This tutorial has an example of calculating the difference and the inverse:

<https://machinelearningmastery.com/machine-learning-data-transforms-for-time-series-forecasting/>



**Manuel** August 13, 2021 at 10:01 pm #

REPLY ↗

Hi, I understand how to difference and perform the inverse to the whole dataset like that tutorial shows.

But I don't know how to perform the inverse difference. What do I have to do exactly in that case.

Thanks.



**Adrian Tam** August 14, 2021 at 3:30 pm #

I believe you are talking about the difference provided as  $(y_1 - y_0)$ ,  $(y_2 - y_1)$ , etc. And you can get the inverse difference, you know what to do. For example, if you get  $y_1$ , add  $(y_2 - y_1)$  to get  $y_2$ , etc.

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Manuel** August 16, 2021 at 6:28 am #

Hi, I understand how to perform the inverse difference, but I don't understand the difference between these two cases:

1. Making the inverse difference to the differenced values:

```
value = inverse_difference(series, differenced[i], len(series)-i)
```

(this appears in Transform Time Series to Stationary in this blog)

2. Making the inverse difference to the result of the prediction:

```
yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)
```

(this appears in the complete example of the LSTM in this blog)

I would like to know how to perform the inverse difference in the second case.

Thanks



**Adrian Tam** August 17, 2021 at 7:38 am #

If I understand you correctly, the situation is like this: We have data from time 1 to N known and trained the LSTM

[Start Machine Learning](#)

difference. So with history as input, we get the difference term only (assume predict for next-period only). And what is the value of N+1? We need to take  $x[N]$  and add the output to it, because  $yhat = x[N+1]-x[N]$ . What about N+2? If we assume  $x[N+1]$  as we obtained from inverse difference is true and feed into the model again, we get  $yhat = x[N+2]-x[N+1]$ , so we can find  $x[N+2]$  by inverse differencing again.

Hope this helps.



**Tom Whitehead** September 14, 2021 at 11:05 am #

REPLY ↗

Hi Jason, finding your tutorials on LSTM so helpful. I have one quick question on the stateful parameter. If the forecast horizon follows on directly training, is it better to maintain the state of the model state it has built up will be useful in making predictions?

If yes, could I simply rearrange the following code:

```
for i in range(nb_epoch):
 model.fit(X, y, epochs=1, batch_size=batch_size, v)
 model.reset_states()
```

to bring the `reset_states` method call before `fit`, so that it is maintained by the model:

```
for i in range(nb_epoch):
 model.reset_states()
 model.fit(X, y, epochs=1, batch_size=batch_size, v)
```

## Start Machine Learning X

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Adrian Tam** September 14, 2021 at 1:35 pm #

REPLY ↗

Your approach works. The problem of doing so is that you need to be very carefully to control when your training ends. Usually we don't. Therefore it is safer just reset the state and feed it some data (even it was used in training) to re-establish the state and do prediction.



**Tom Whitehead** September 14, 2021 at 4:20 pm #

REPLY ↗

Got it, thanks Adrian.



**Miya** September 17, 2021 at 7:46 am #

REPLY ↗

Hi Jason, I have a question, this example can predict the value at t+1. But if I want to predict the value from t+1 to t+5 , How should i change the code?

Start Machine Learning



**Adrian Tam** September 19, 2021 at 6:06 am #

REPLY ↗

Call `y = model.predict(x)` each time, get the output `y` and put it as the last row of `x`, then call `predict()` again. This is the simplest way to do multi-step predictions. But be careful that as you are using prediction output as input, the longer you predict, the less accurate it should be.



**Manuel Ruiz Molina** October 1, 2021 at 3:42 am #

REPLY ↗

Hi, can you help me to inverse the differencing on forecast? I don't know what values do I have to sum.

Thanks



**Adrian Tam** October 1, 2021 at 12:21 pm #

Not to do sum, but to do cumulative sum. Differencing is `np.diff(x)` and you can get back to `x[0]+y.cumsum()`



**Manuel** October 1, 2021 at 9:50 pm #

I understand that this solution is for a normal forecasting, but I am doing a one step forecasting like in this post, so in my case "y" will have only one element.

In this case, what is the difference?

I am having problems performing the inverse differencing when doing a one step forecasting, I would be very grateful if you would help me.

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Adrian Tam** October 6, 2021 at 7:06 am #

REPLY ↗

Your `y` in that case would be  $y=x(t+1)-x(t)$  which  $x(t+1)$  is what you're forecasting. Hence you need to find  $x(t)$  which is the last value in your input series, and add to your `y`



**Timirlan** November 9, 2021 at 1:18 am #

REPLY ↗

Jason, thank you very much for the lessons!

I have a time series db of car drivers aka bank checks for gasoline. How can a long break in purchase be diagnosed as leaving a gas station?

I would like to predict the time interval between pur

Start Machine Learning

Thanks.



**Adrian Tam** November 14, 2021 at 1:40 pm #

REPLY ↩

Not sure I understand correctly, but try to process the data to be number of times a driver visit a gas state and the time since the last visit. This might be easier for a model to predict.



**Timirlan** November 9, 2021 at 1:44 am #

REPLY ↩

in other words I have these series for buyer 1

2018-01-02 100  
2018-01-12 155  
2018-02-04 150  
2018-02-22 115

...

buyer 2

2019-11-01 111  
2018-11-21 98  
2018-11-29 122  
2018-12-03 150

...

if you are doing daily or weekly summing then there will be a series with spaces

## Start Machine Learning X

You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Adrian Tam** November 14, 2021 at 1:39 pm #

REPLY ↩

How about monthly summing?



**JOEP PELA** December 9, 2021 at 8:16 am #

REPLY ↩

Does differencing remove seasonality and if it doesn't how do I go about removing it.



**Adrian Tam** December 10, 2021 at 4:18 am #

REPLY ↩

If you talk about seasonality, you already assumed a model for the data. If you want to remove seasonality, you can do it by another model (e.g., using SARIMAX function from statsmodels package) before feeding into LSTM network.

Start Machine Learning



**James Carmichael** December 22, 2021 at 9:03 am #

REPLY ↗

Hi Joep...Differencing is a popular and widely used data transform for making time series data stationary.

Please review the following tutorial for more details.

[https://machinelearningmastery.com/remove-trends-seasonality-difference-transform-python/#:~:text=Differencing%20is%20a%20method%20of,%2C%20so%2Dcalled%20temporal%20dependence.&text=Differencing%20can%20help%20stabilize%20the,or%20reducing\)%20trend%20and%20seasonality.](https://machinelearningmastery.com/remove-trends-seasonality-difference-transform-python/#:~:text=Differencing%20is%20a%20method%20of,%2C%20so%2Dcalled%20temporal%20dependence.&text=Differencing%20can%20help%20stabilize%20the,or%20reducing)%20trend%20and%20seasonality.)



**Anupam Dungdung** April 9, 2022 at 4:02 am #

Hi Jason! Thanks for the wonderful tutorial do to make forecasts outside of the available data are for training and 12 is for testing the model. So I observation? (Ex. supposedly for next 12 months)



**James Carmichael** April 9, 2022 at 7:39 pm #

Hi Anupam...You would modify the nu

```
split data into train and test X = series.values
```

## Start Machine Learning

You can master applied Machine Learning **without math or fancy degrees**.

Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

START MY EMAIL COURSE



**Anupam Dungdung** April 10, 2022 at 5:49 pm #

REPLY ↗

Hello James. I am sorry but I couldn't quite get what you are trying to say.

## Leave a Reply

Name (required)

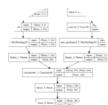
Start Machine Learning

Email (will not be published) (required)[SUBMIT COMMENT](#)**Welcome!**

I'm *Jason Brownlee* PhD  
and I **help developers** get results with **machine learning**.  
[Read more](#)

**Never miss a tutorial:****Picked for you:**

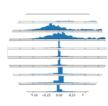
[How to Develop LSTM Models for Time Series Forecasting](#)



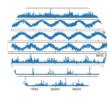
[How to Develop Convolutional Neural Network Models for Time Series Forecasting](#)



[Multi-Step LSTM Time Series Forecasting Models for Power Usage](#)



[1D Convolutional Neural Network Models for Human Activity Recognition](#)



[Multivariate Time Series Forecasting with LSTMs in Keras](#)

**Picked for you:**

[Simple Genetic Algorithm From Scratch in Python](#)



[Curve Fitting With Python](#)

## Start Machine Learning X

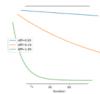
You can master applied Machine Learning  
**without math or fancy degrees.**

Find out how in this *free* and *practical* course.

 Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

[START MY EMAIL COURSE](#)



Simulated Annealing From Scratch in  
Python



How to Choose an Optimization Algorithm



Differential Evolution from Scratch in  
Python

## Start Machine Learning X

You can master applied Machine Learning **without math or fancy degrees.** Find out how in this *free* and *practical* course.

Email Address

I consent to receive information about services and special offers by email. For more information, see the [Privacy Policy](#).

**START MY EMAIL COURSE**

Loving th

The Deep Learning for Time Series EBook is where you'll find the *Really Good* stuff.

>> SEE WHAT'S INSIDE

---

© 2022 Machine Learning Mastery. All Rights Reserved.

[LinkedIn](#) | [Twitter](#) | [Facebook](#) | [Newsletter](#) | [RSS](#)

[Privacy](#) | [Disclaimer](#) | [Terms](#) | [Contact](#) | [Sitemap](#) | [Search](#)

**Start Machine Learning**